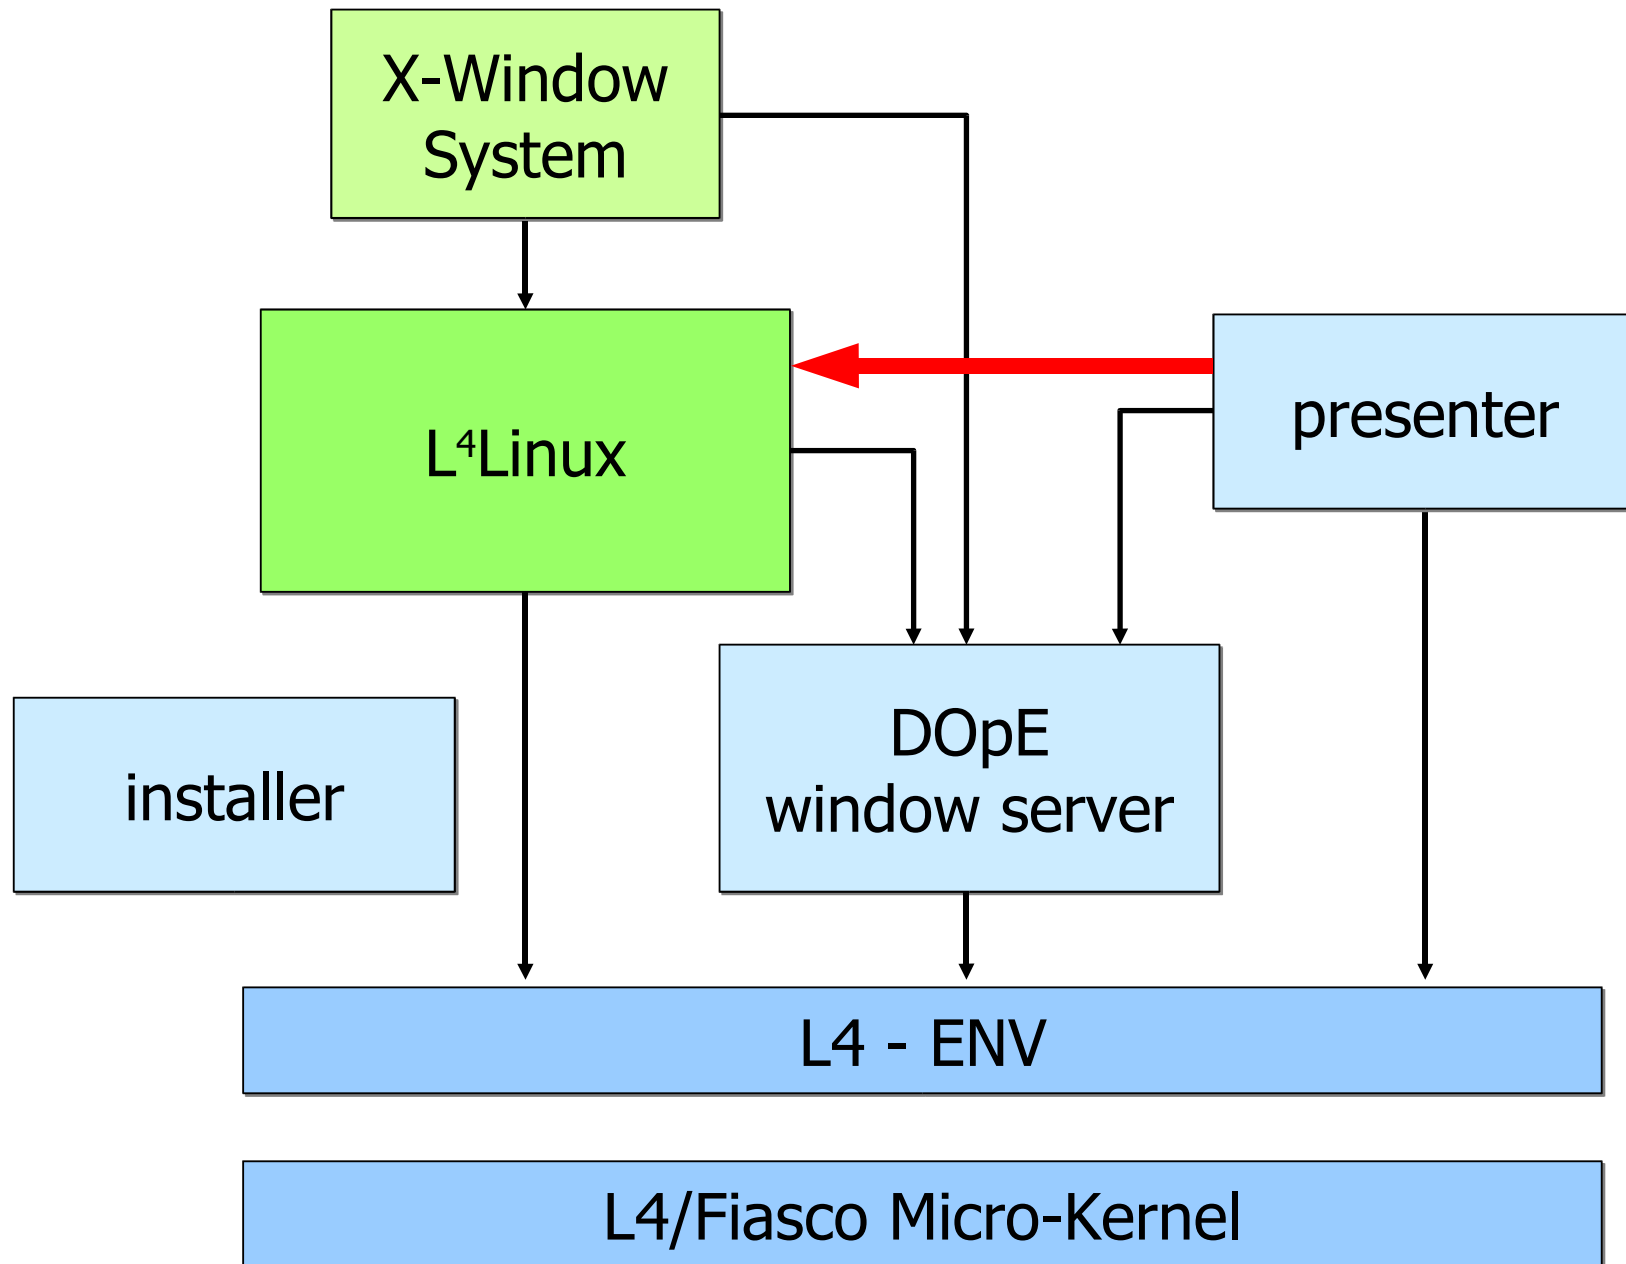
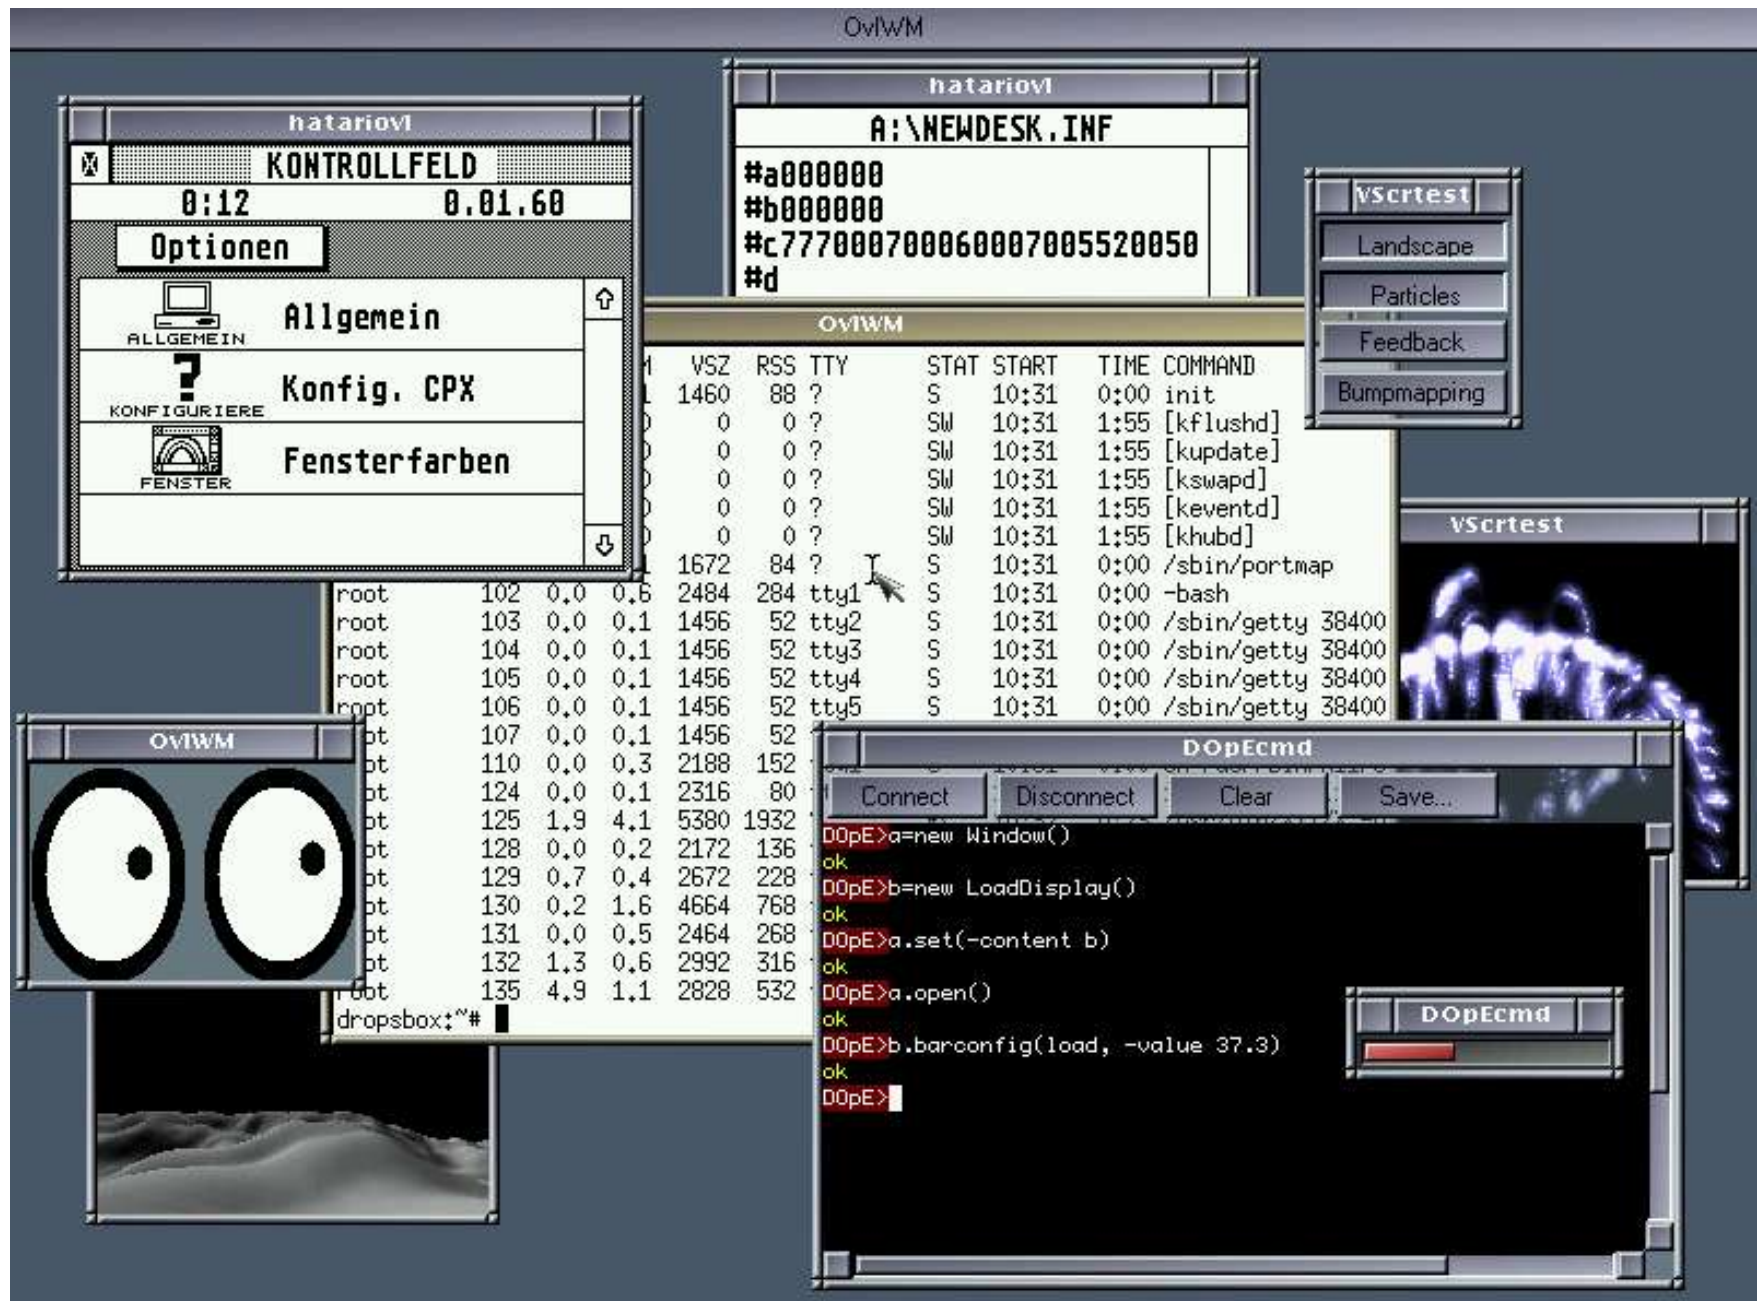


NIZZA

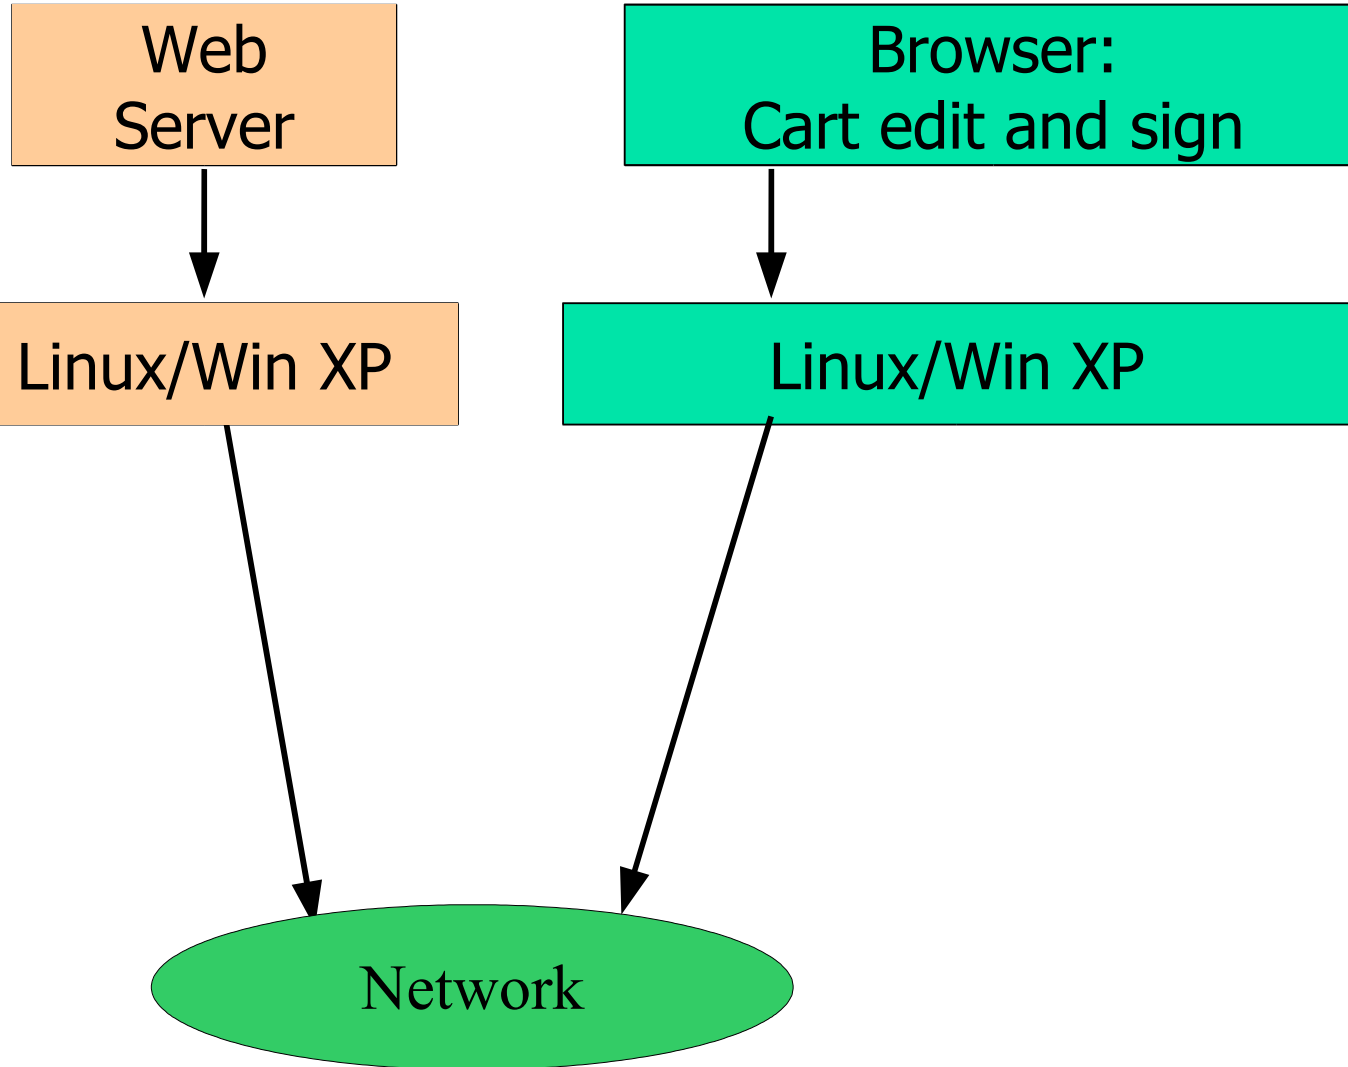
Hermann Härtig
TU Dresden ♦ OS
August 2004




Screen Shot



Internet Transaction



Your password(s), credit card number, ...



see:

Understanding Data Lifetime
via Whole System Simulation
Jim Chow, Ben Pfaff, Tal
Garfinkel, Kevin Christopher,
and Mendel Rosenblum,
Stanford University
Usenix Security 04

Lenin's Transaction

Web
Server

Browser:
Cart edit

**Cart
Sign**

Linux/Win XP

Linux/Win XP

Network

Lenin's Transaction

Web
Server

Browser:
Cart edit

**Cart
Sign**

Linux/Win XP

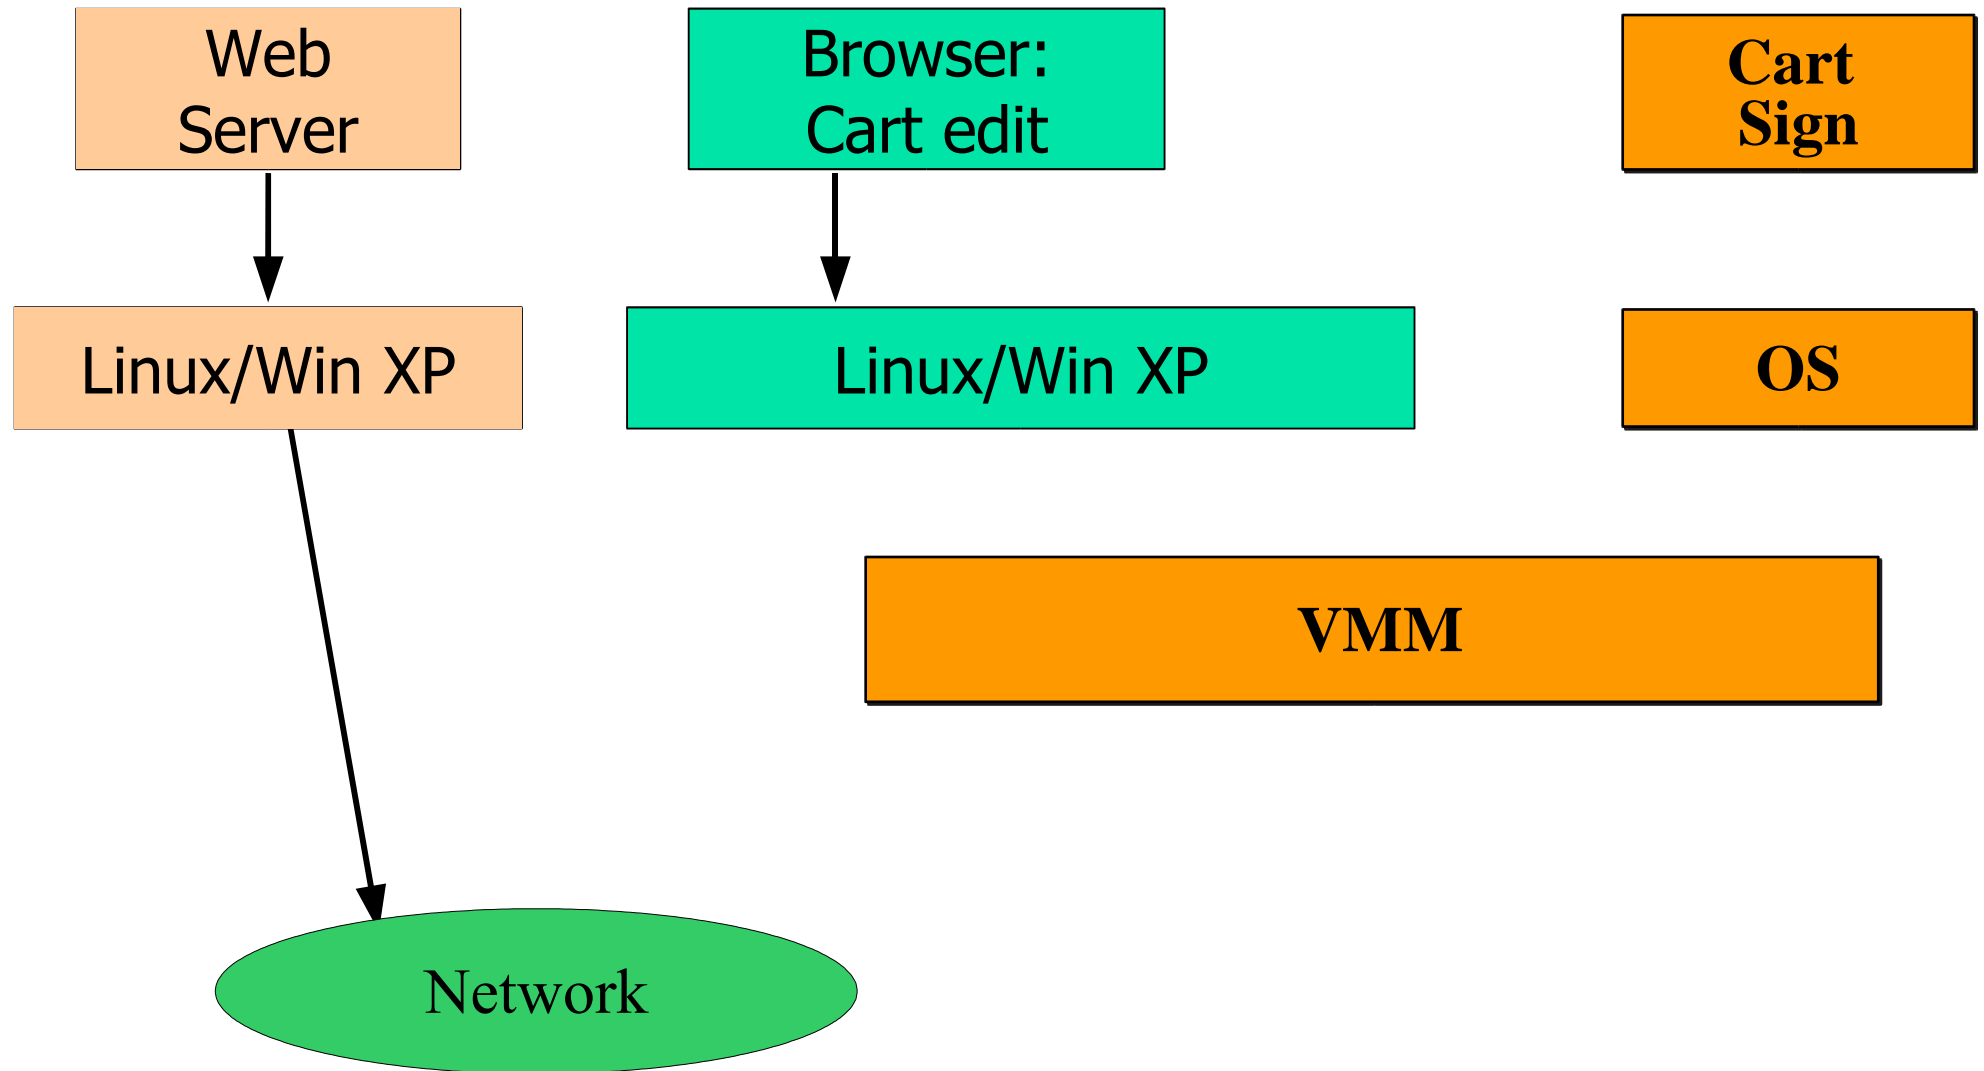
Linux/Win XP

OS

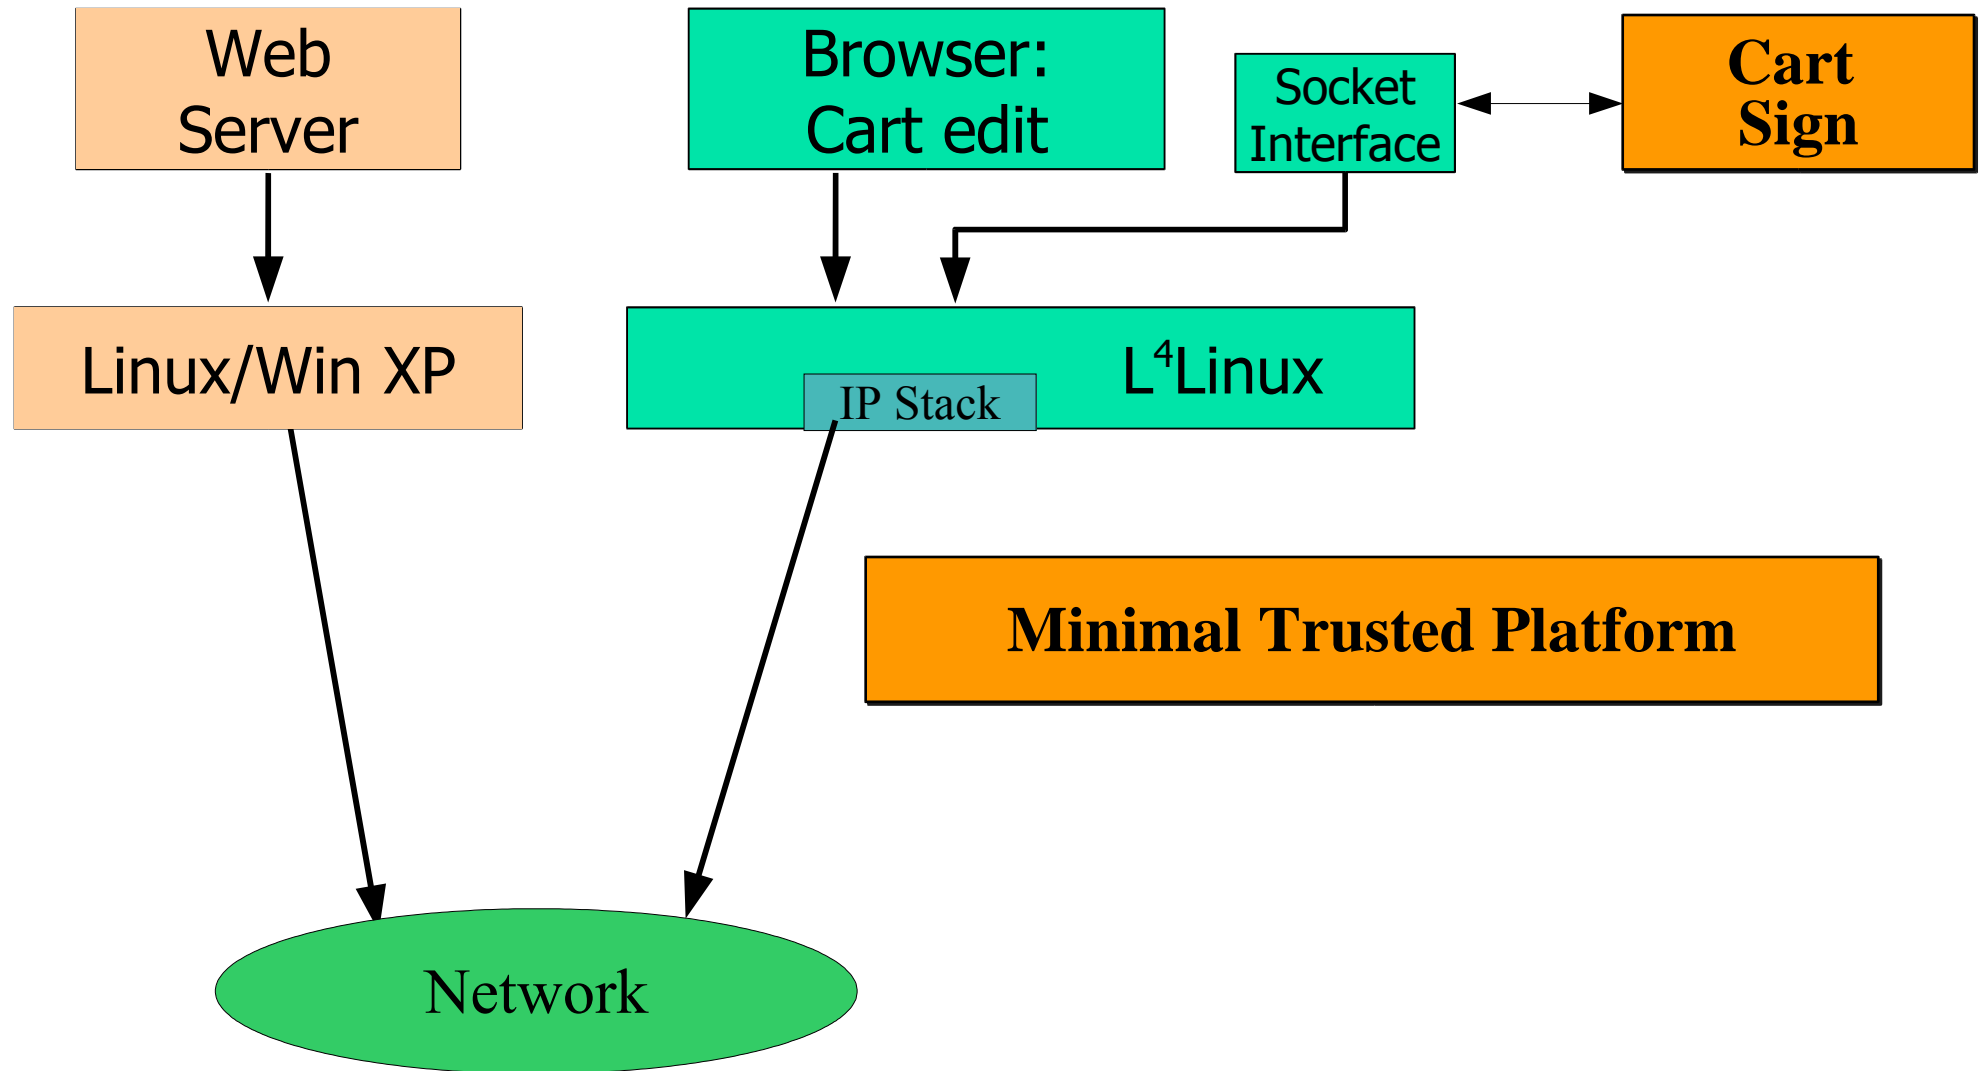
**dedicated
hardware**

Network

Lenin's Transaction



Lenin's Transaction



Lenin's Transaction

Untrusted

Web
Server

Browser:
Cart edit

Socket
Interface

**Cart
Sign**

Linux/Win XP

IP Stack

L⁴Linux

Minimal Trusted Platform

Network

Lenin's Transaction

Untrusted

Web
Server

Browser:
Cart edit

Socket
Interface

**Cart
Sign**

Linux/Win XP

IP Stack

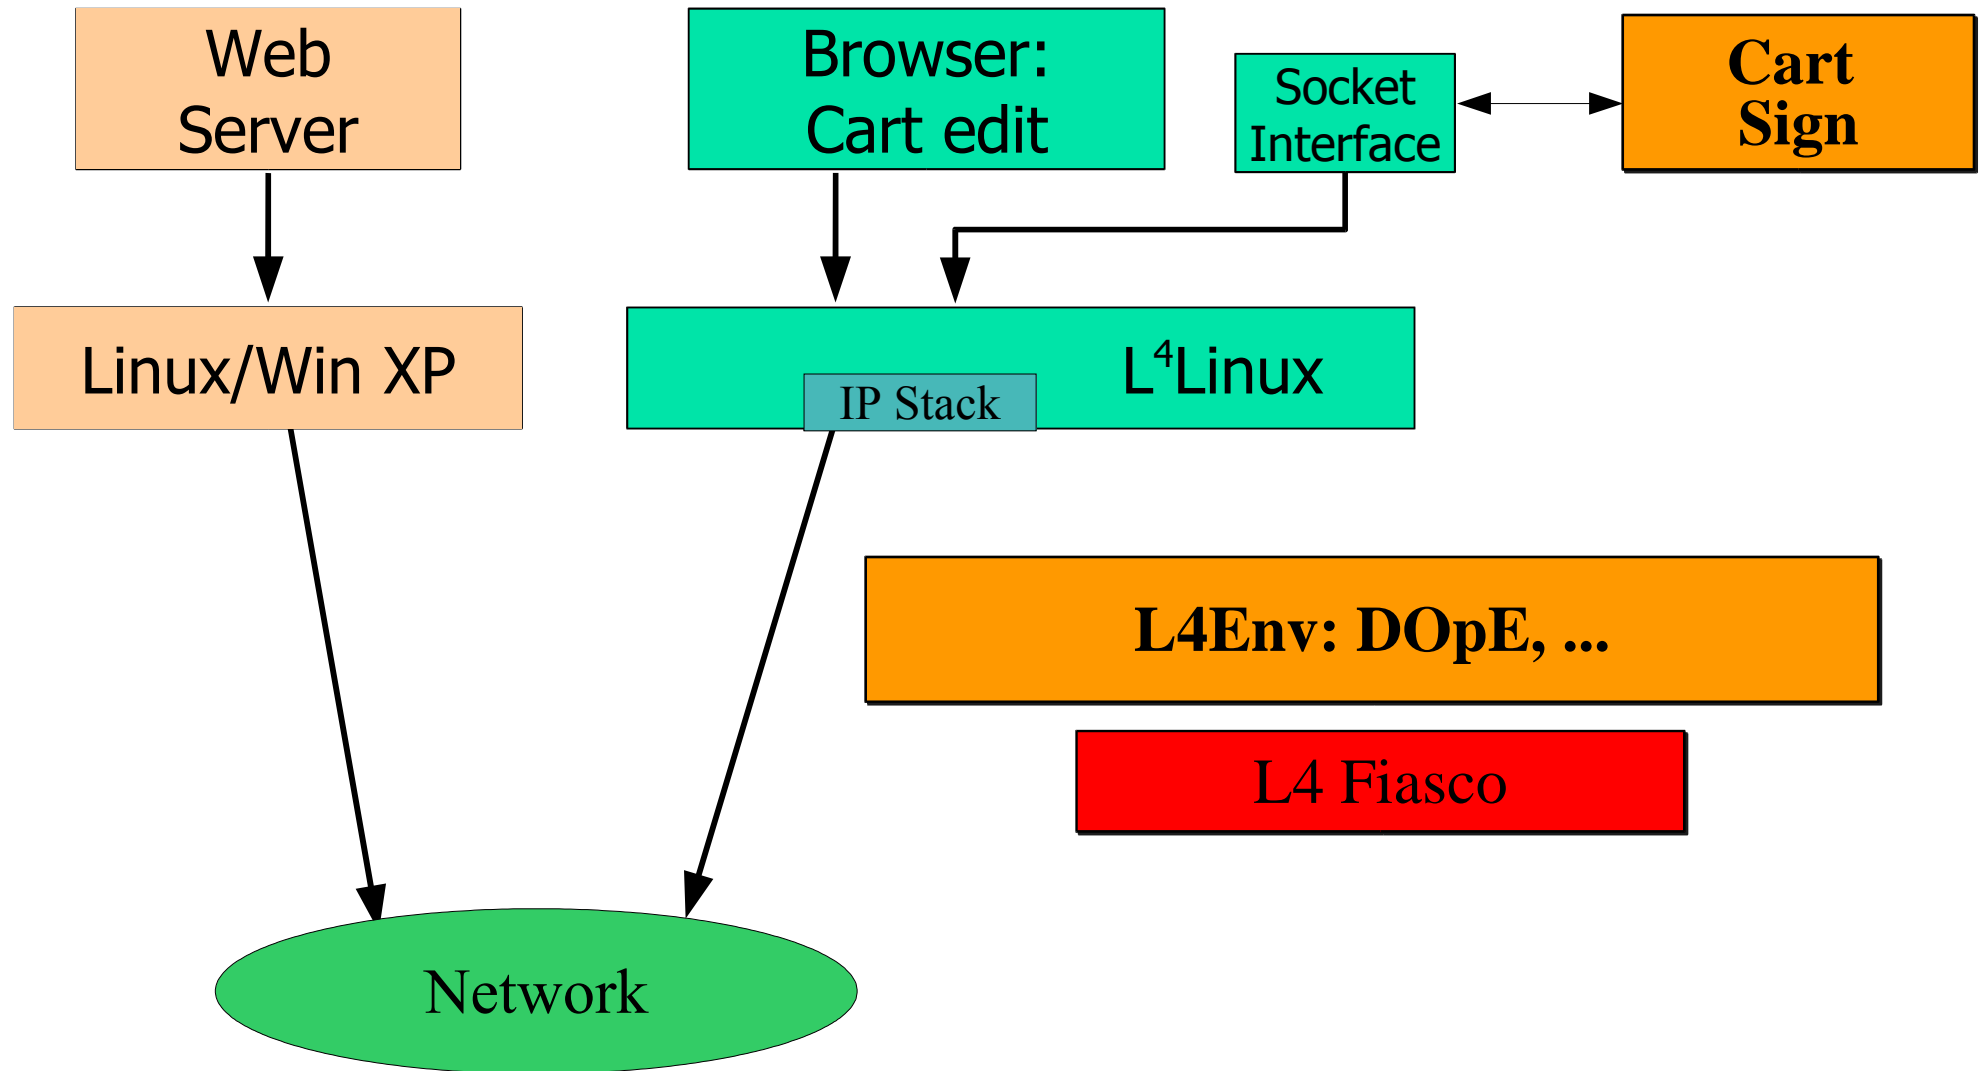
L⁴Linux

Minimal Trusted Platform

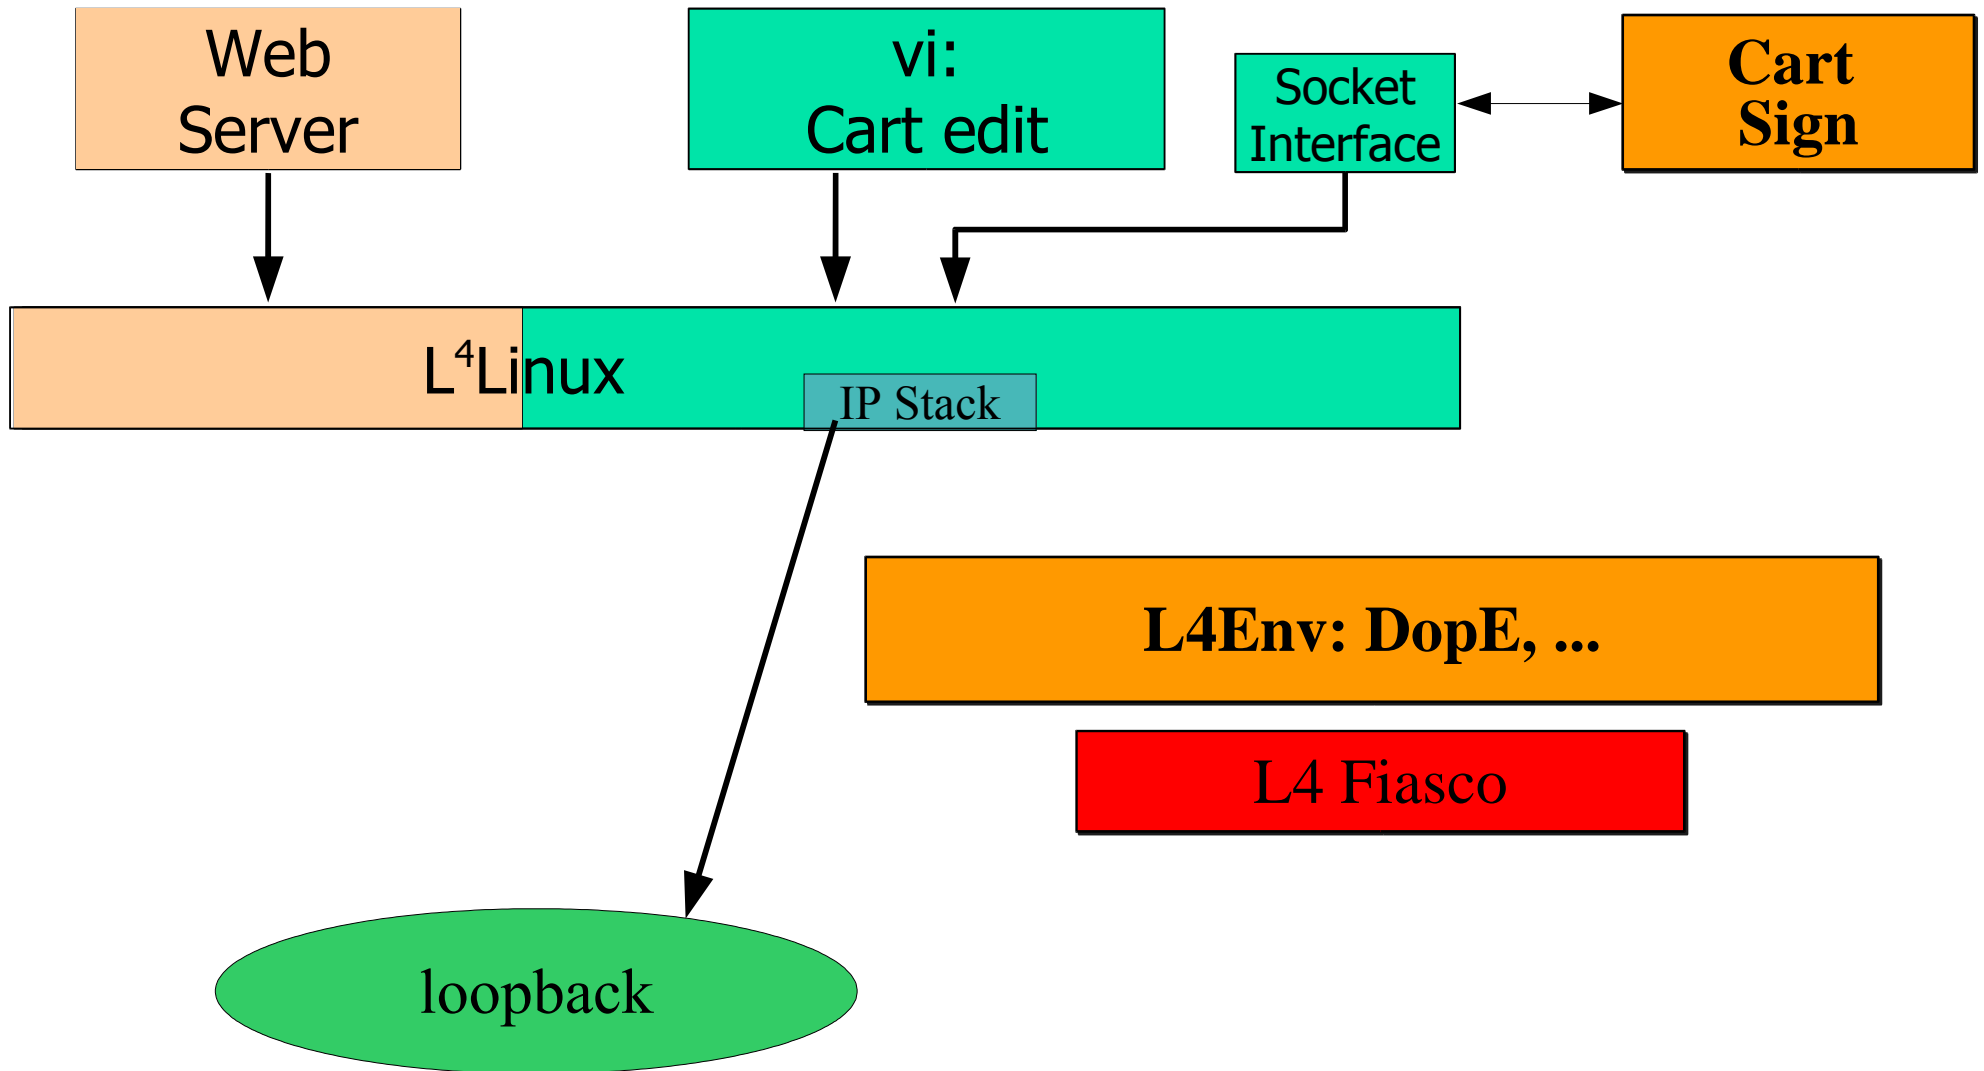
Network

password, credit
card number, keys, ...

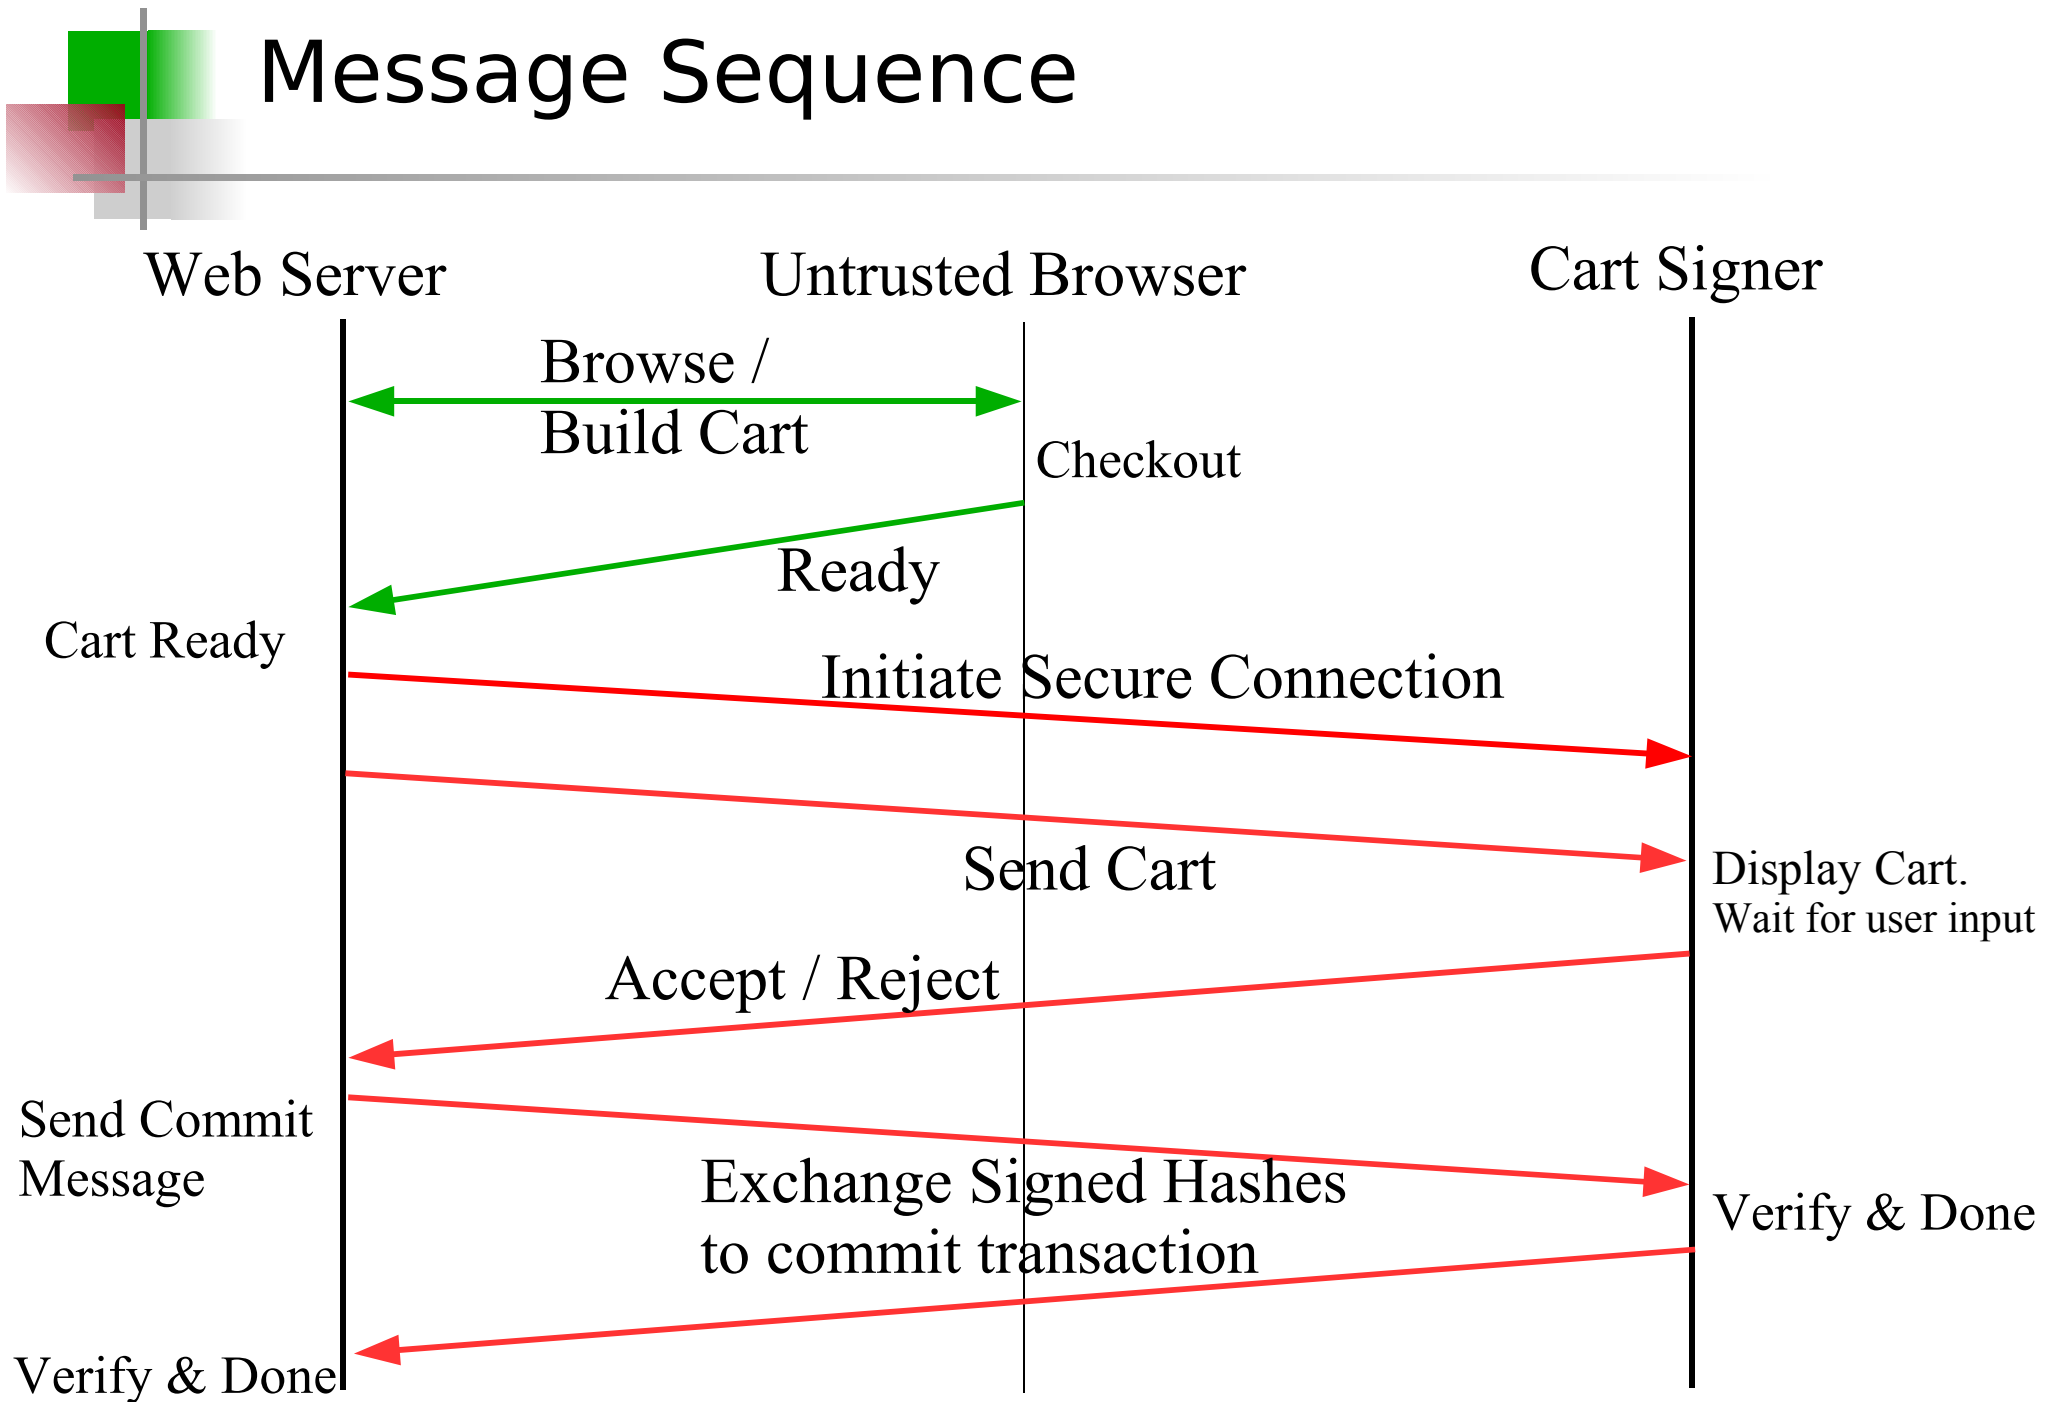
Lenin's Demo



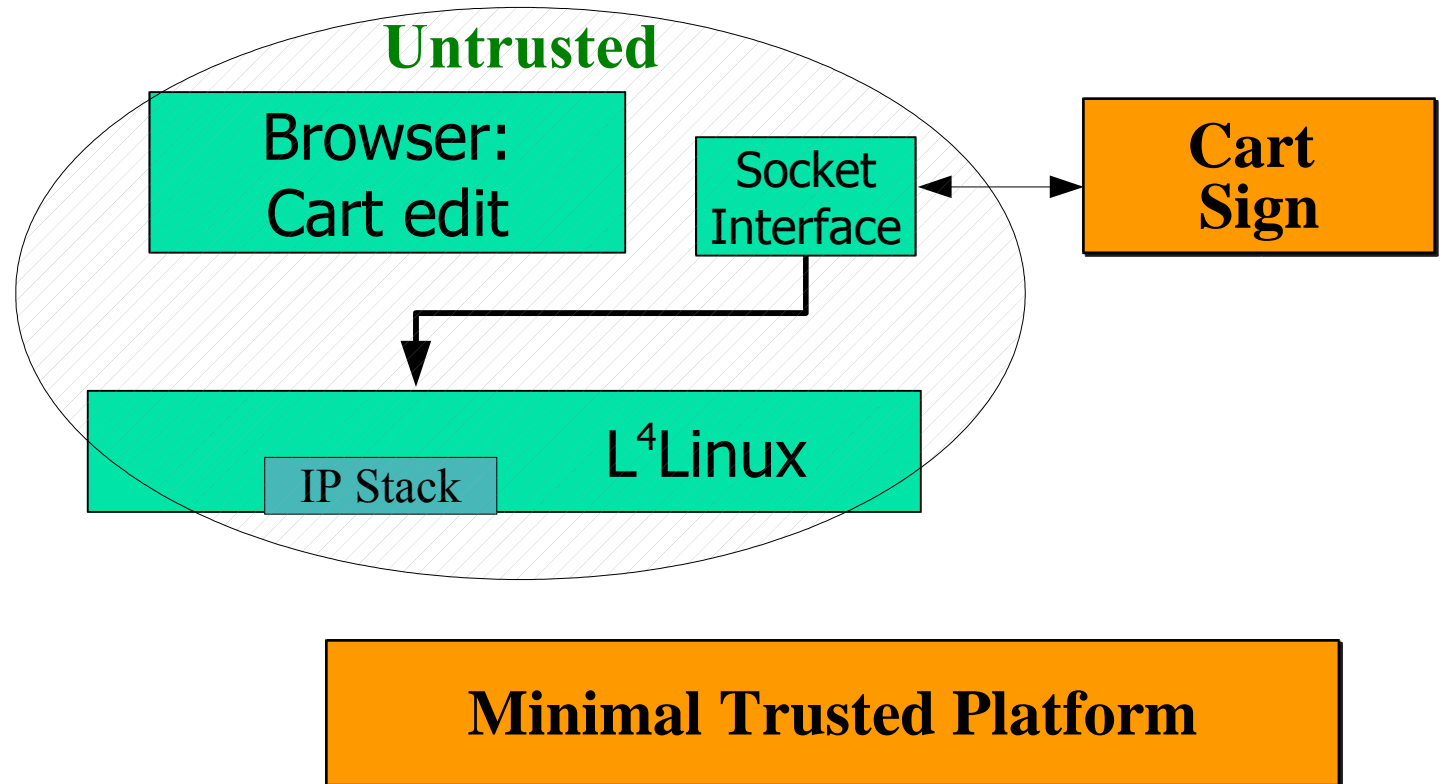
Lenin's Demo



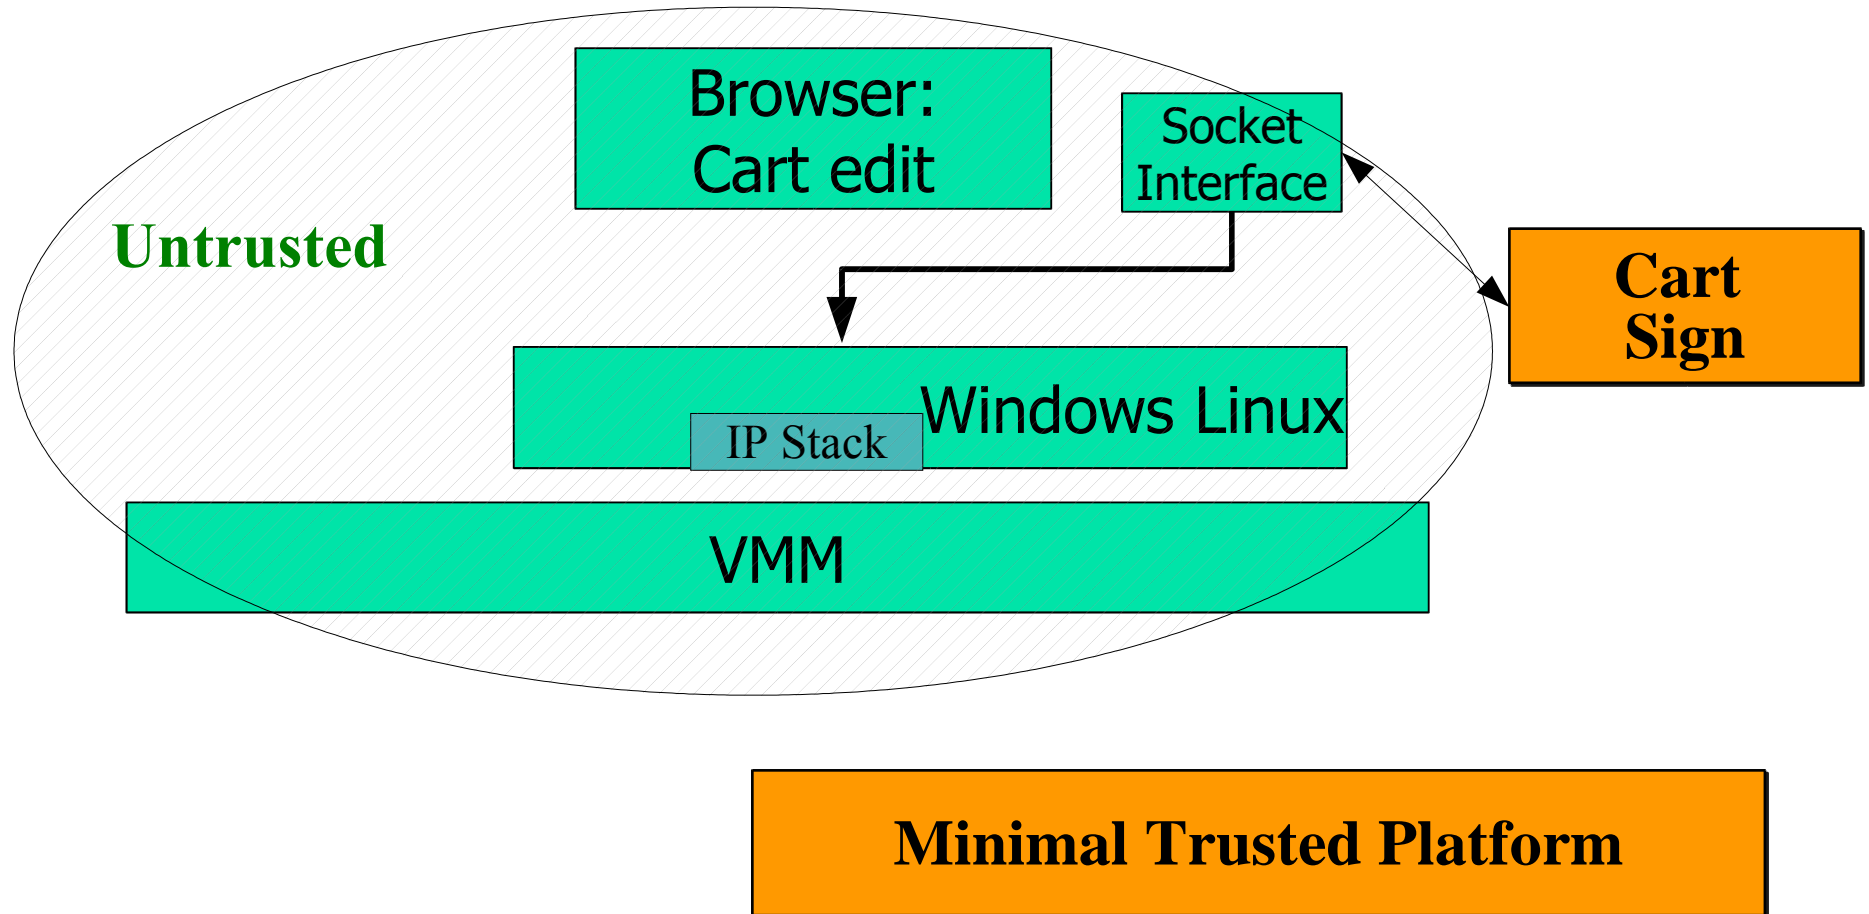
Message Sequence



Challenge: Untrusted VMM



Challenge: Untrusted VMM





Outline

- propaganda
- some names
- security and system security objectives
- design principles
- architecture and components
 - use cases
 - components: current and future
- Nizza vs. Virtual Machine Monitors and other related work
- technical risks



Names

- TUD♦OS Technische Universität Dresden OS
 - DROPS Dresden Real-Time OS
 - Nizza Security Architecture
 - Micro-Sina A Nizza Application (use case)
 - L⁴Env set of servers and libraries
 - DOpE Window manager for DROPS and Nizza
 - L⁴Linux Linux kernel as user-level server
- L4 a micro-kernel interface
- L4/Fiasco, L4/Pistacchio: L4 implementations



Objectives: Security

- confidentiality
no unauthorized access to information
- integrity
no unauthorized, **unnoticed** modification of information
- recoverability
no permanent damage to information
- availability
timeliness of service



Objectives: System Security

- Secure and unsecure applications



Objectives: System Security

- Secure and unsecure applications
- Compatibility:
 - Legacy applications
 - Legacy/Fashionable Hardware



Objectives: System Security

- Secure and unsecure applications
- Compatibility
- Flexible sandboxing



Objectives: System Security

- Secure and unsecure applications
- Compatibility
- Flexible sandboxing
- Resource Control



Objectives: System Security

- Secure and unsecure applications
- Compatibility
- Flexible sandboxing
- Resource Control
- Small TCB:
complexity acceptable, if for a small group
 - Each member fully understands interaction of all components.
 - Each component is understood by one member.



Principles

- Trusted Computing Base: per application
- platform:
 - small set of small components (servers, ...)
 - small interfaces
 - select components of platform per application



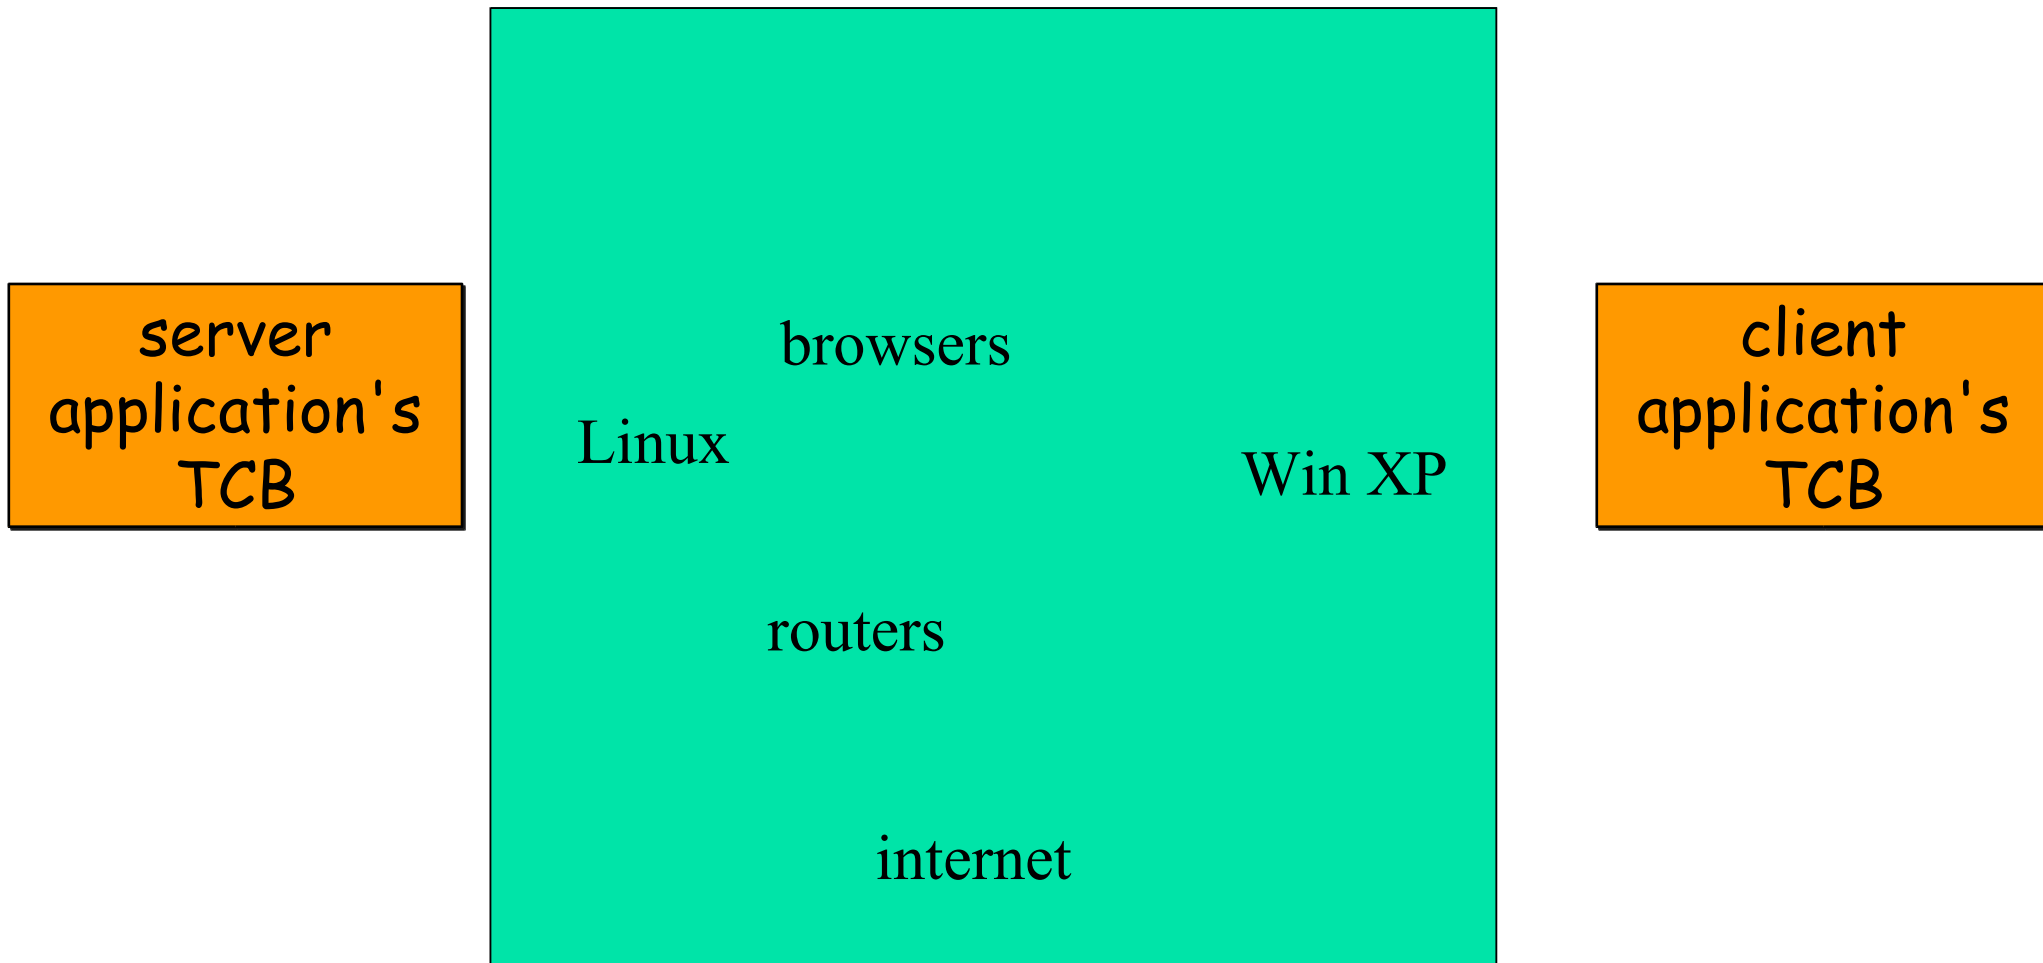
Principles, continued

- split applications and services:
sensitive part in/on trusted platform
and *other* part
- reuse legacy for *other* part
→ trusted wrappers / tunneling



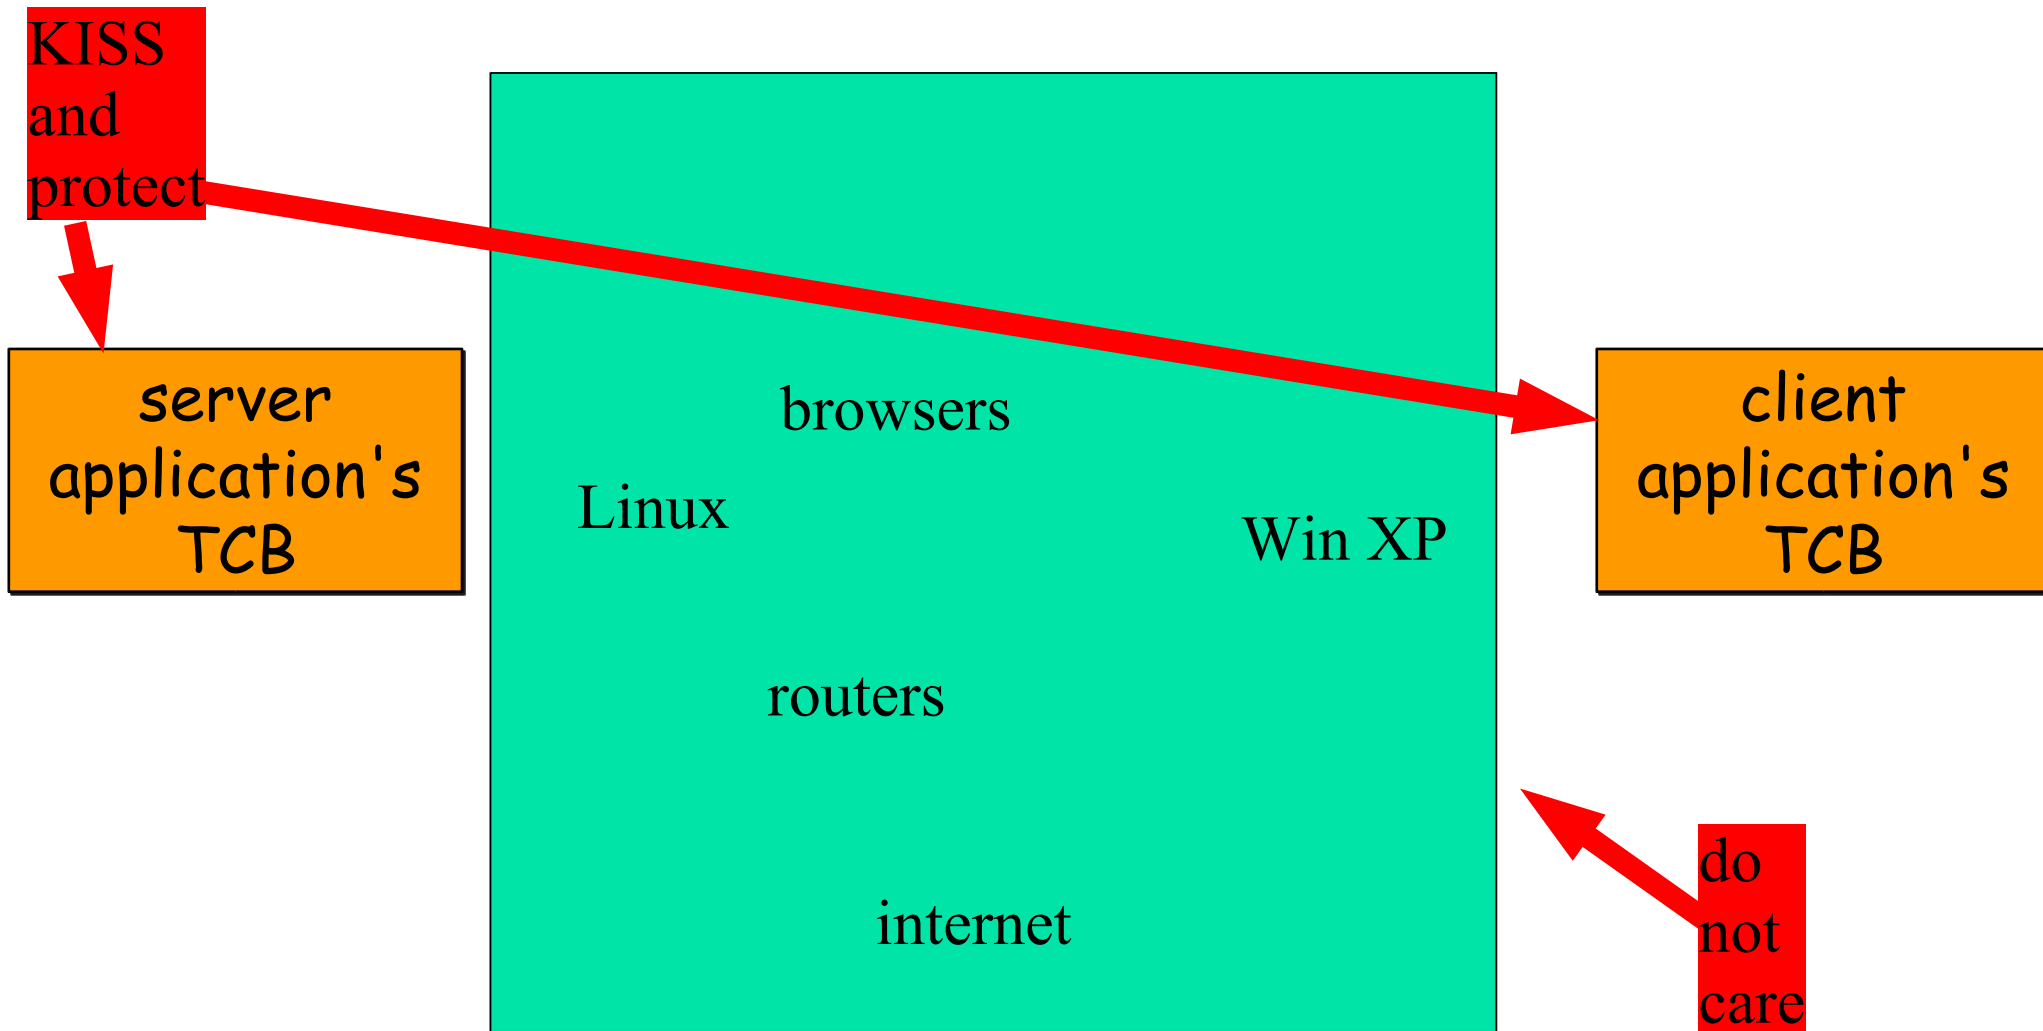
Principles

- push end-to-end argument to the extreme



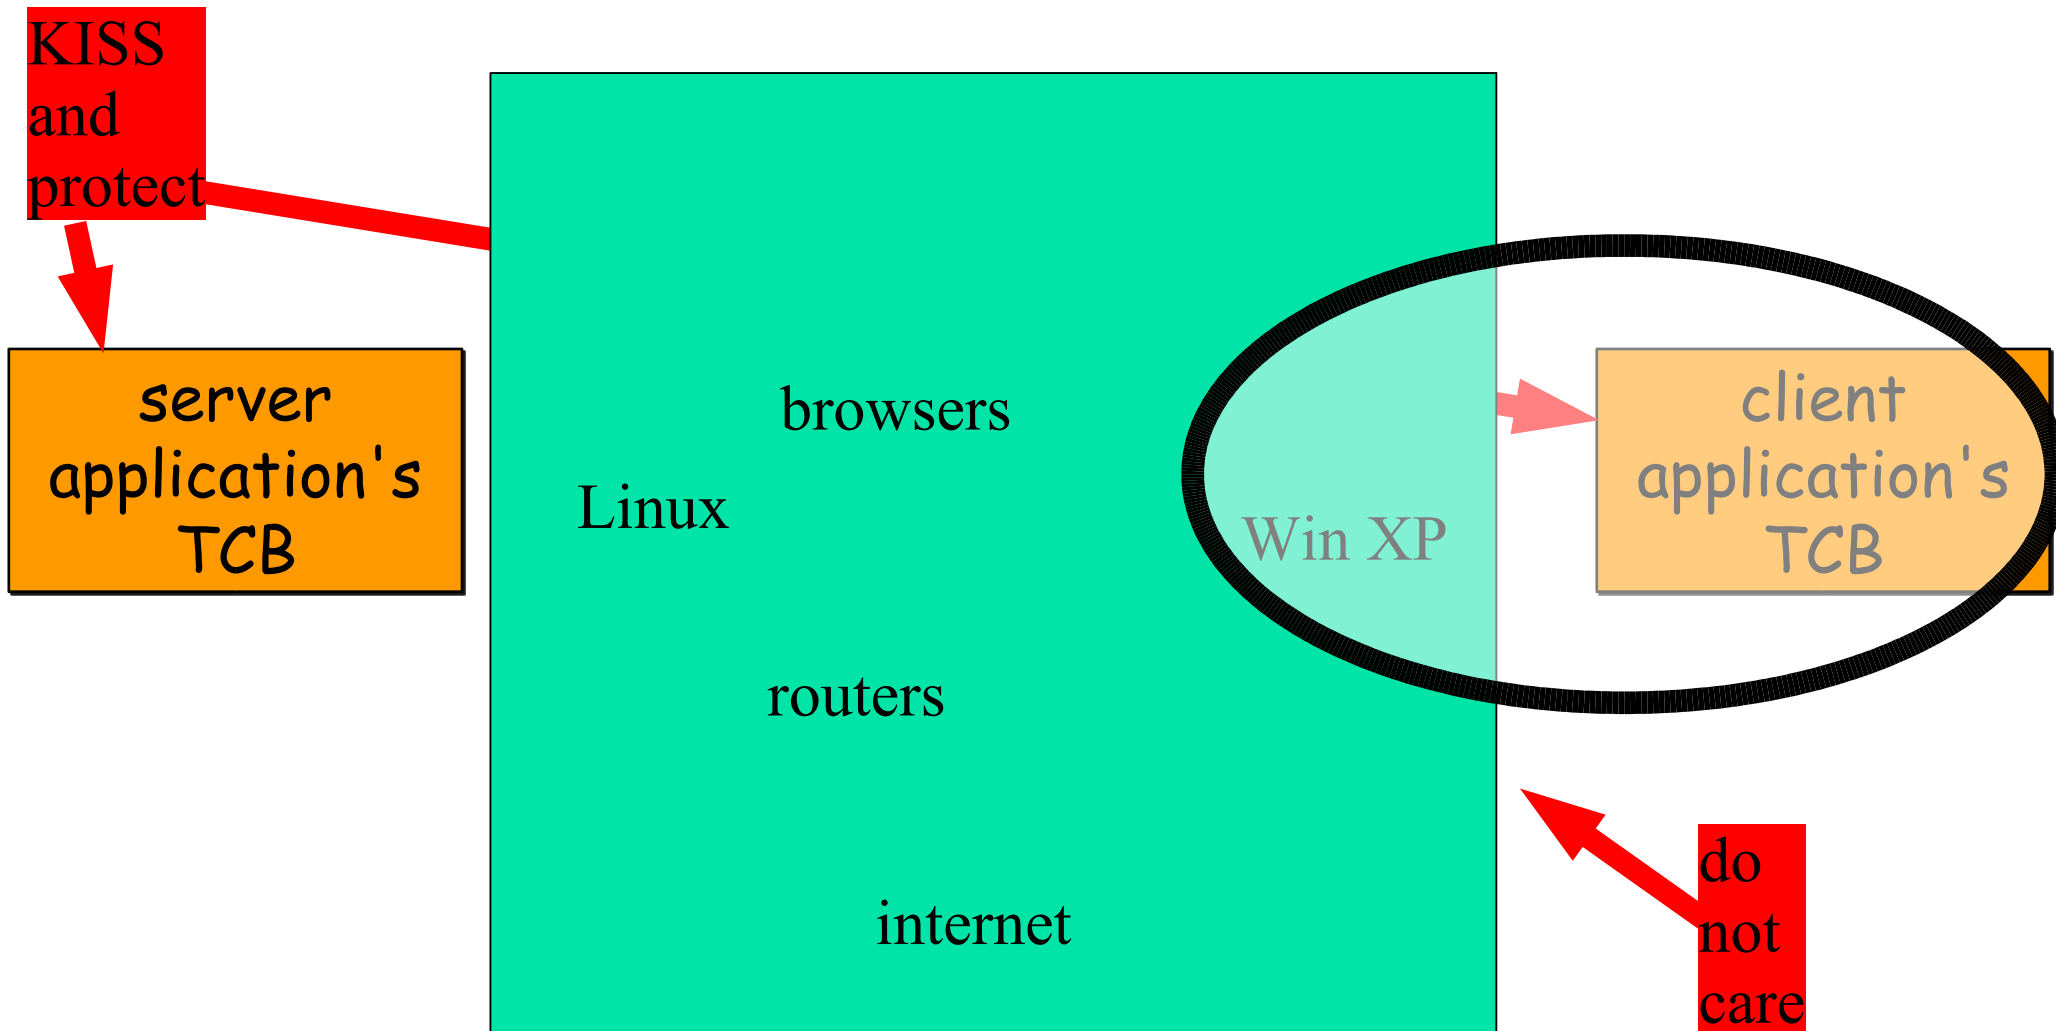
Principles

- push end-to-end argument to the extreme



Principles

- push end-to-end argument to the extreme





Principles and Techniques

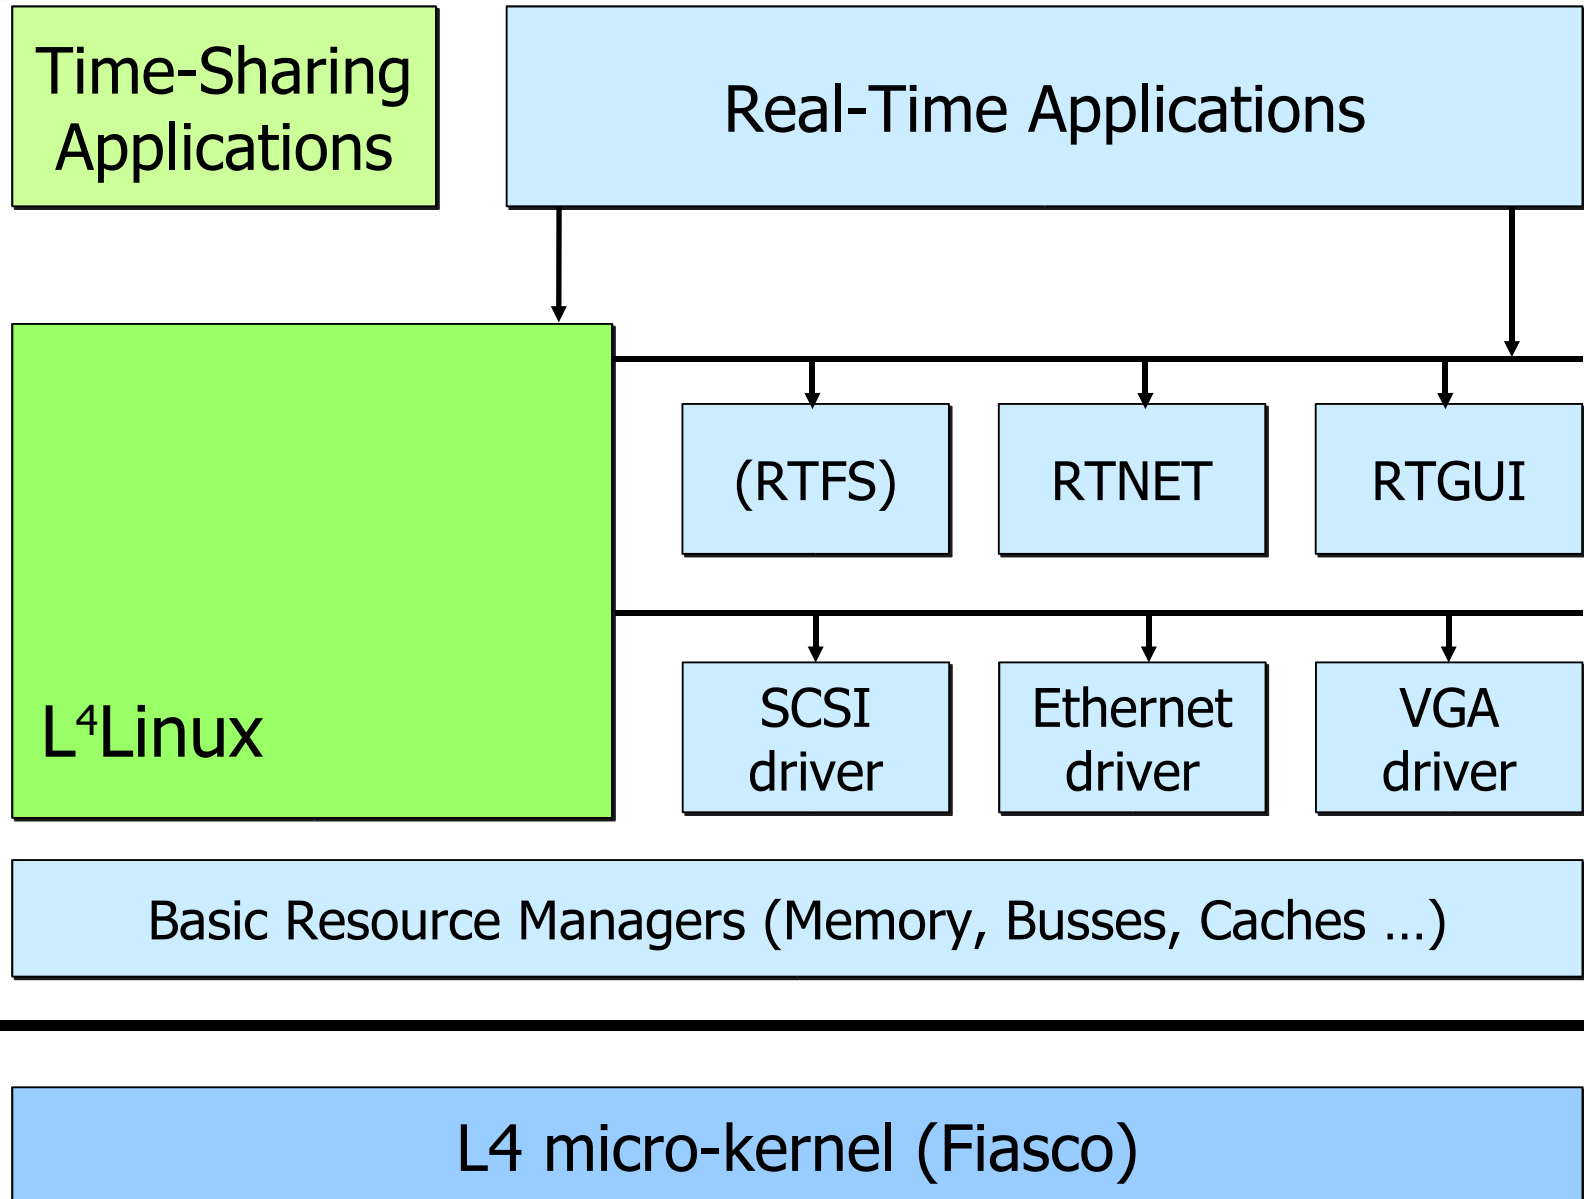
- micro kernel:
 - separates legacy from sensitive partitions
 - separates components of small platform
 - provides mechanisms for access control
mediates communication
- contract-based access control
- secure booting / attestation +
trusted path to user

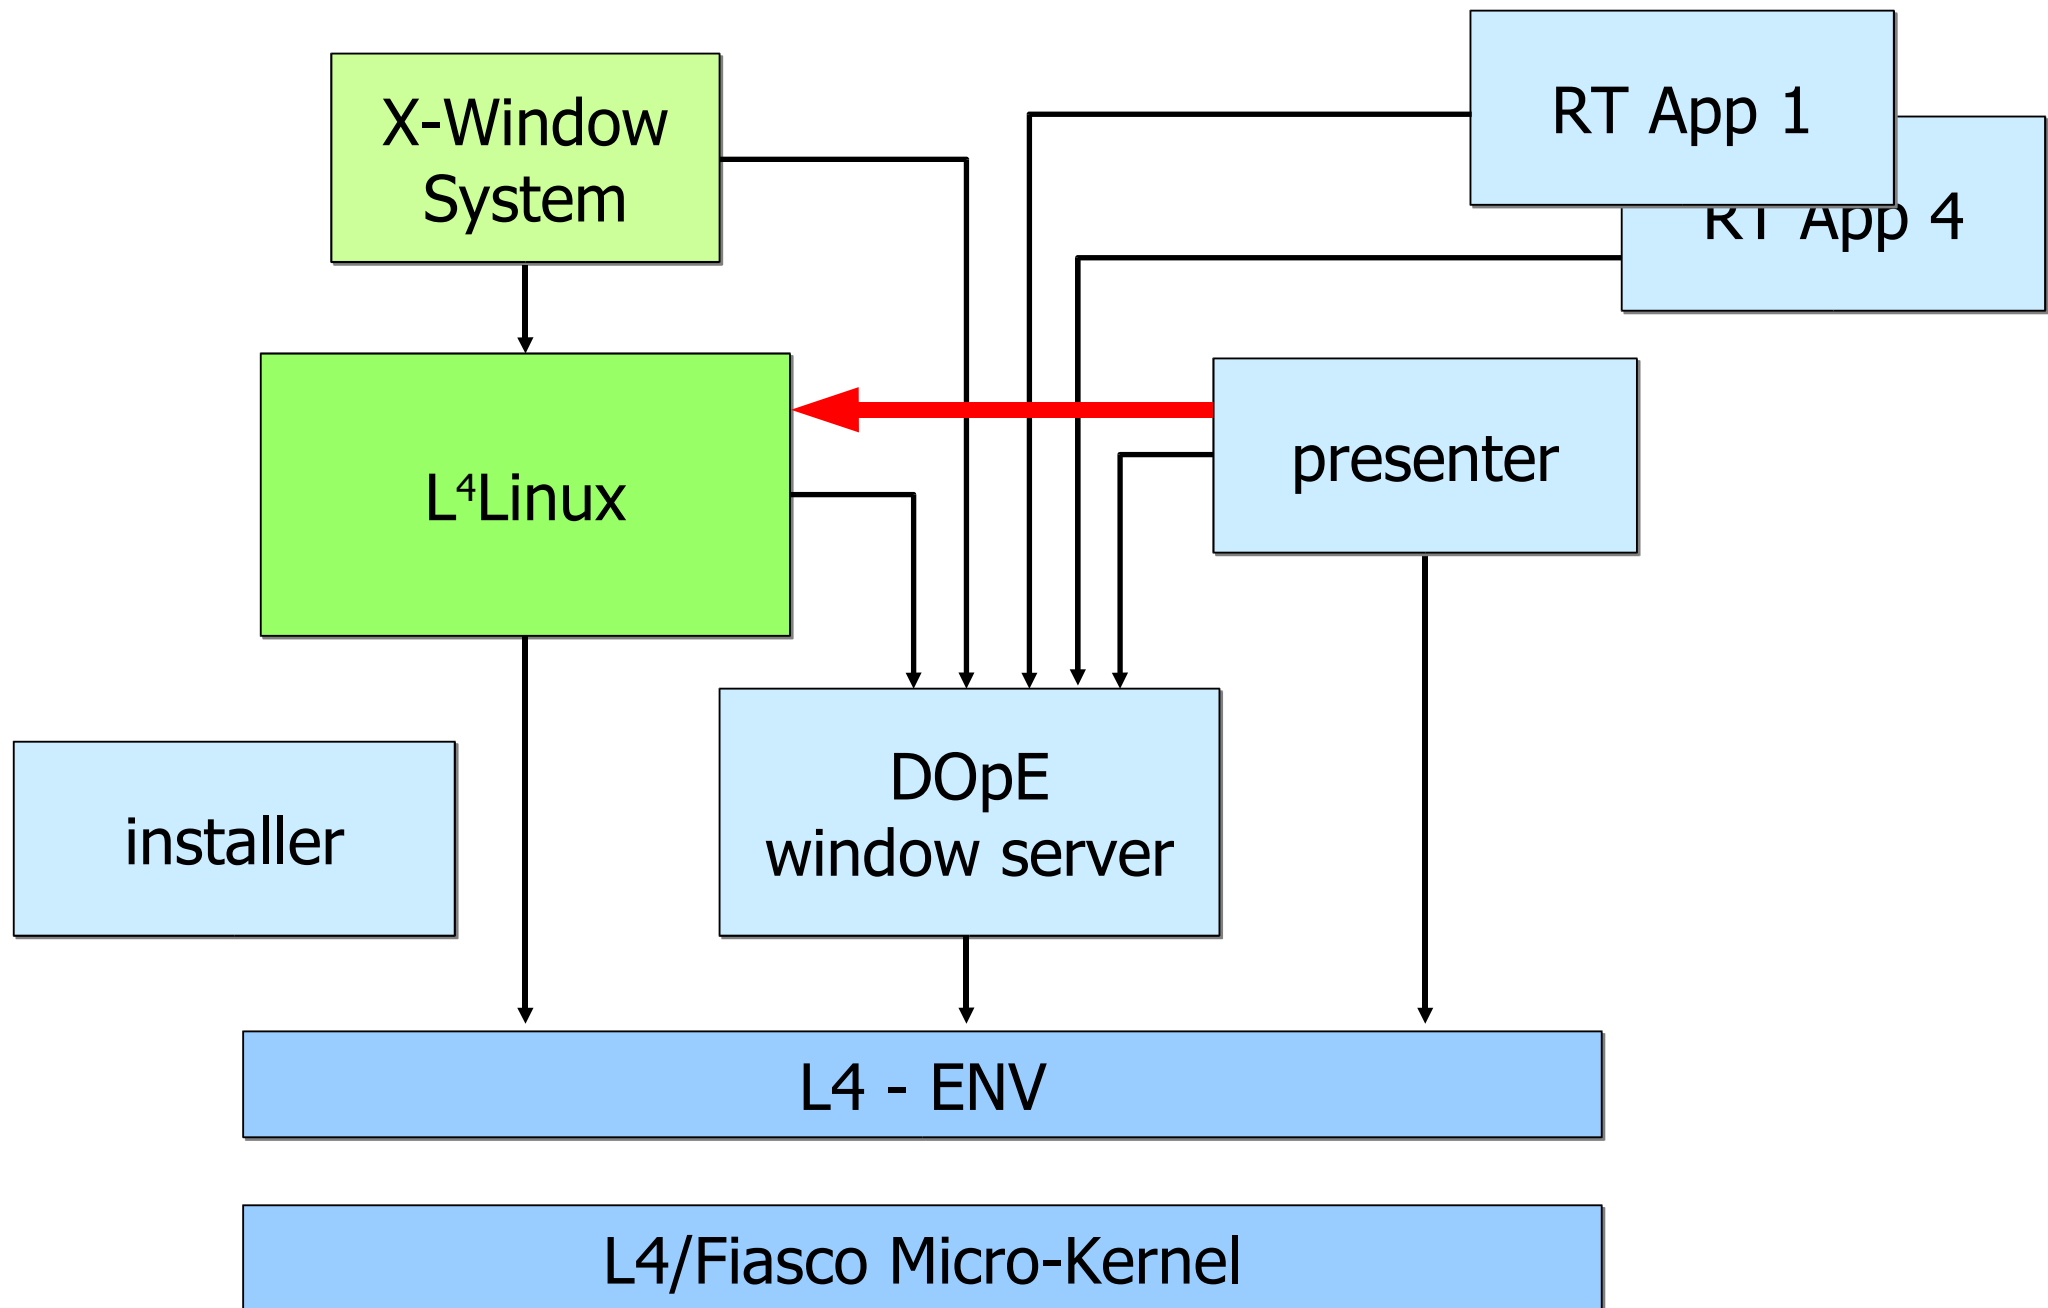


Architecture + Components

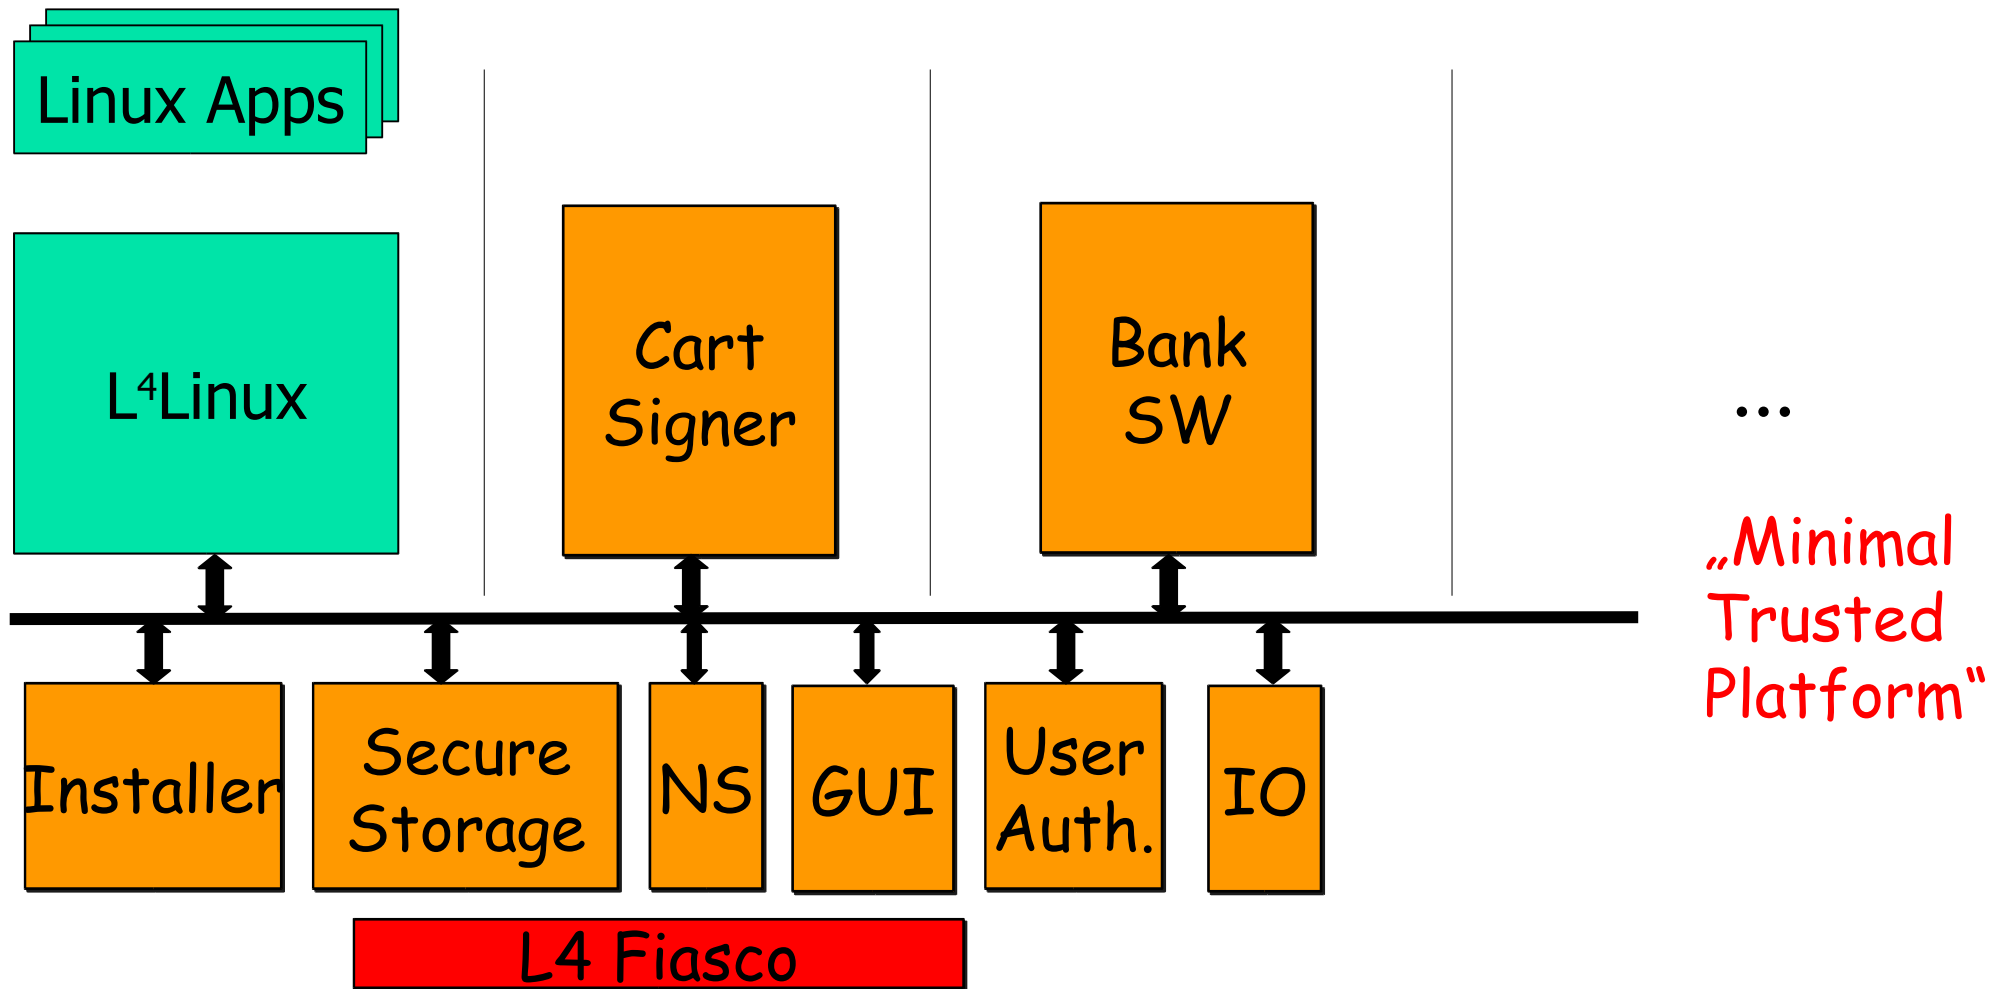
- the origin: Dresden Real-time OPerating System
- Nizza architecture
- use cases
- some Nizza components

Starting Point: DResden Real-Time OS

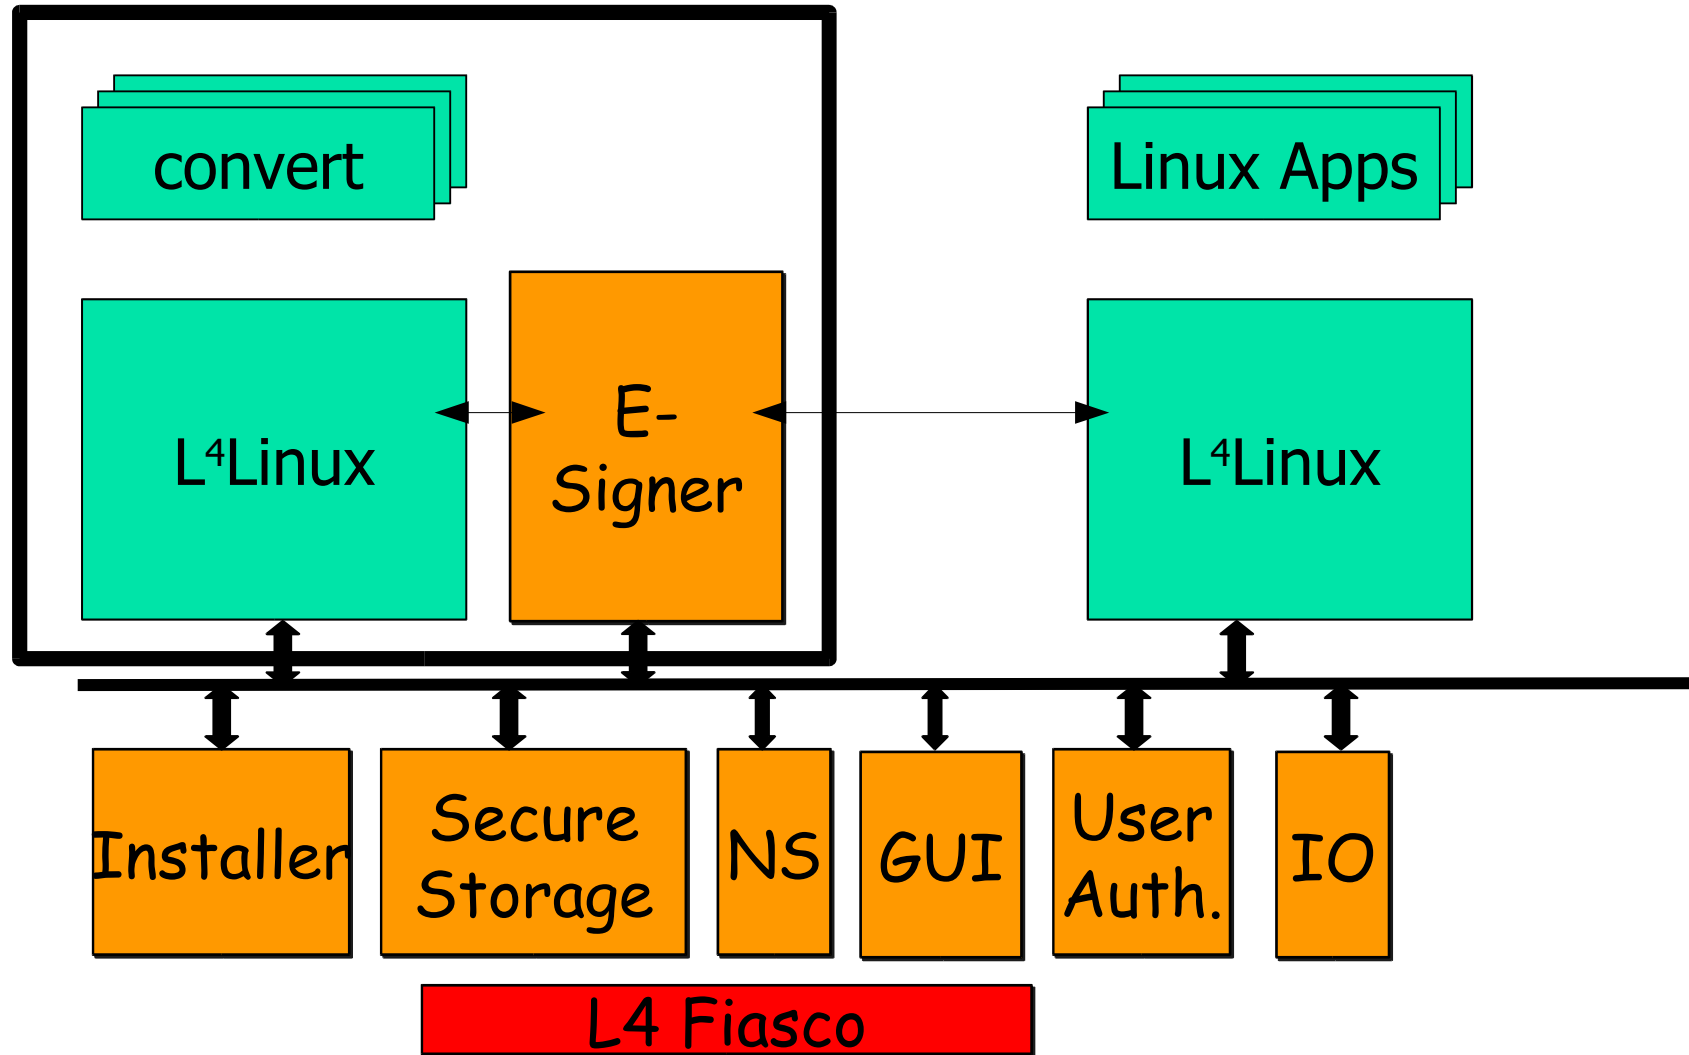




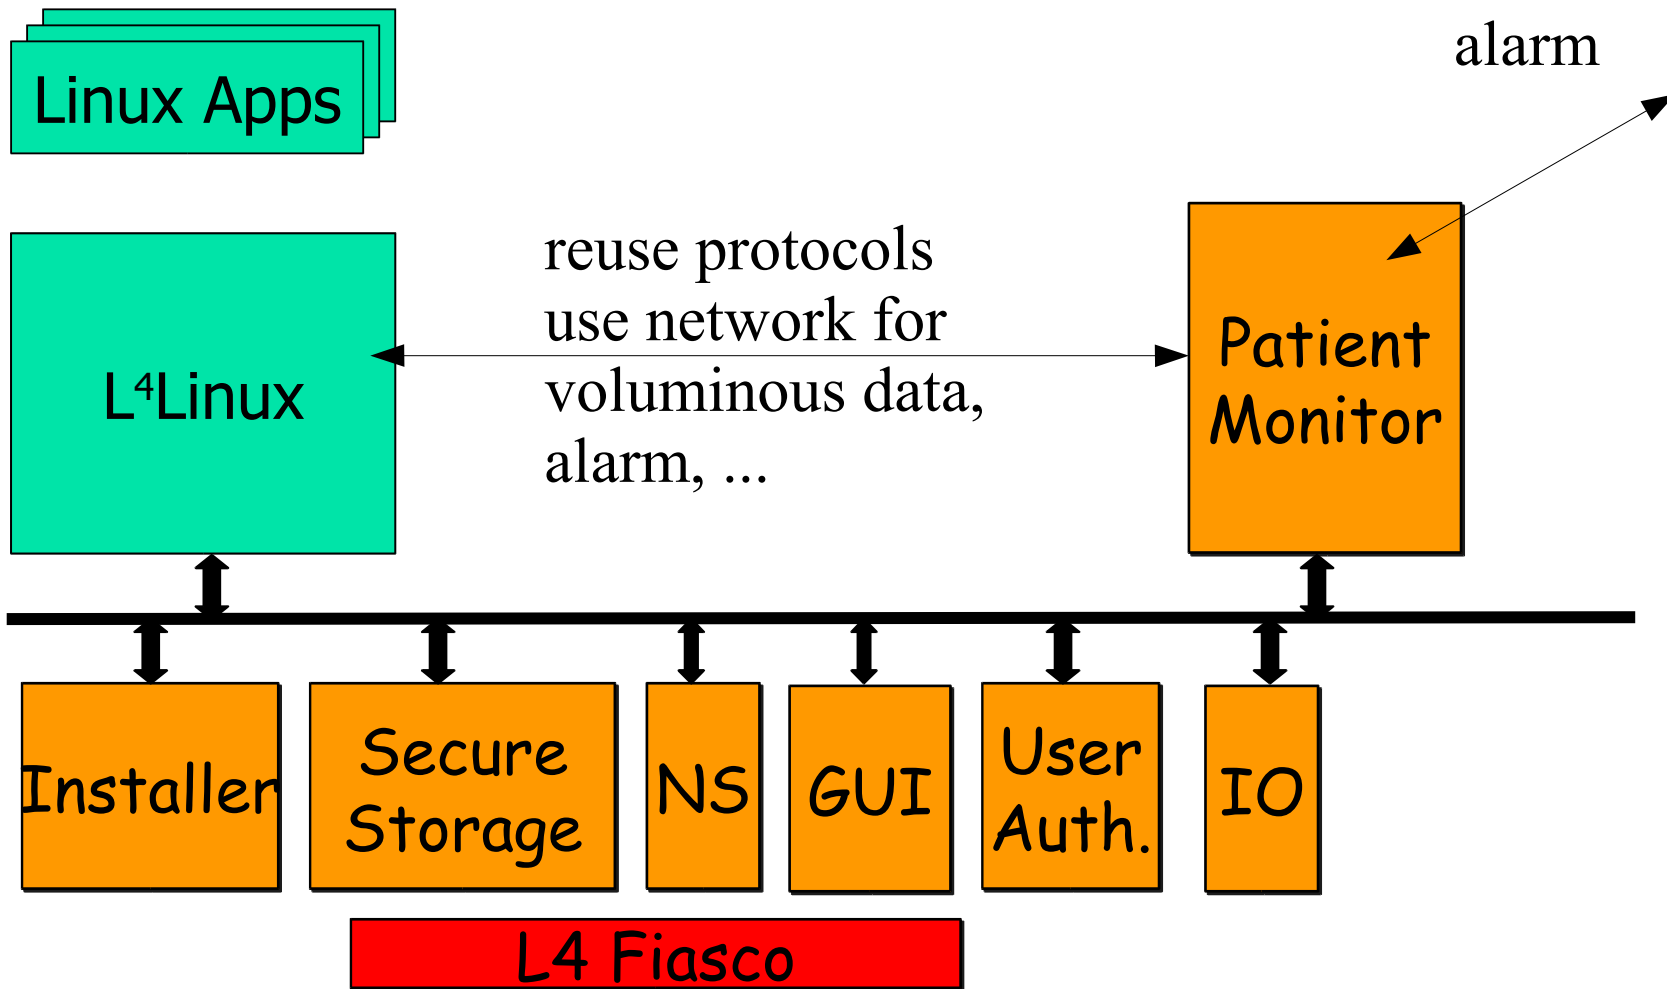
NIZZA Architecture



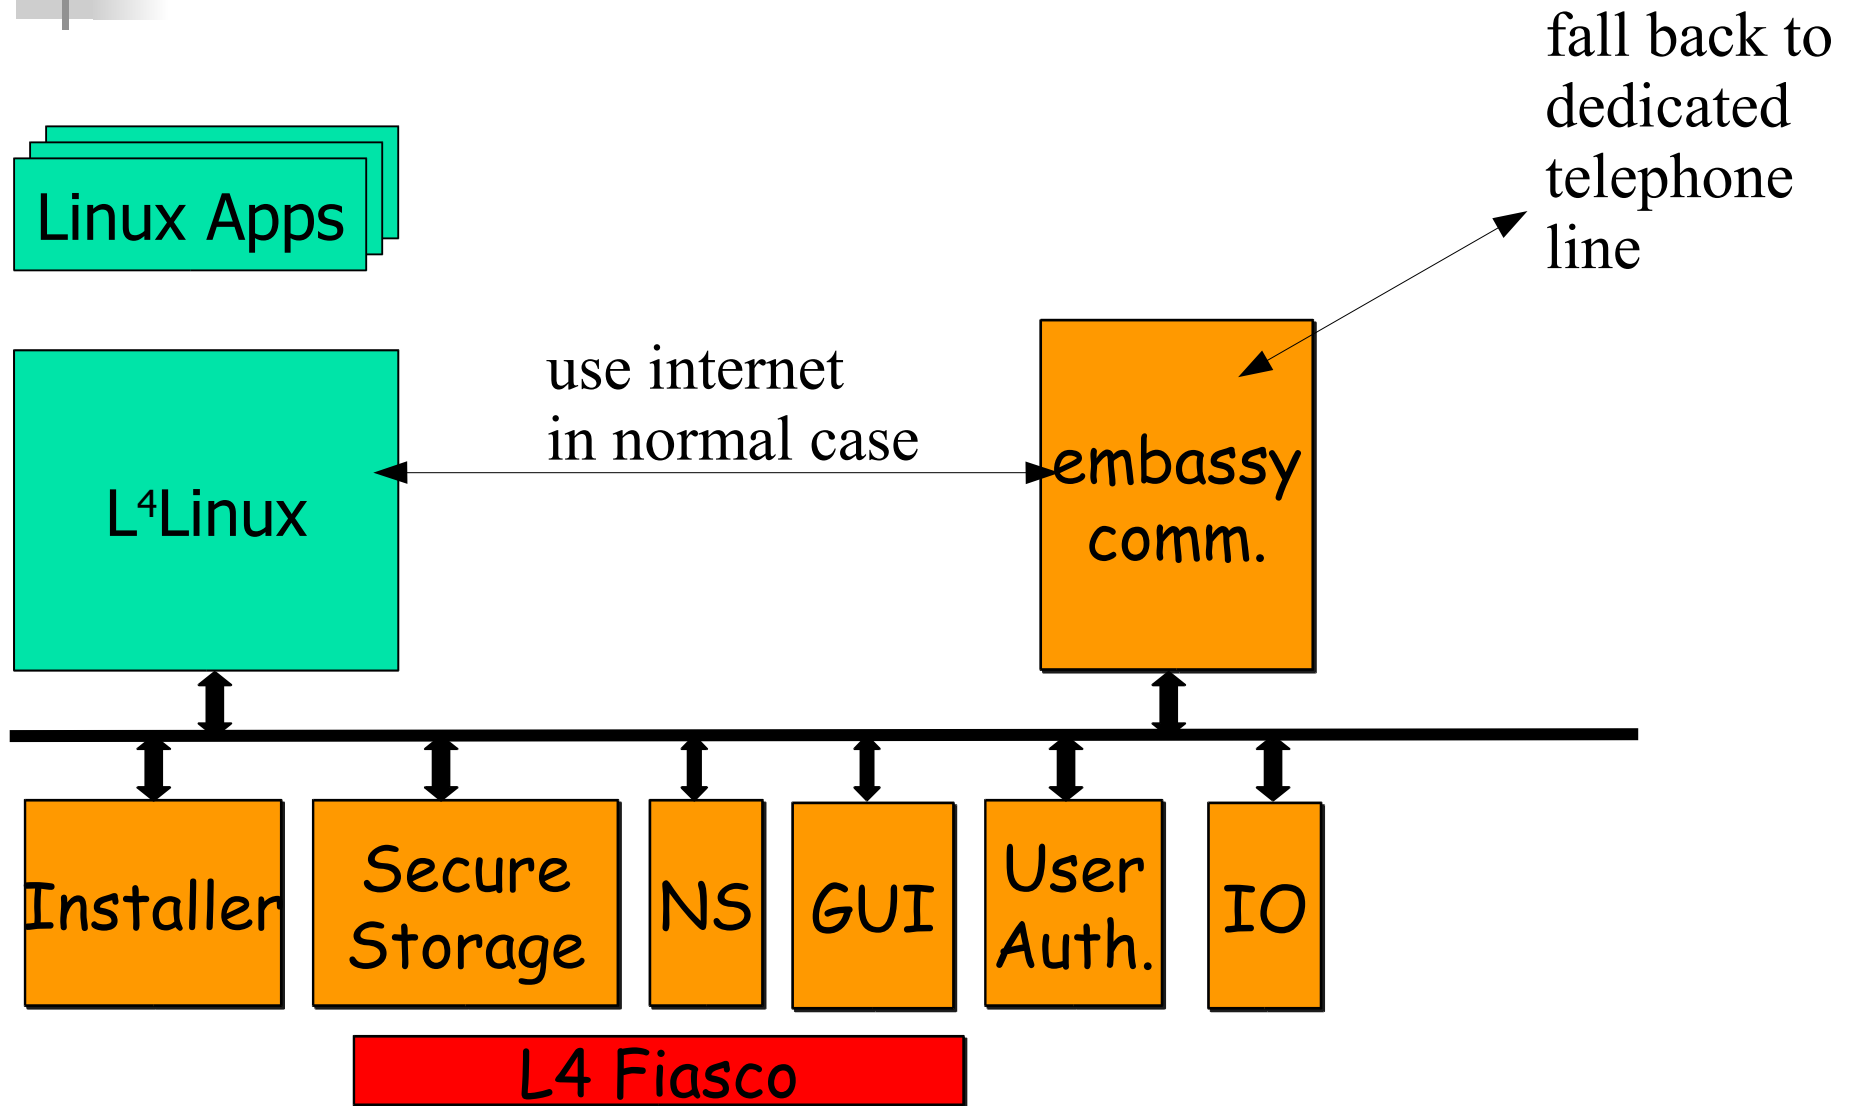
More Use Cases



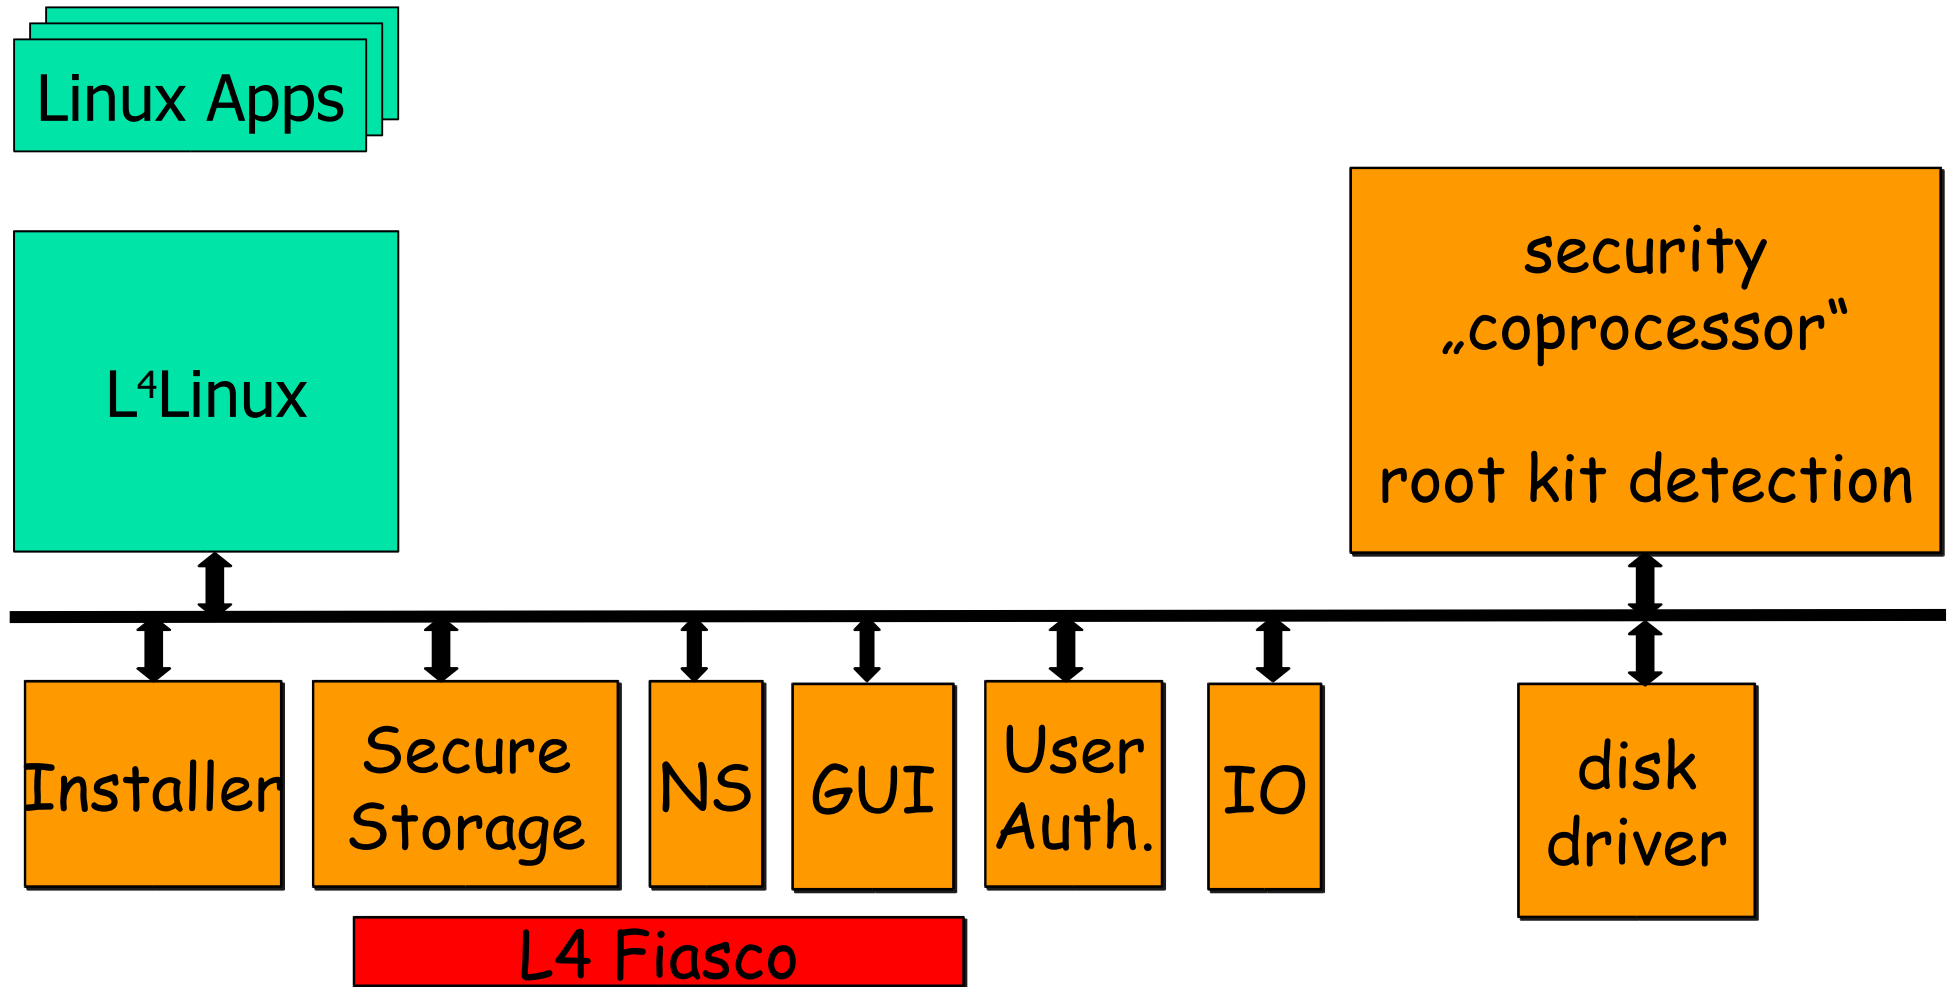
More Use Cases




More Use Cases



More Use Cases



More Use Cases





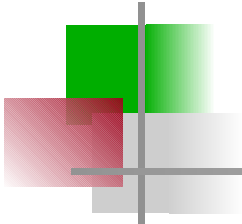
Architecture + Components

- the origin: Dresden Real-time OPerating System
- Nizza architecture
- use cases
- **some Nizza components**
 - L4 micro kernel: present and future
 - L⁴Linux: encapsulation and reuse
 - secure booting + trusted path
 - secure storage with small TCB (future)



L4 Micro-Kernel: evolution

- the original:
address spaces, threads, IPC
- L4/Fiasco-RT: real-time:
periodic threads, fine grained scheduling
support, etc
- **L4/Fiasco-X.e:**
unified access control +
kernel resource management
- L4-Next: virtualization support




L4 micro kernel (Jochen Liedtke)

fundamental abstractions


- address spaces (separation)
- threads
- inter process communication (IPC)
 - explicit
 - interrupts
 - faults and mappings

Drivers at User Level



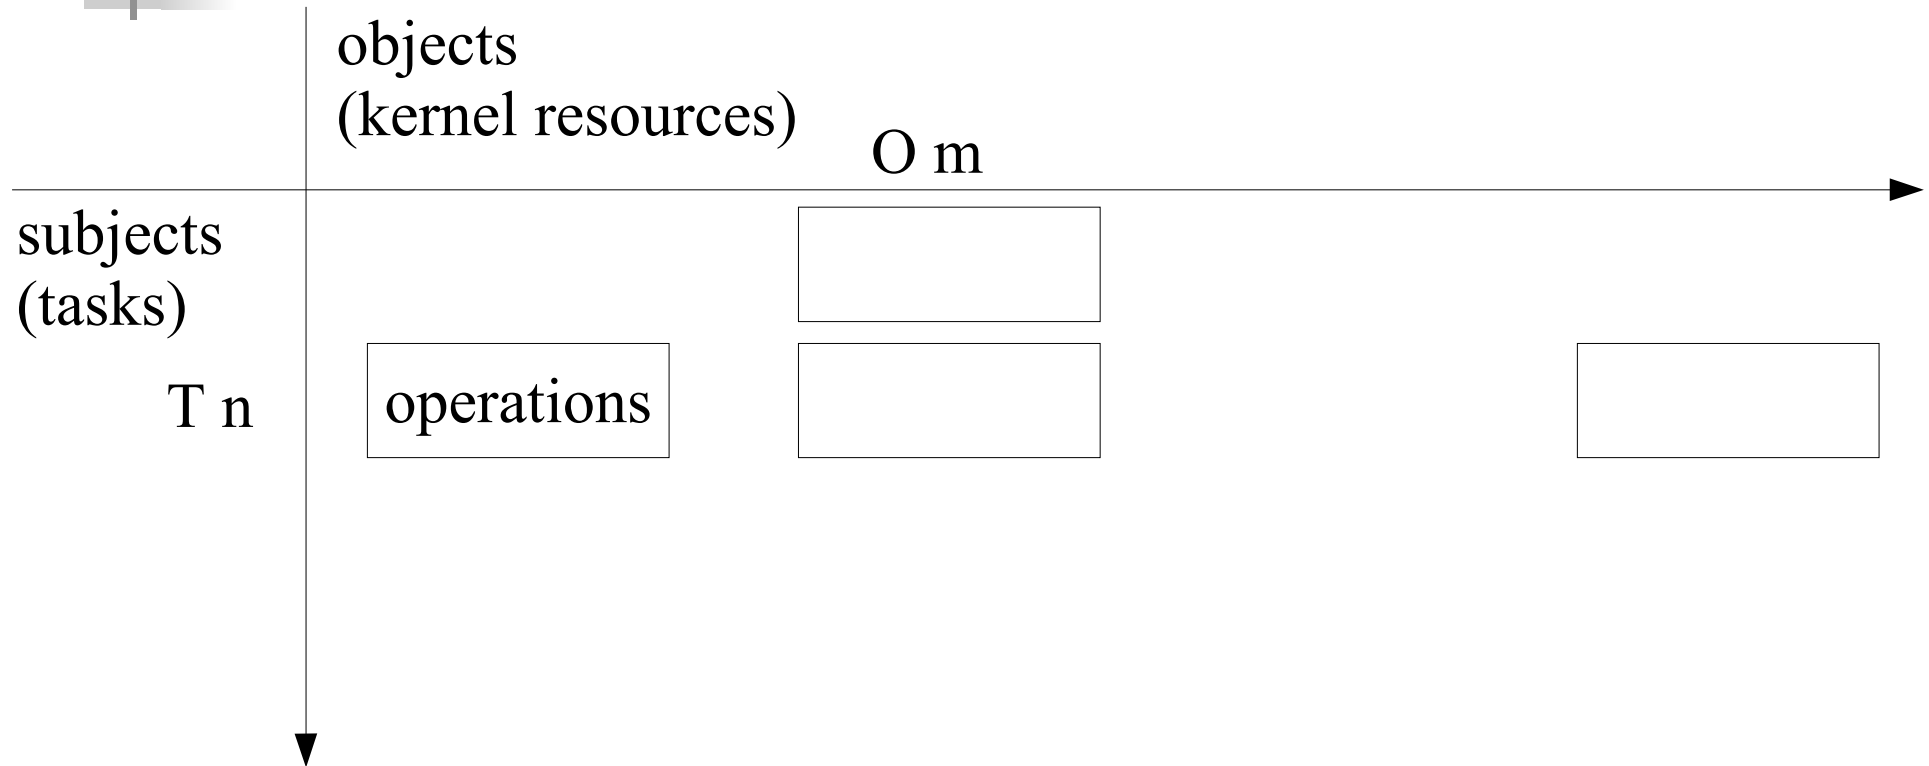
- **IO ports:** part of the user address space
- **interrupts:** messages from hardware

Memory Pages

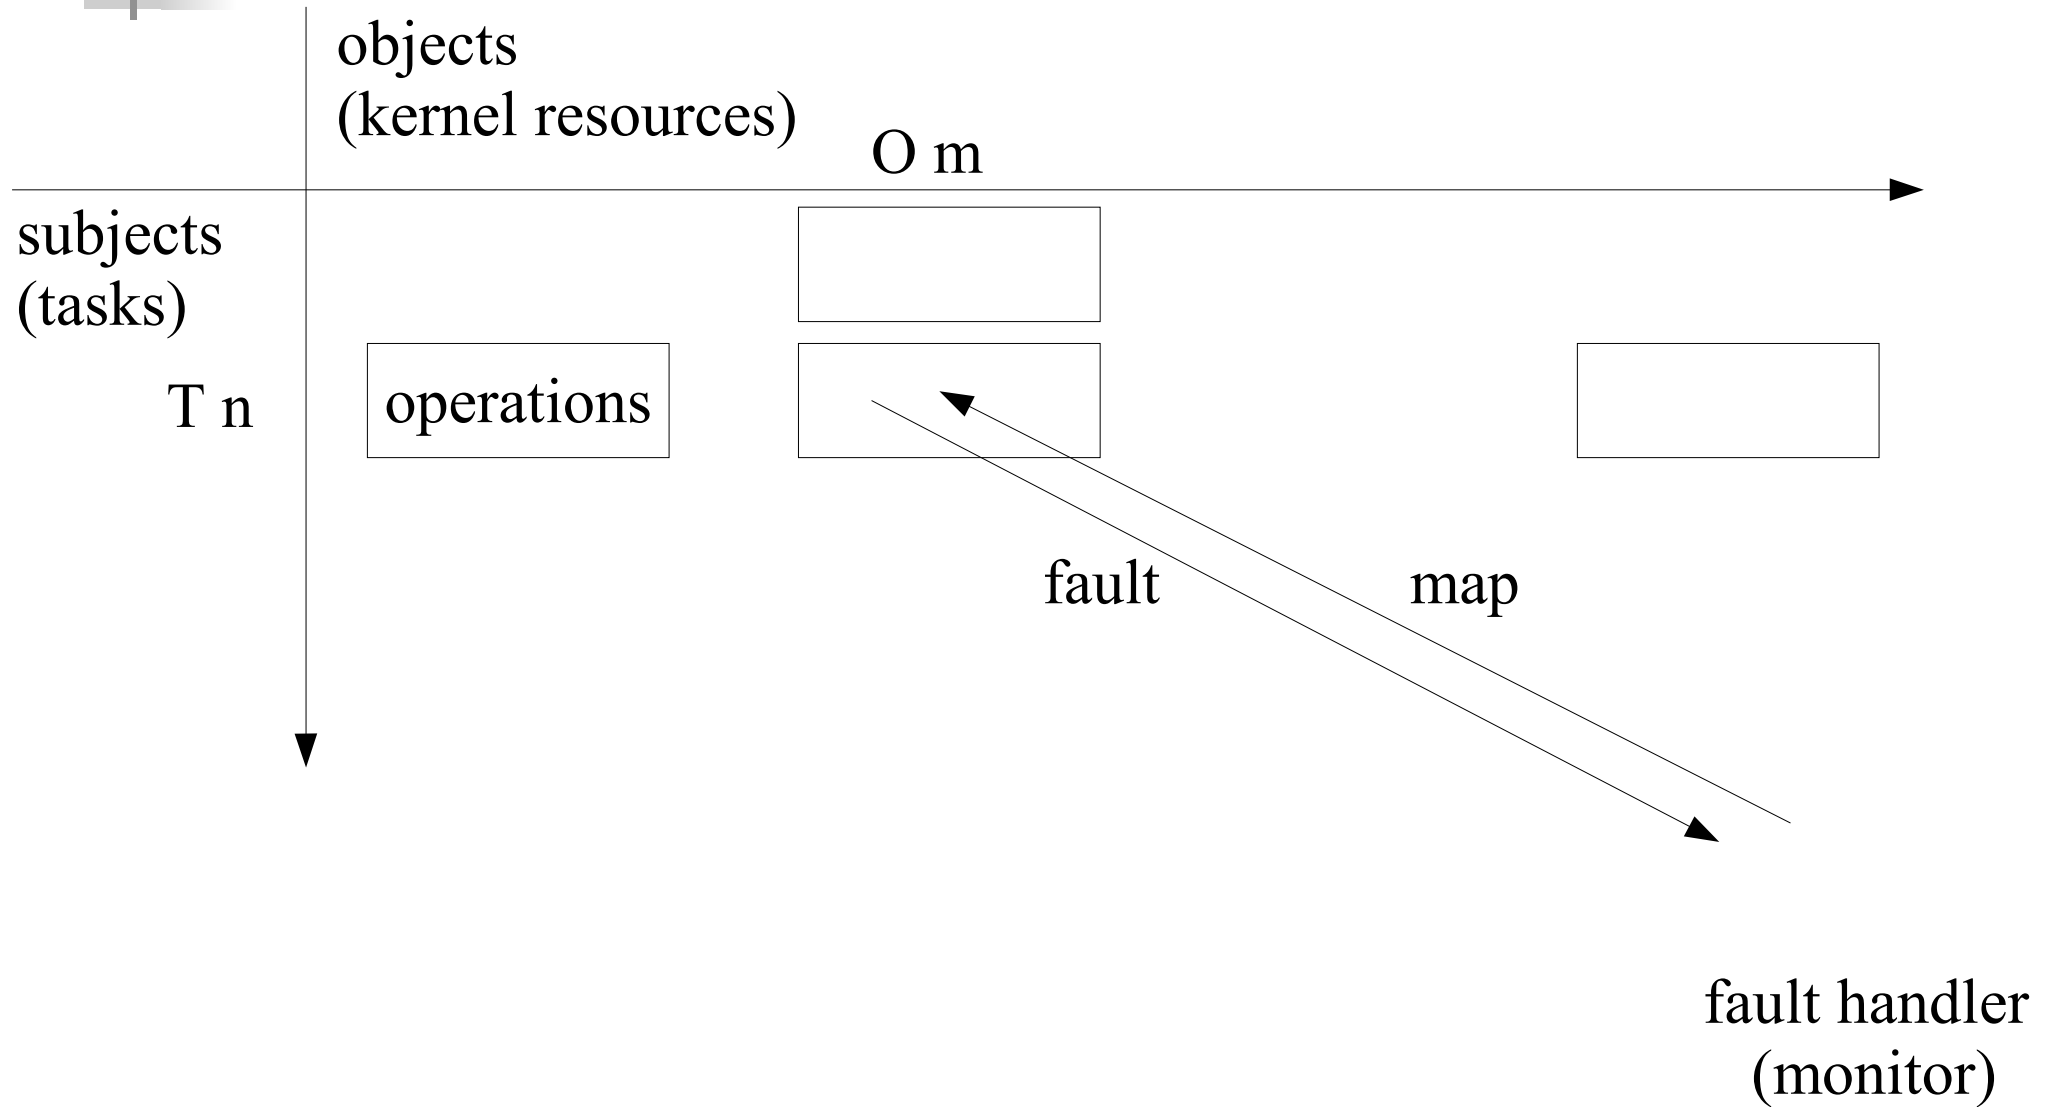




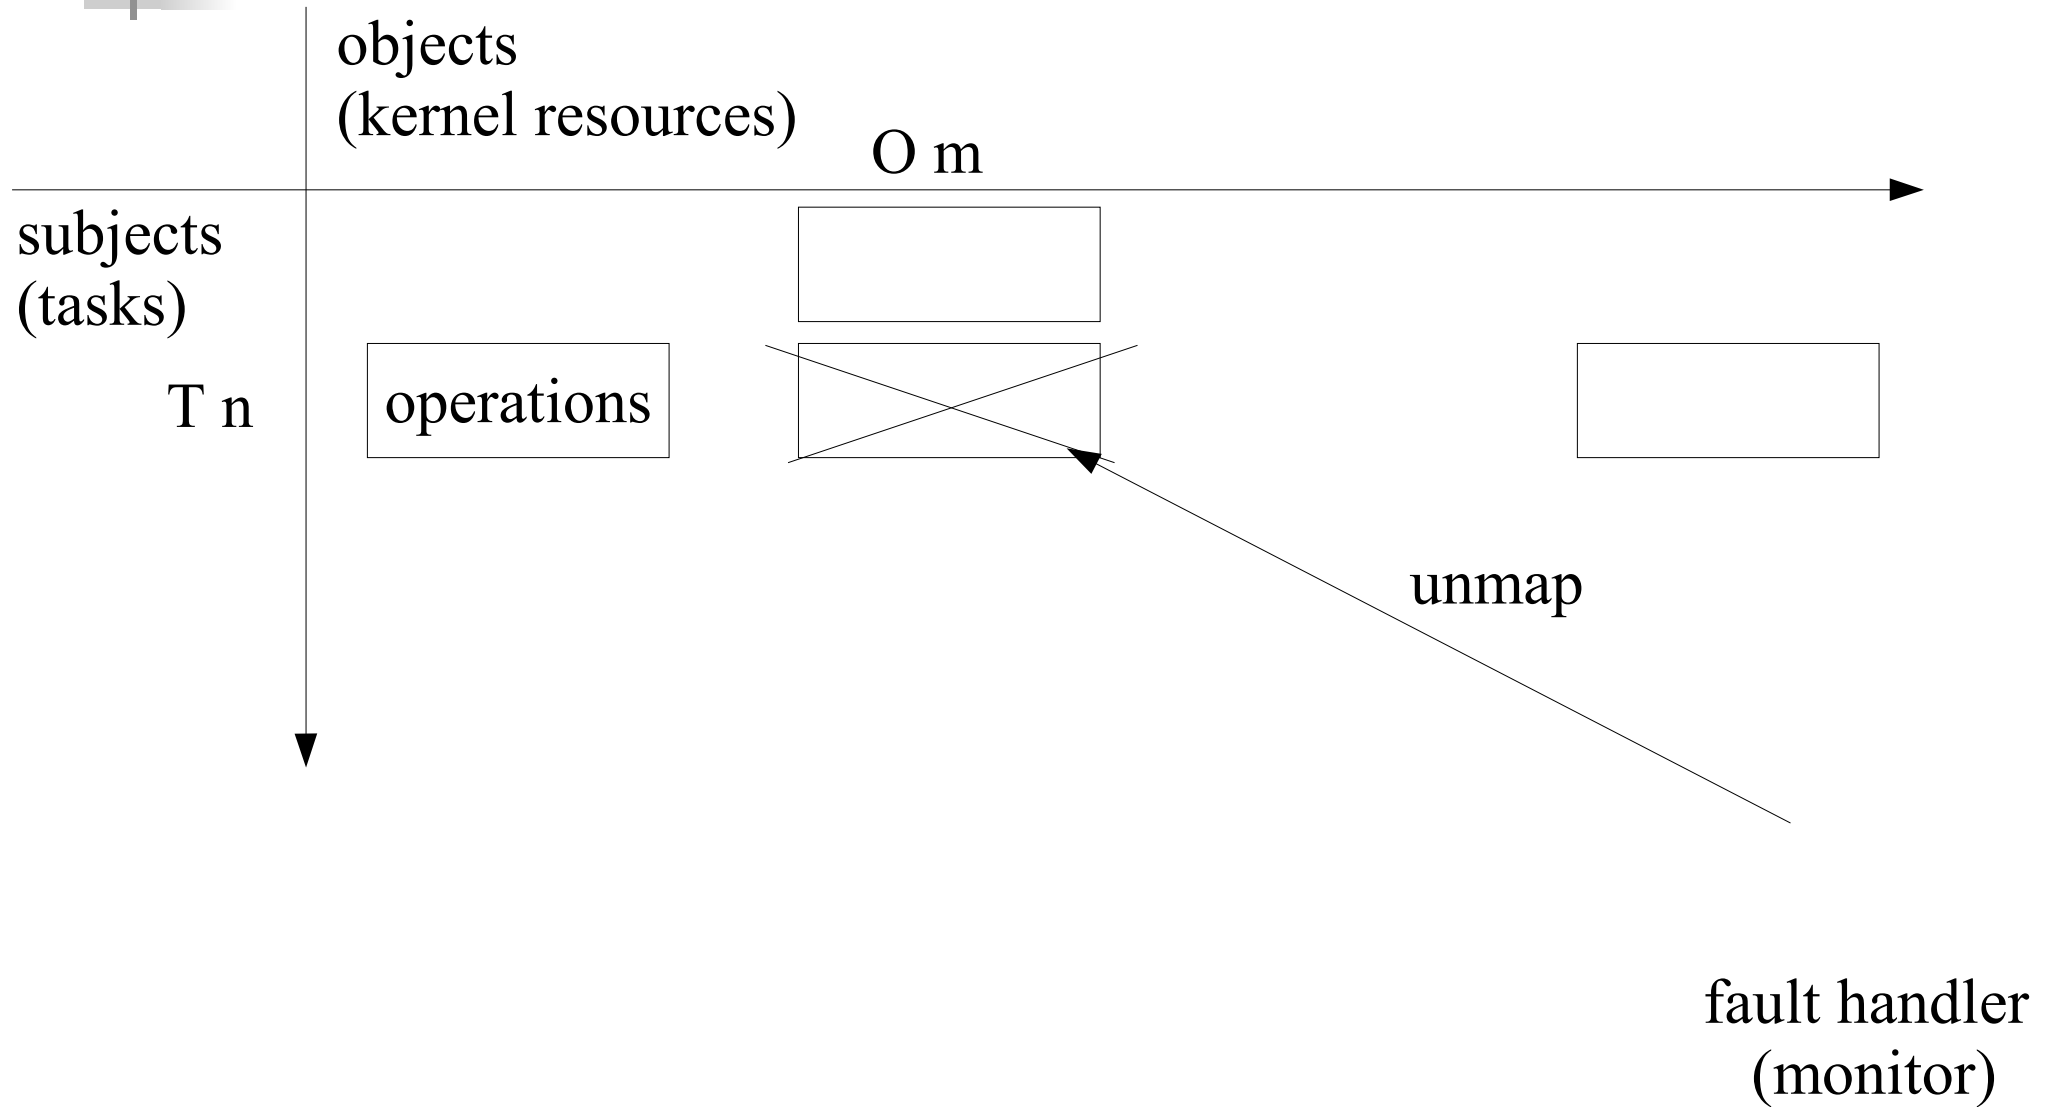
L4 X.e: Unifying Access Control



L4 X.e: Unifying Access Control



L4 X.e: Unifying Access Control





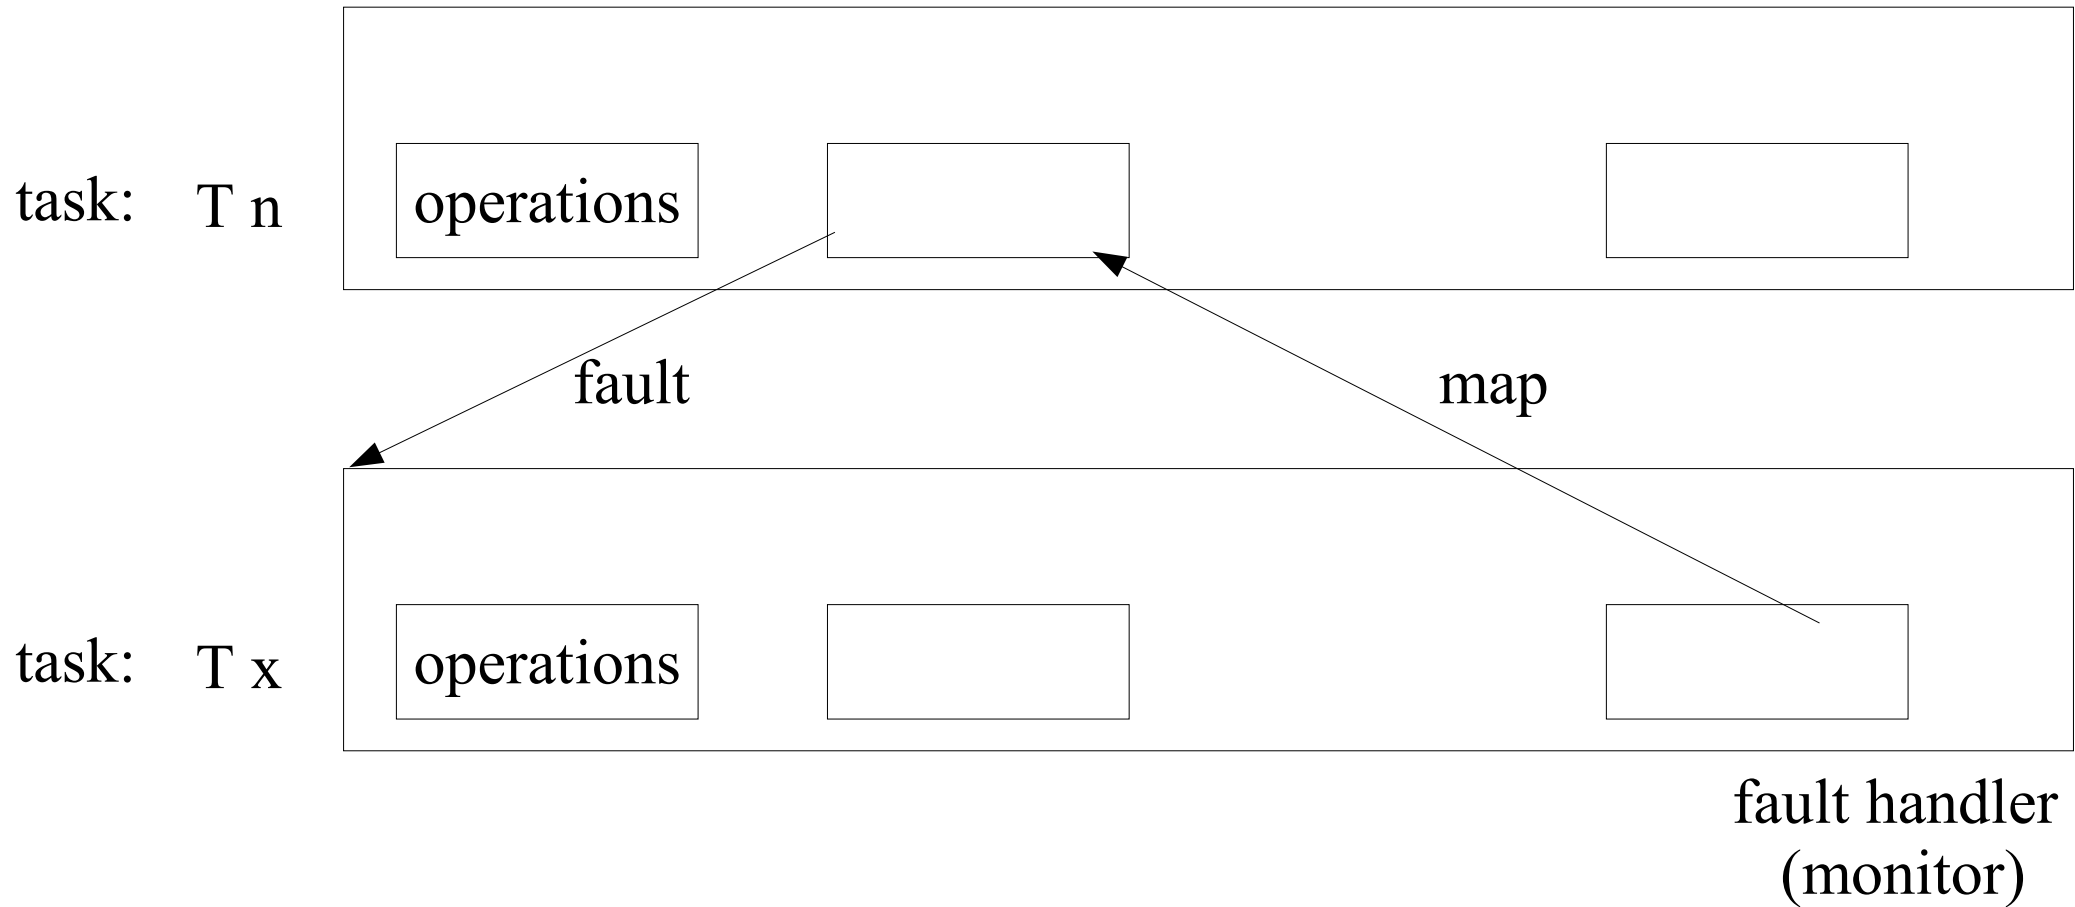
L4 X.e: Unifying Access Control

task: T_n

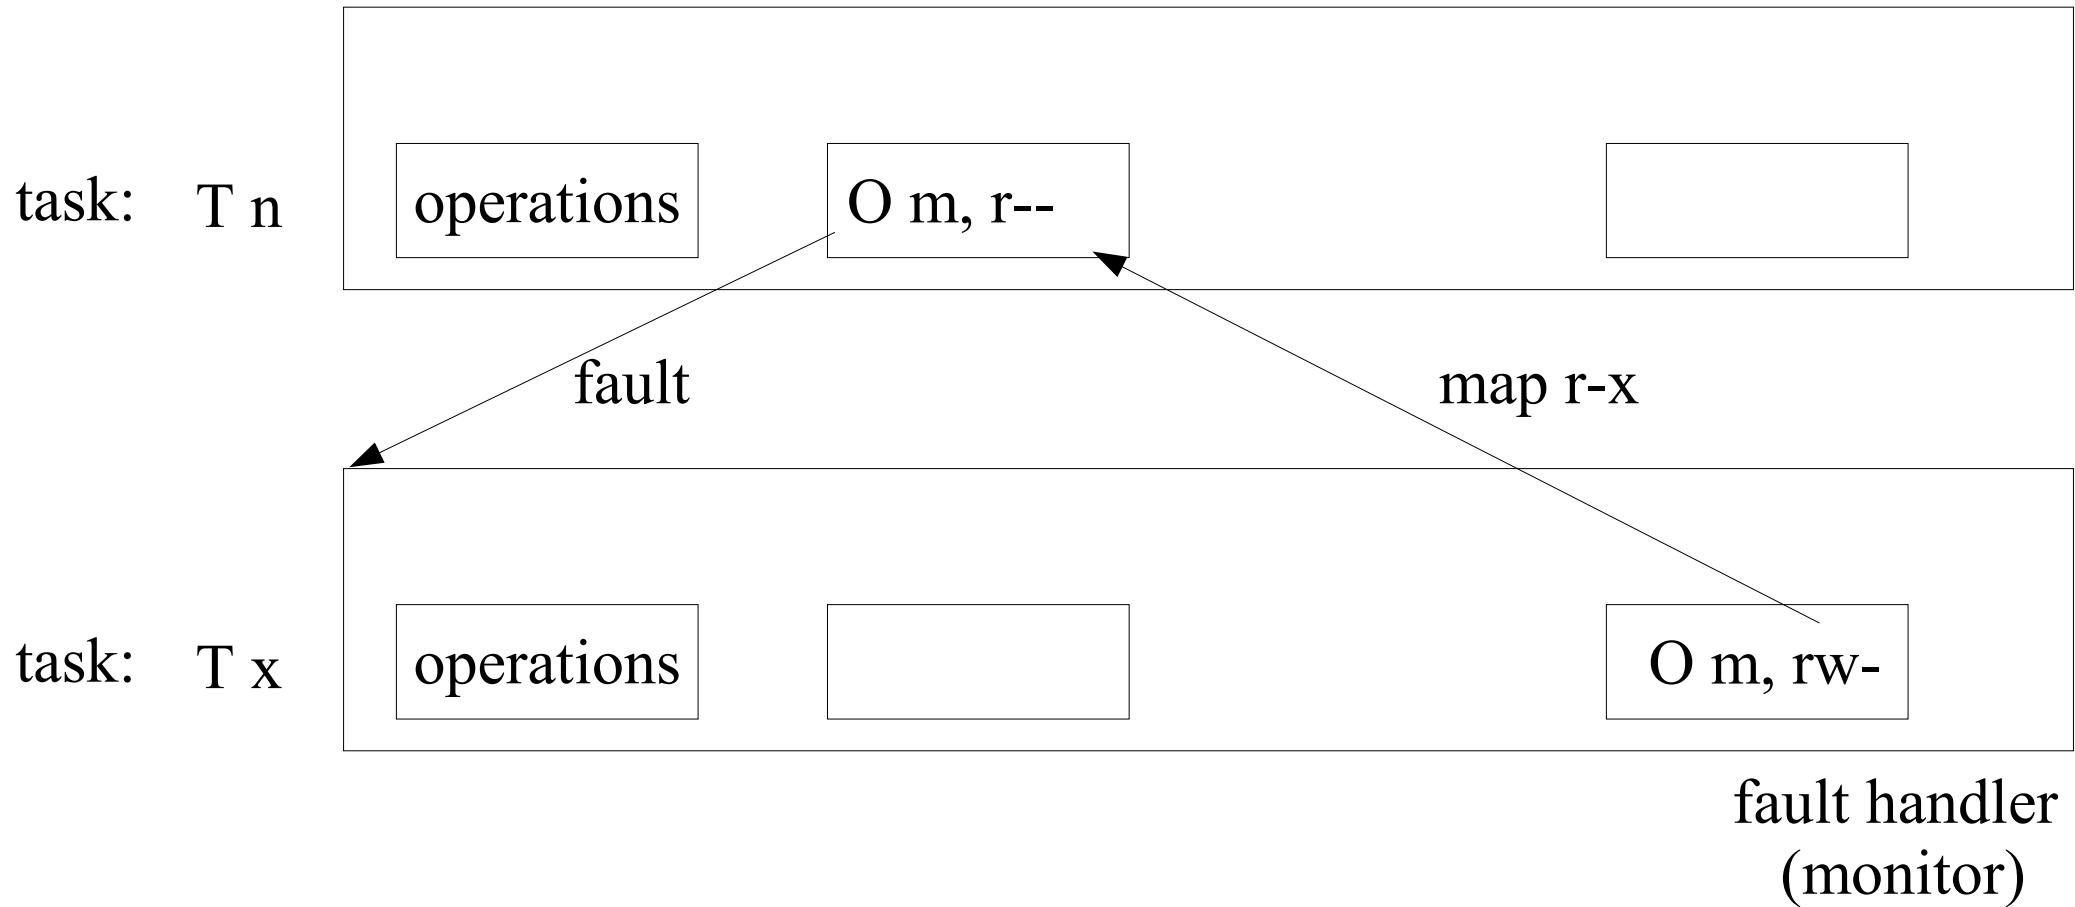


operations


L4 X.e: Unifying Access Control




L4 X.e: Unifying Access Control



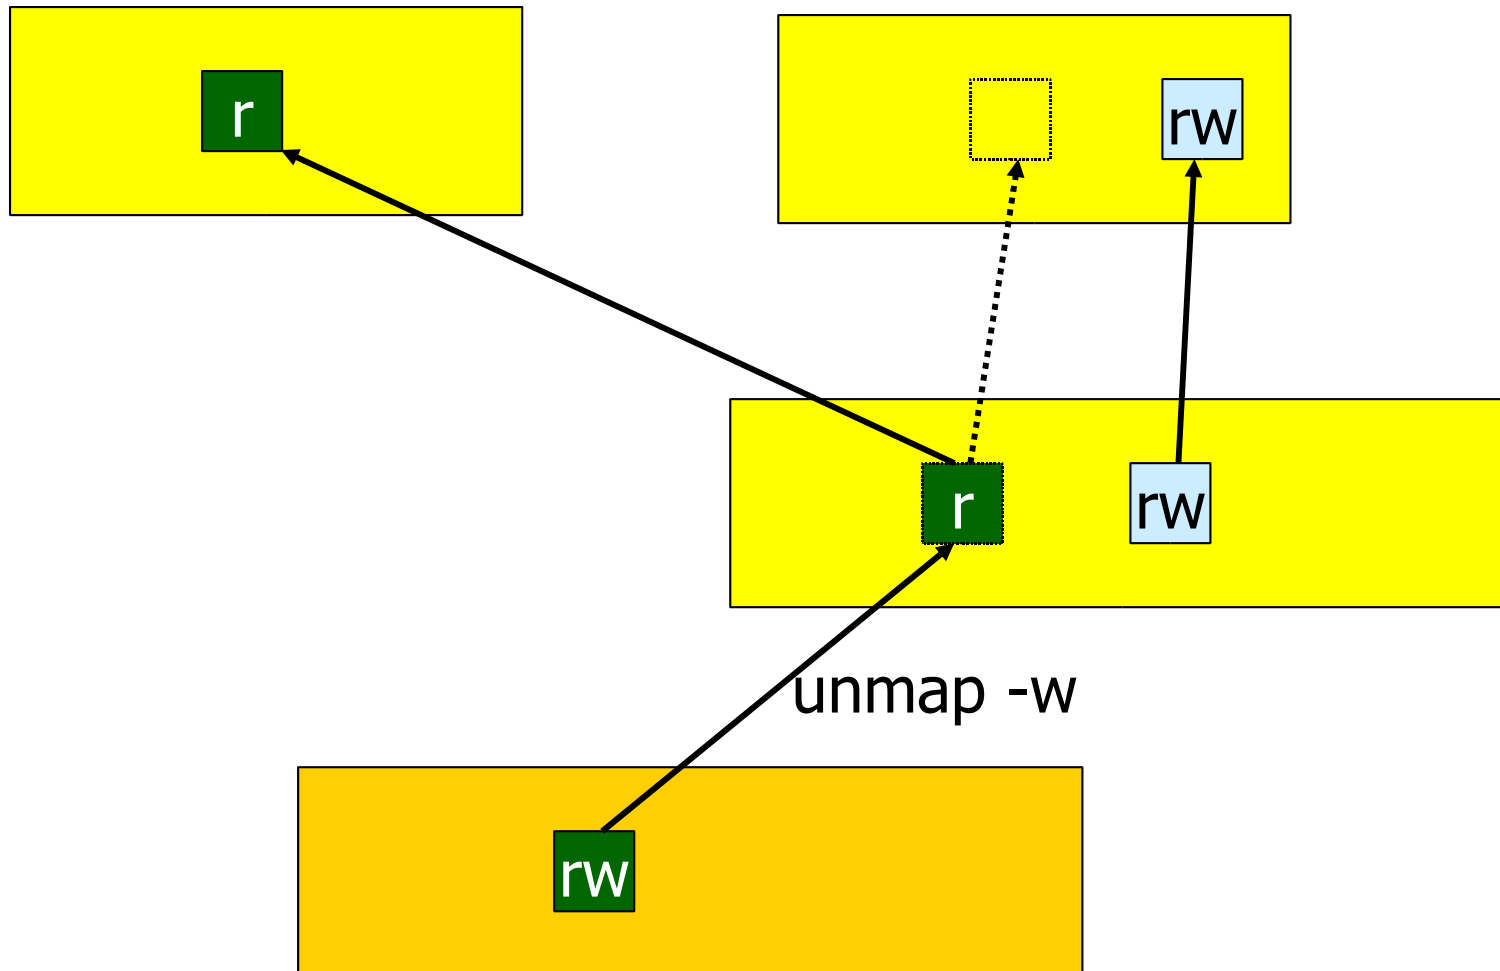
Memory Pages





Fault, Map, Unmap



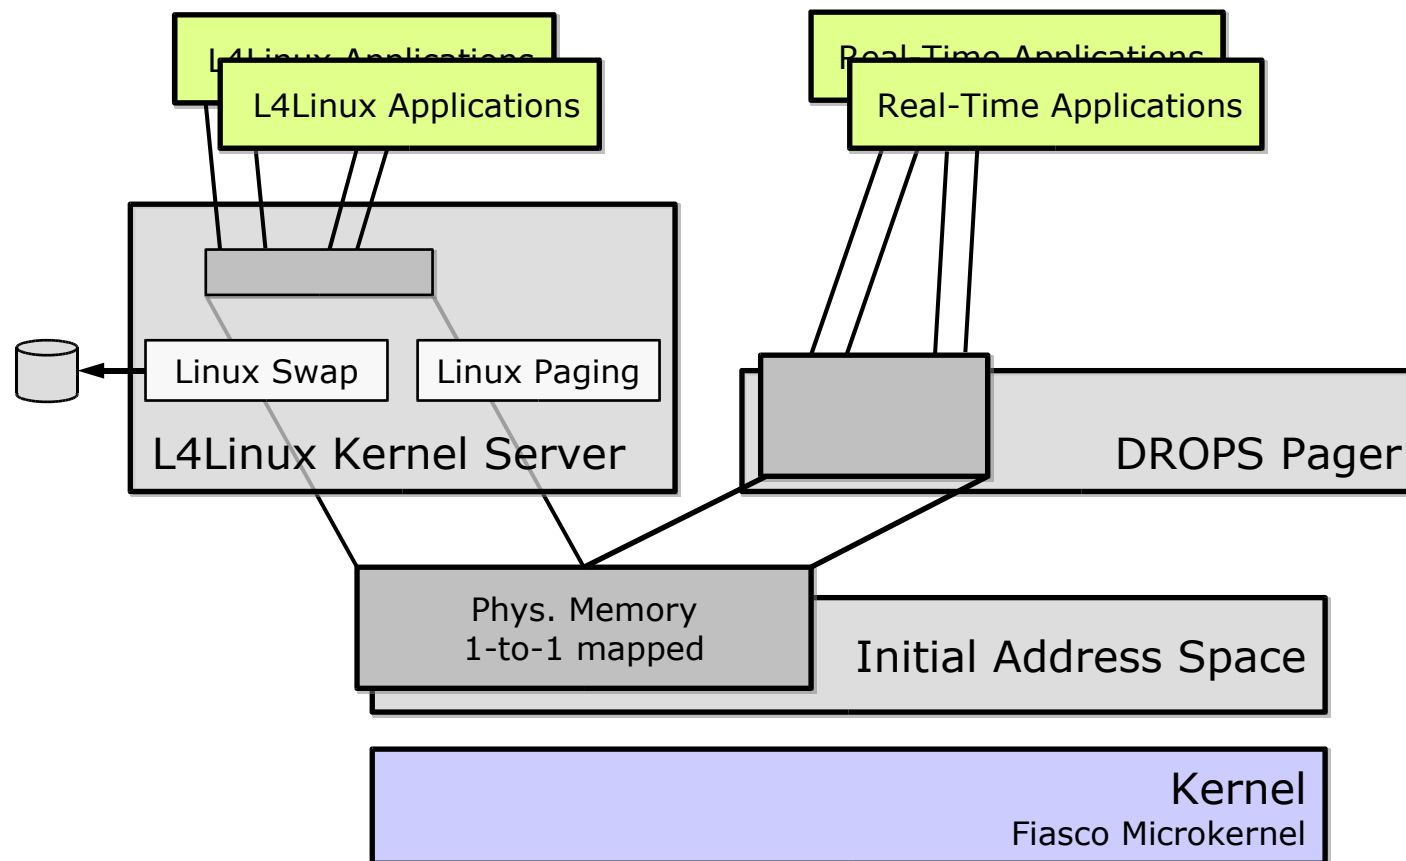
Fault, Map, Unmap




Map Trees



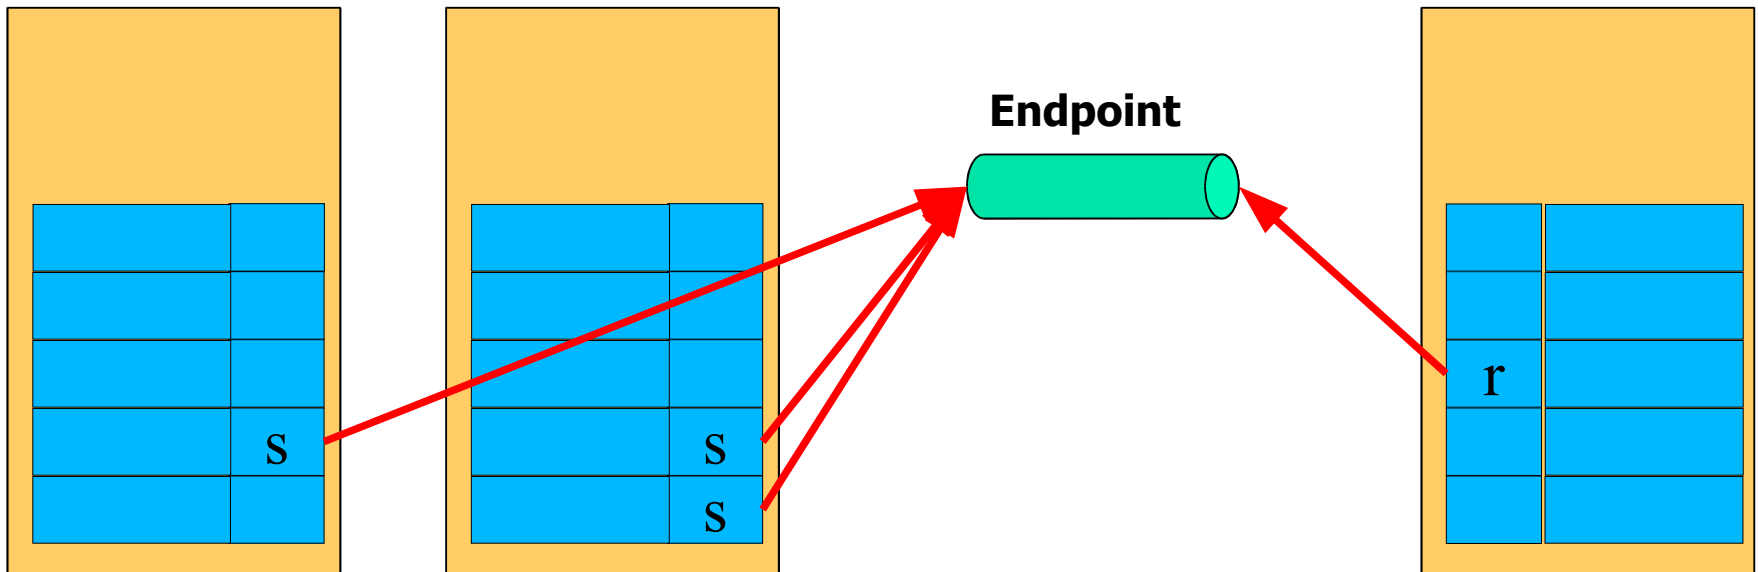
Pager Example



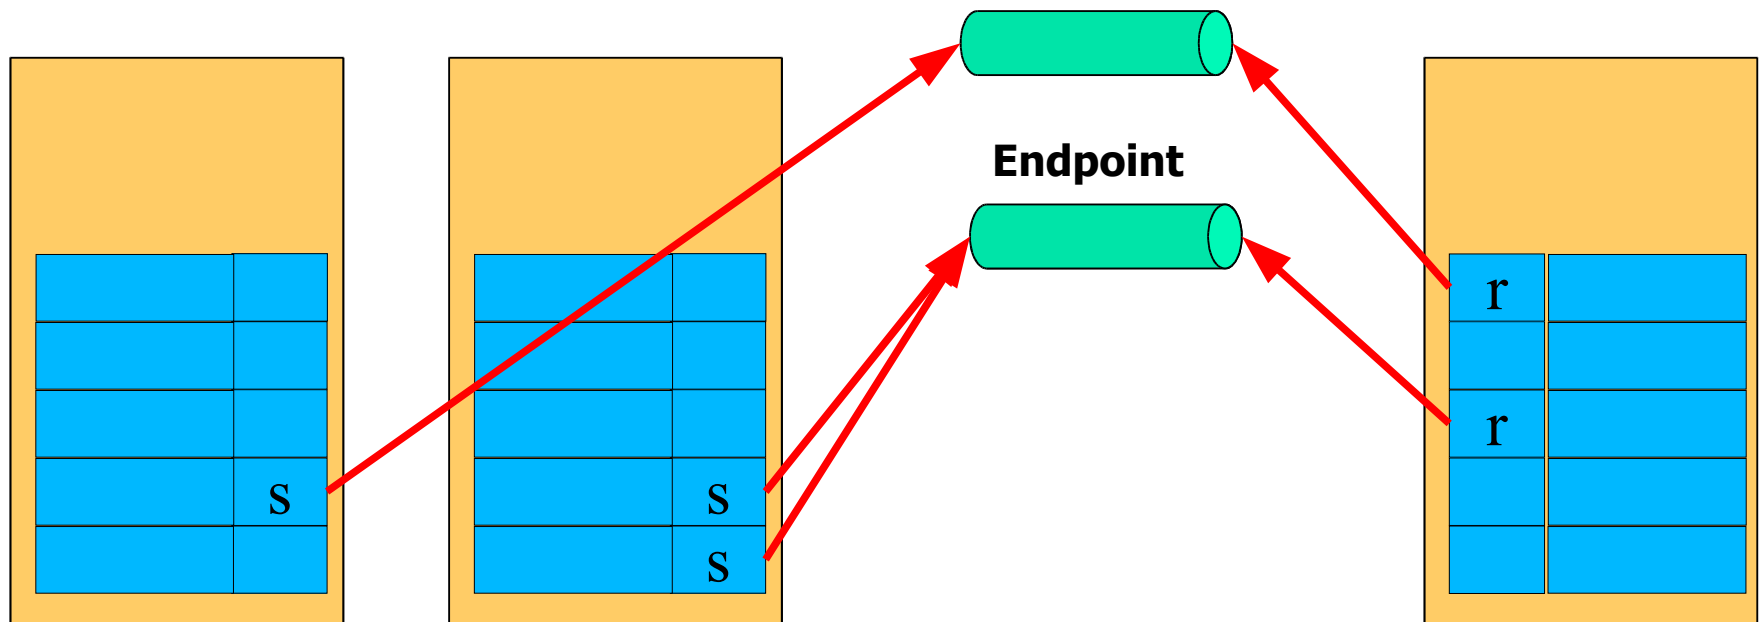
Communication



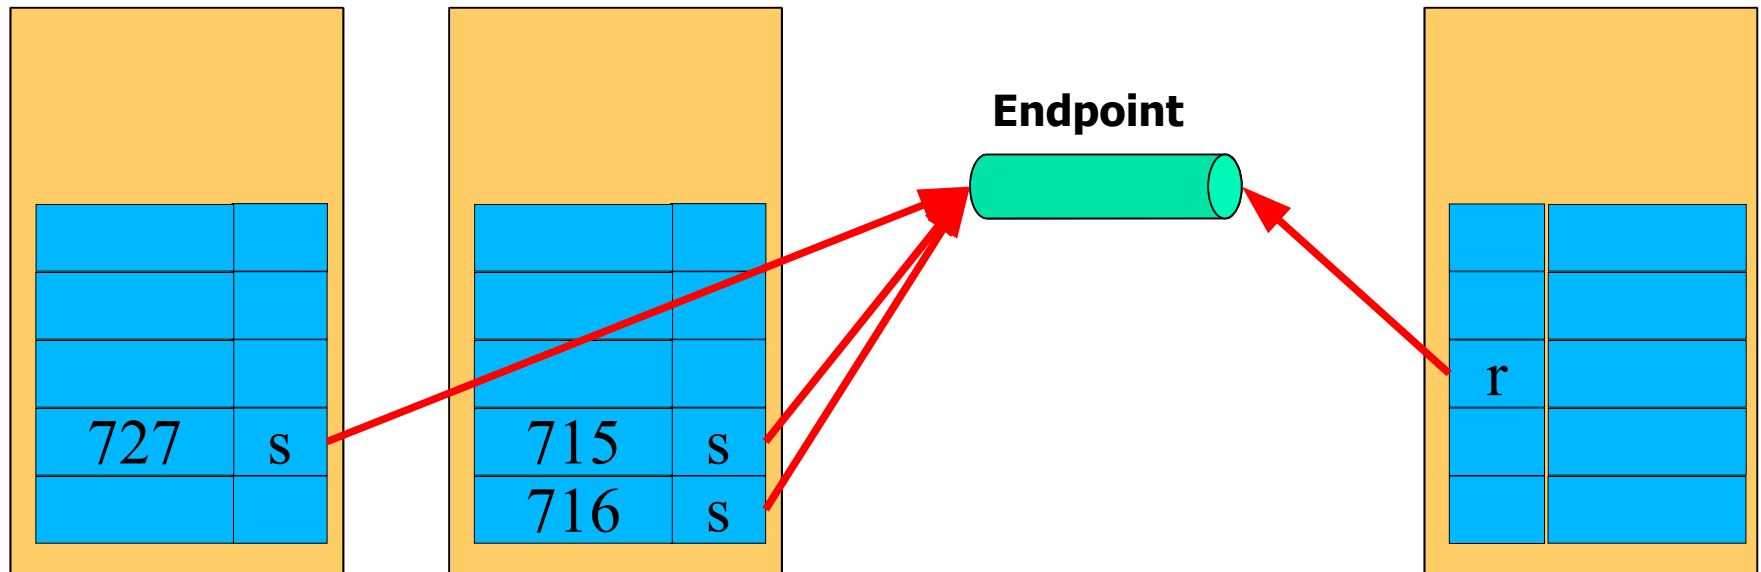
Communication




Communication




Communication




Communication




Communication




Communication

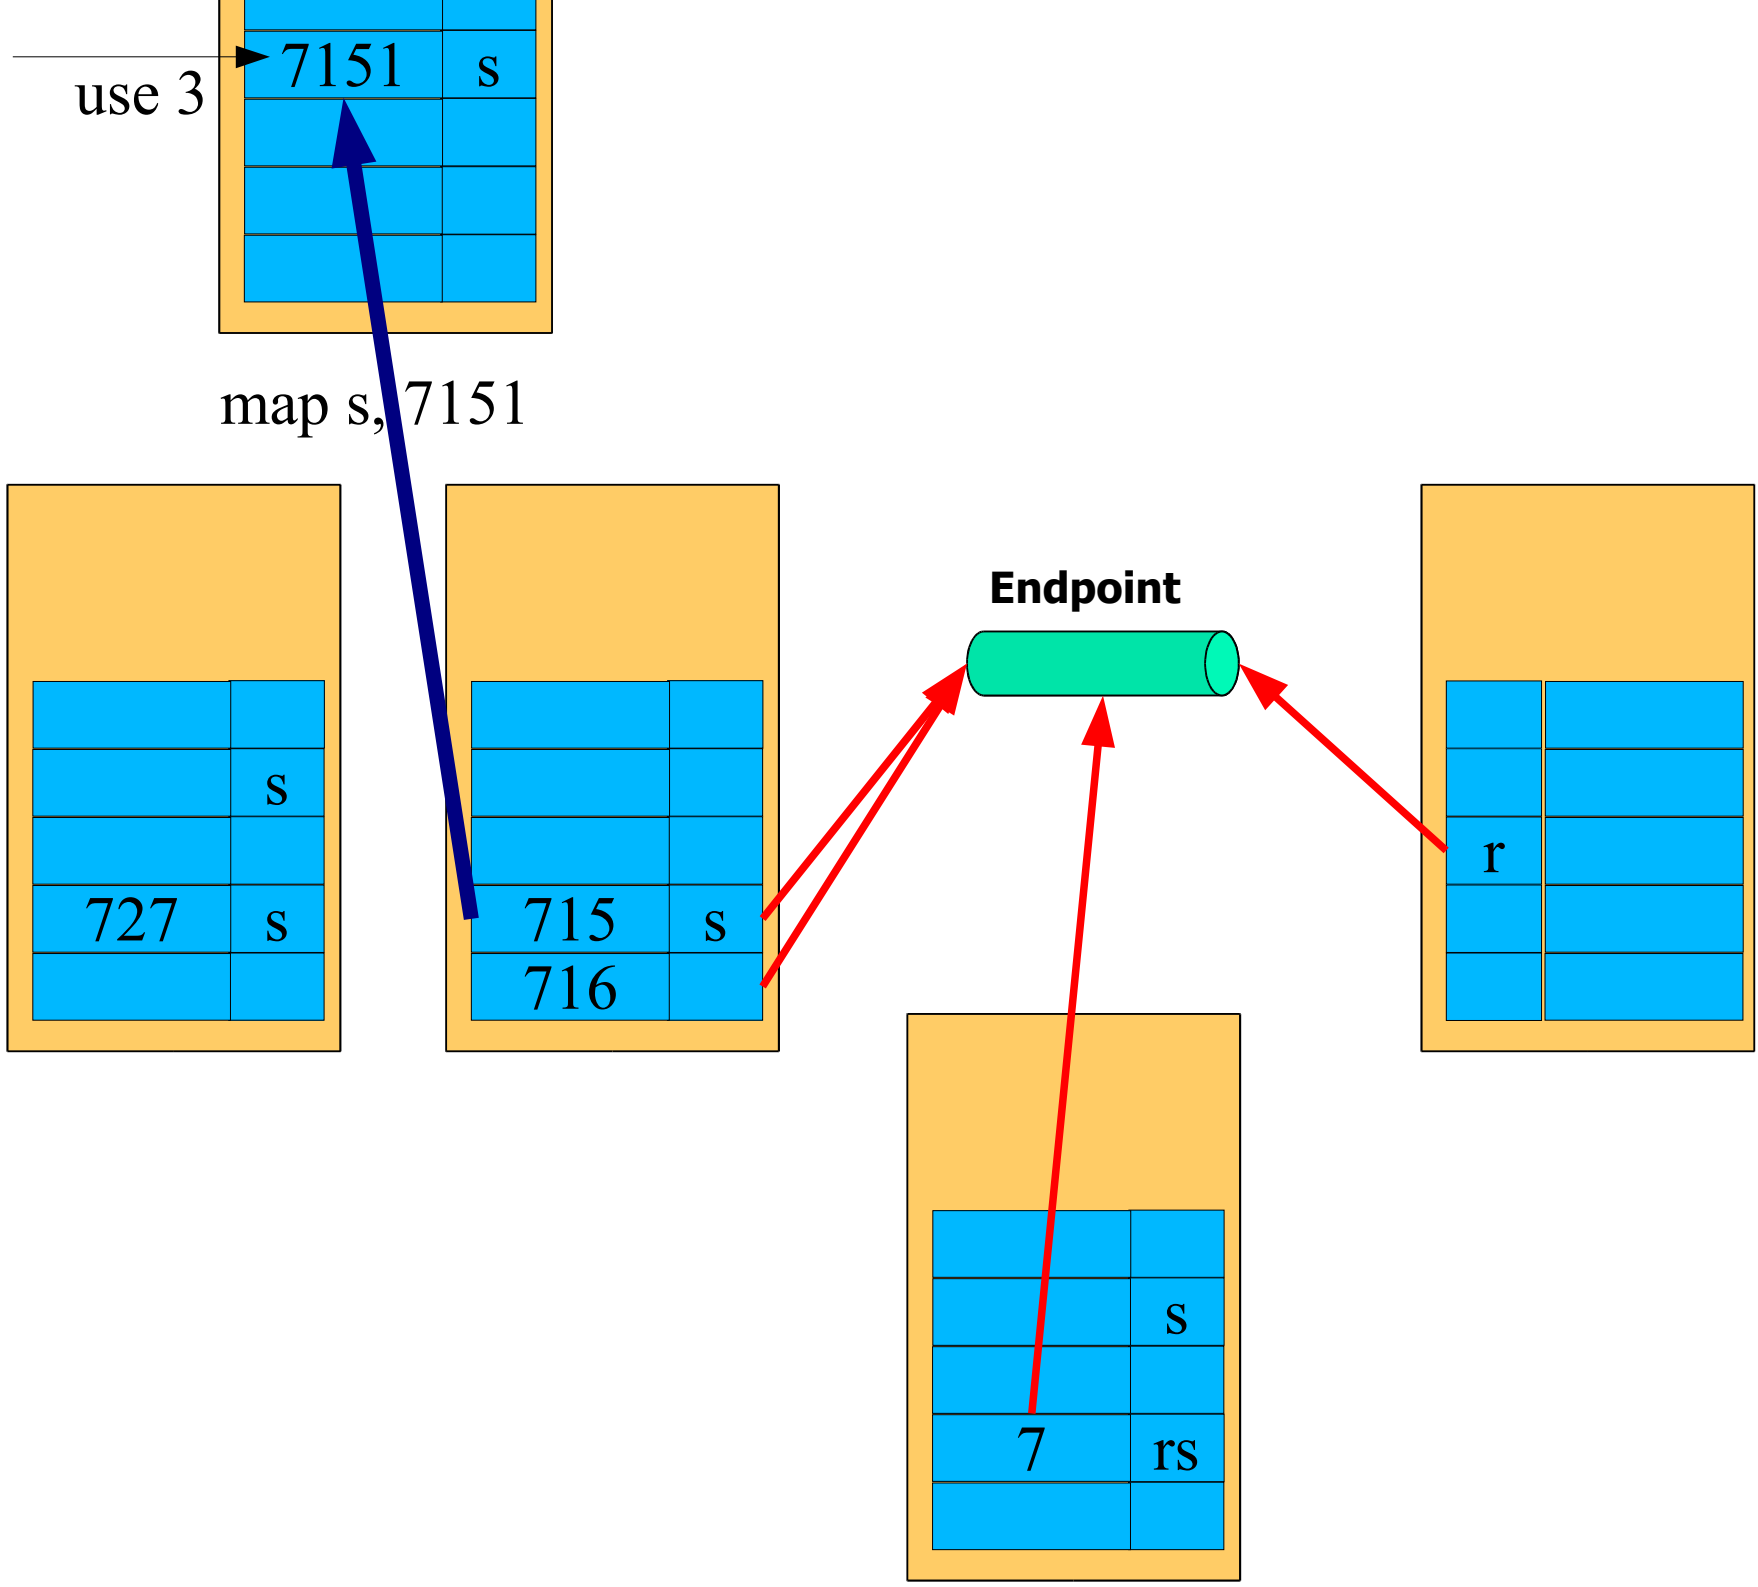



Communication



Communication









L⁴Linux




Linux Apps


L⁴Linux

L4 Fiasco

Pagers



Linux “top halves” as threads





Performance

Time-Sharing
Applications

L⁴Linux

Fiasco

(Härtig, Hohmuth, Liedtke, Schönberg, Wolter:
The Performance of μ -Kernel based Systems,
SOSP 1997)

jobs per minute

simulated load

70


L⁴Linux compared to MACH

Time-Sharing
Applications

L⁴Linux

Fiasco

AIM Suite-VII benchmark - jobs per minute

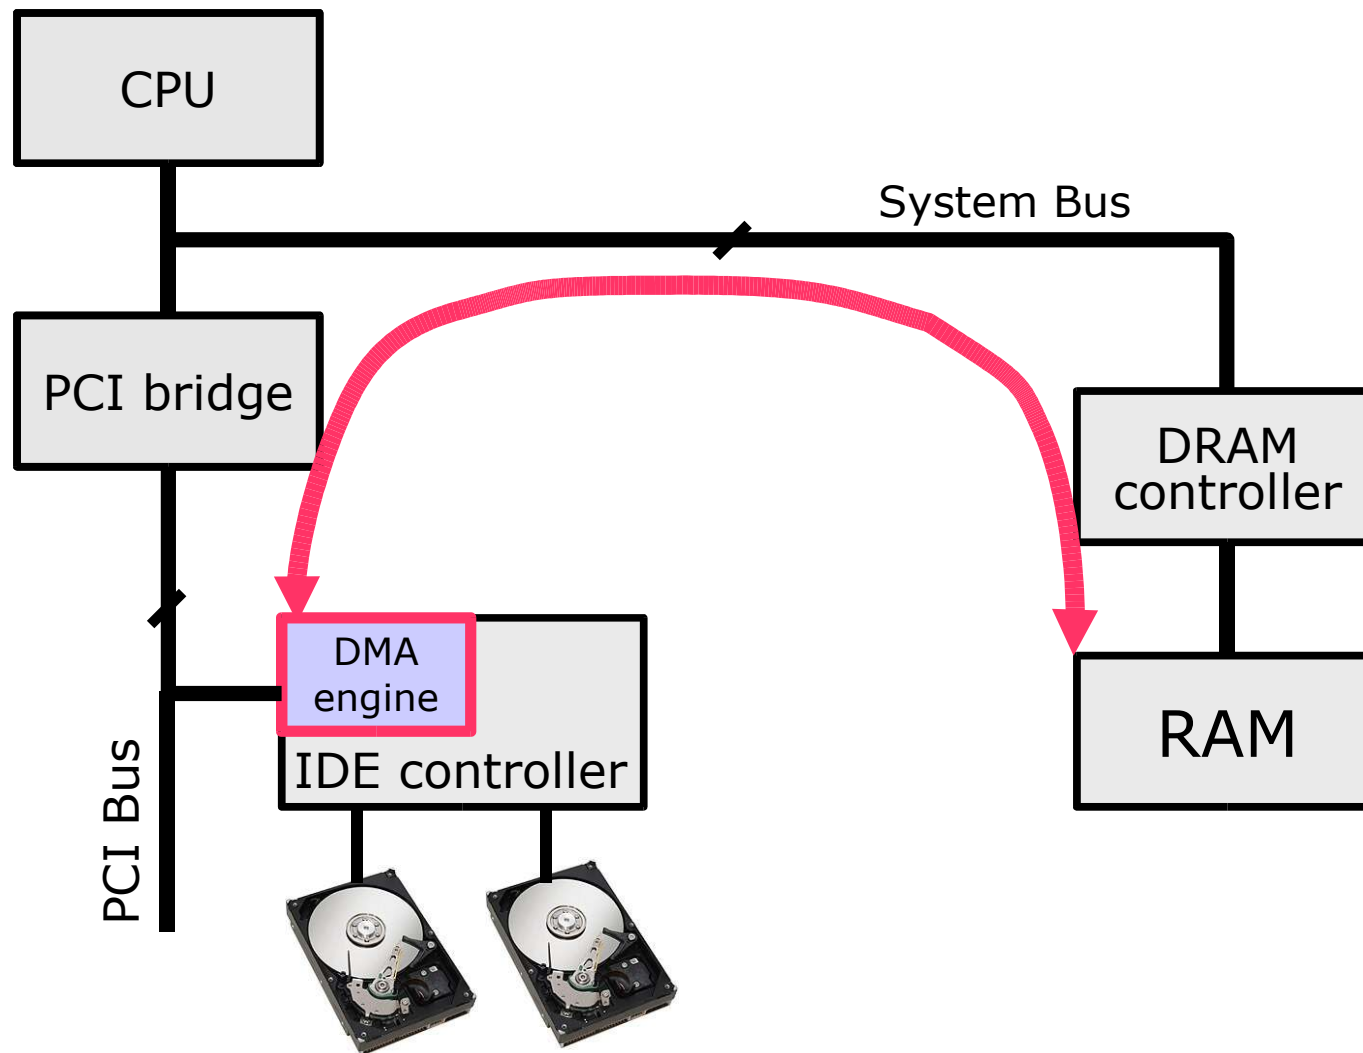




The DMA Problem

- separation is enforced by MMU
 - devices access memory using bus master DMA
 - DMA uses physical addresses (on most architectures)
- malicious devices
(or malicious device drivers, firmware!)
can access/modify all components of a system

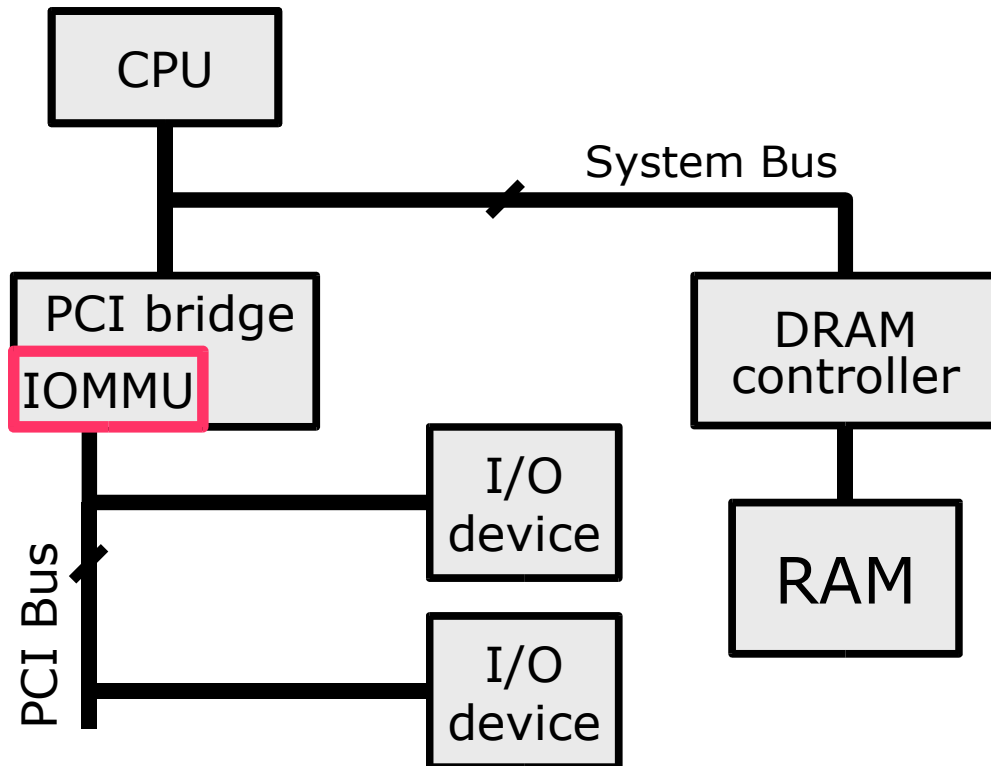
The DMA Problem



IOMMU

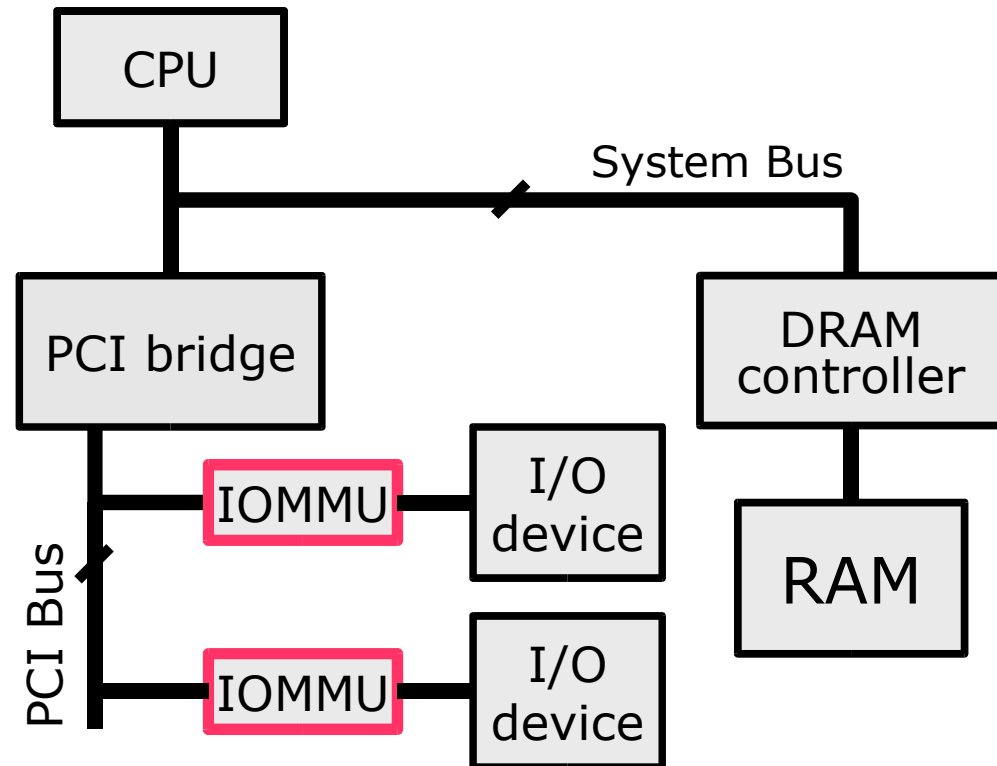
per-bus IOMMU:

- IA64 chipsets
- Opteron



per-device IOMMU:

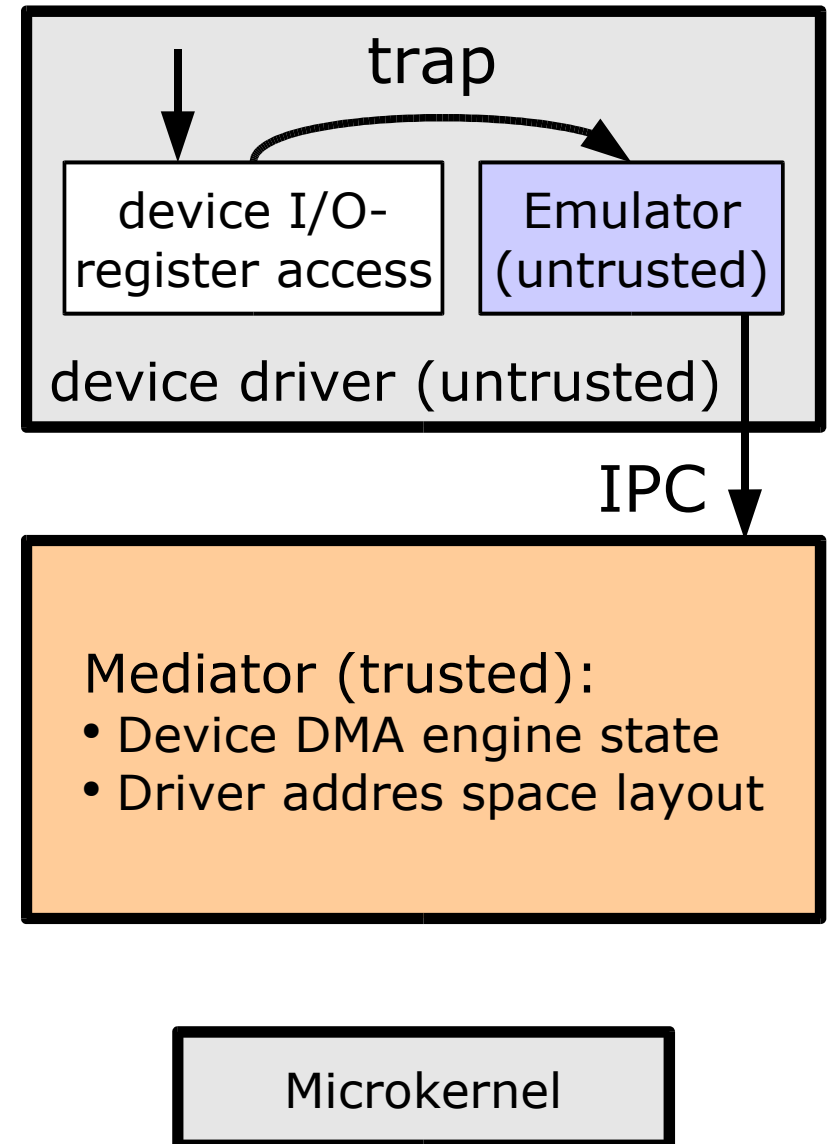
- **n/a** in current hardware!



IOMMU in Software

Make sure that device cannot perform malicious DMA:

- **Trap** read/write access to I/O-registers of the device
- An **emulator** (untrusted) determines size and value to be read/written
- A **mediator** (trusted) checks and performs the access





IOMMU in Software

- **Emulator:**

- in driver's address space → untrusted
- malfunction does only decrease availability of device
- ~ 500 LoC

- **Mediator:**

- trusted, own address space
- specific for a device or a class of devices
- ~ 300 LoC




IOMMU in Software

- Implemented for:
 - Fast Ethernet card (DEC Tulip 21143)
 - ATA Controllers
- Does not work for firmware-programmable devices:
 - private interface between device driver and firmware
 - no protection mechanism between firmware and device


IOMMU in Software: NIC Performance

Pentium-III 800 MHz, Fast Ethernet

TCP Rate [MBit/s]




CPU Load [%]




IOMMU in Software: ATA Performance

Pentium-III 800 MHz, VIA82C586 ATA Controller


Total Time [s]



CPU Load [%]




Secure Booting, Remote Attestation and Trusted Path



Authentic application/system:

- how does the remote bank know ?
- how does the local client/user know ?

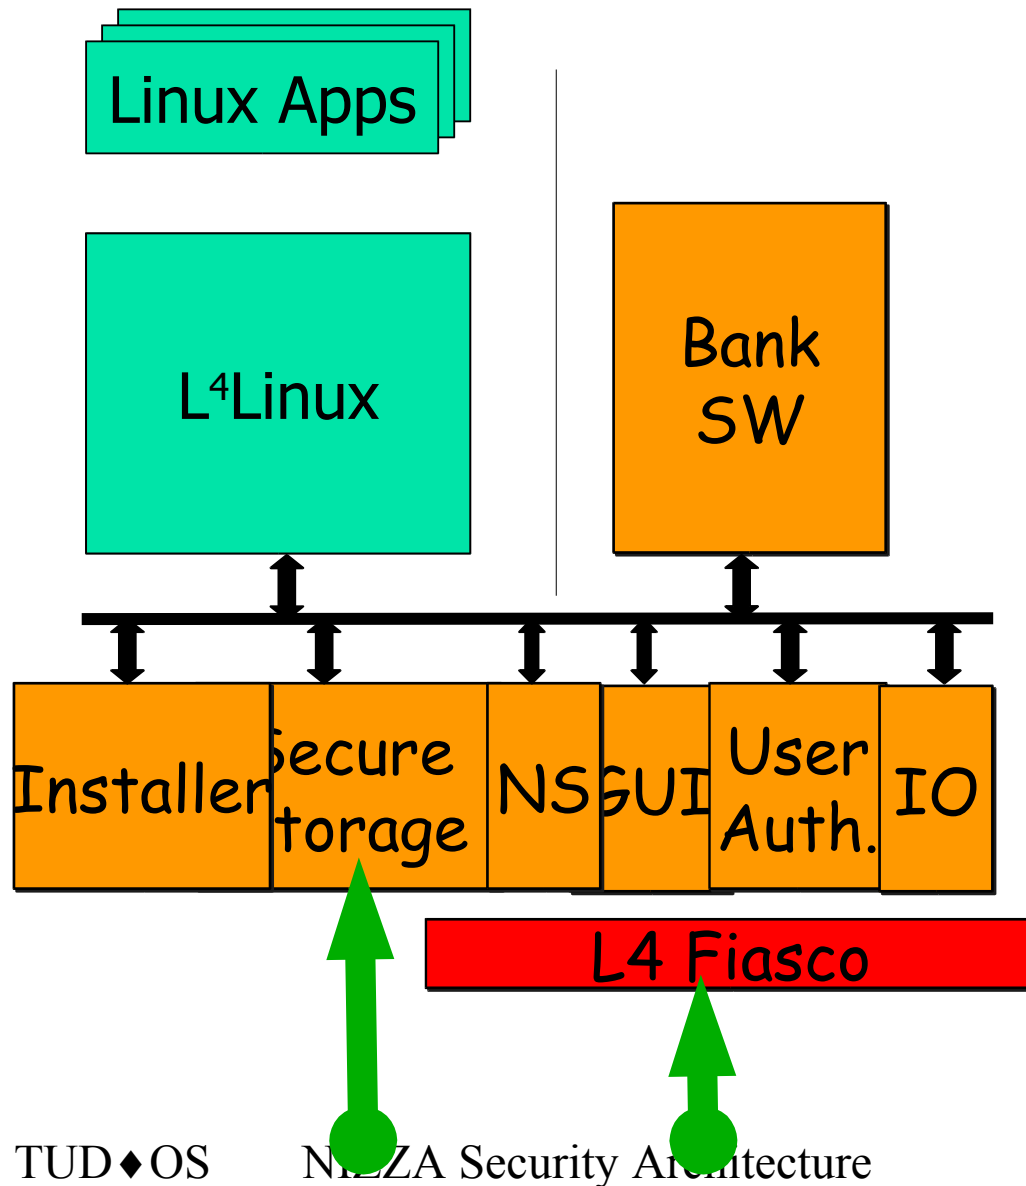
Secure Booting, Remote Attestation and Trusted Path



Authentic application/system:

- how does the remote bank know ?
- how does the local client/user know ?

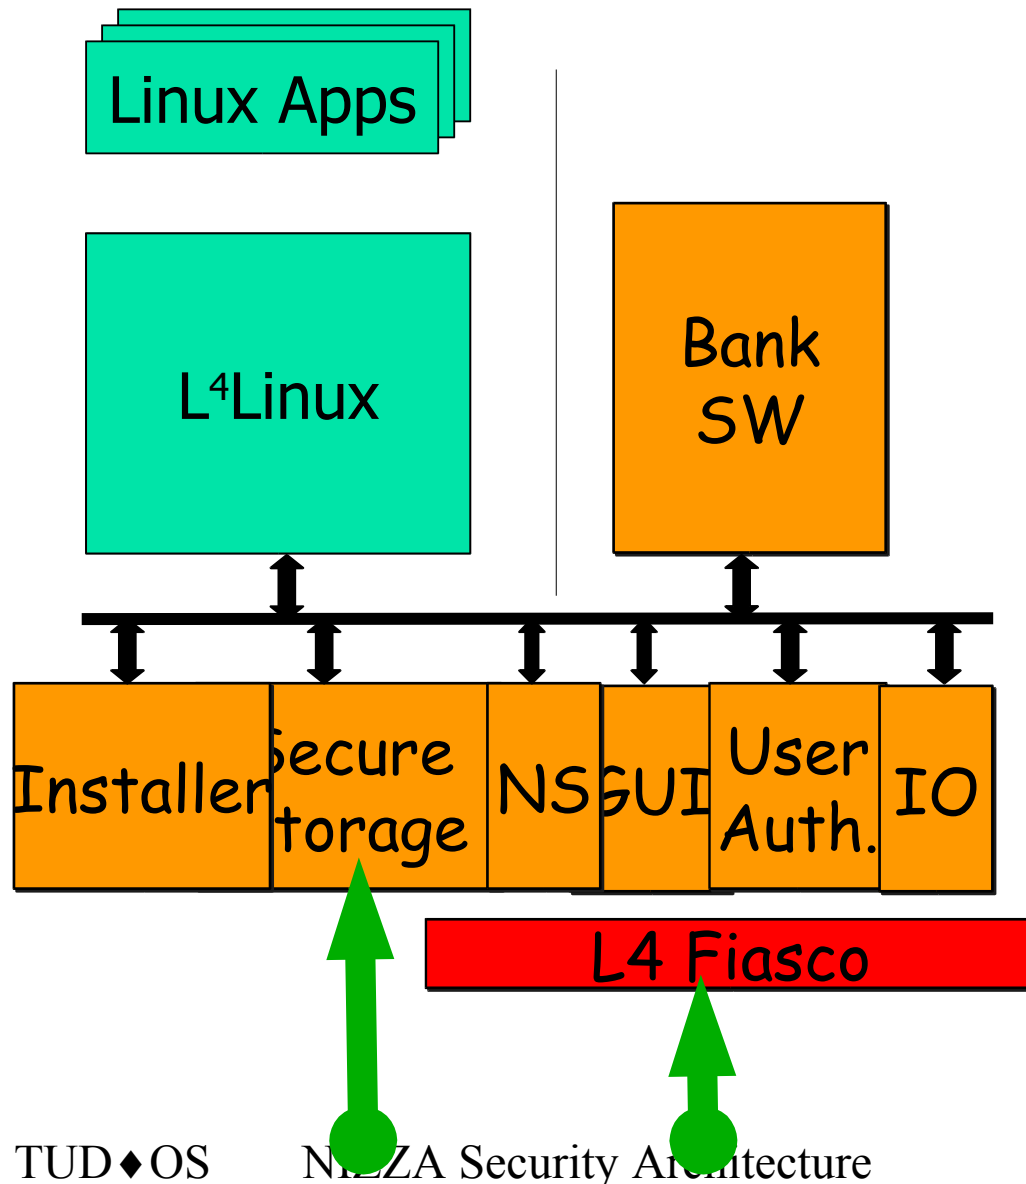
Secure Booting, Remote Attestation and Trusted Path



Authentic application/system:

- how does the remote bank know ?
 - attestation protocol up to Nizza trusted platform
 - mediate other communication thru trusted installer

Secure Booting, Remote Attestation and Trusted Path



Authentic application/system:

- how does the local client/user know ?
 - attestation protocol up to Nizza trusted platform
 - indicate “red/green”
 - handover to DOpE



Secure Storage with small TCB (future)

- objectives:
 - security
 - confidentiality, integrity,
 - recoverability
 - availability
 - system security
 - small TCB
 - attacks:
 - theft/loss of device
 - full penetration of L4Linux

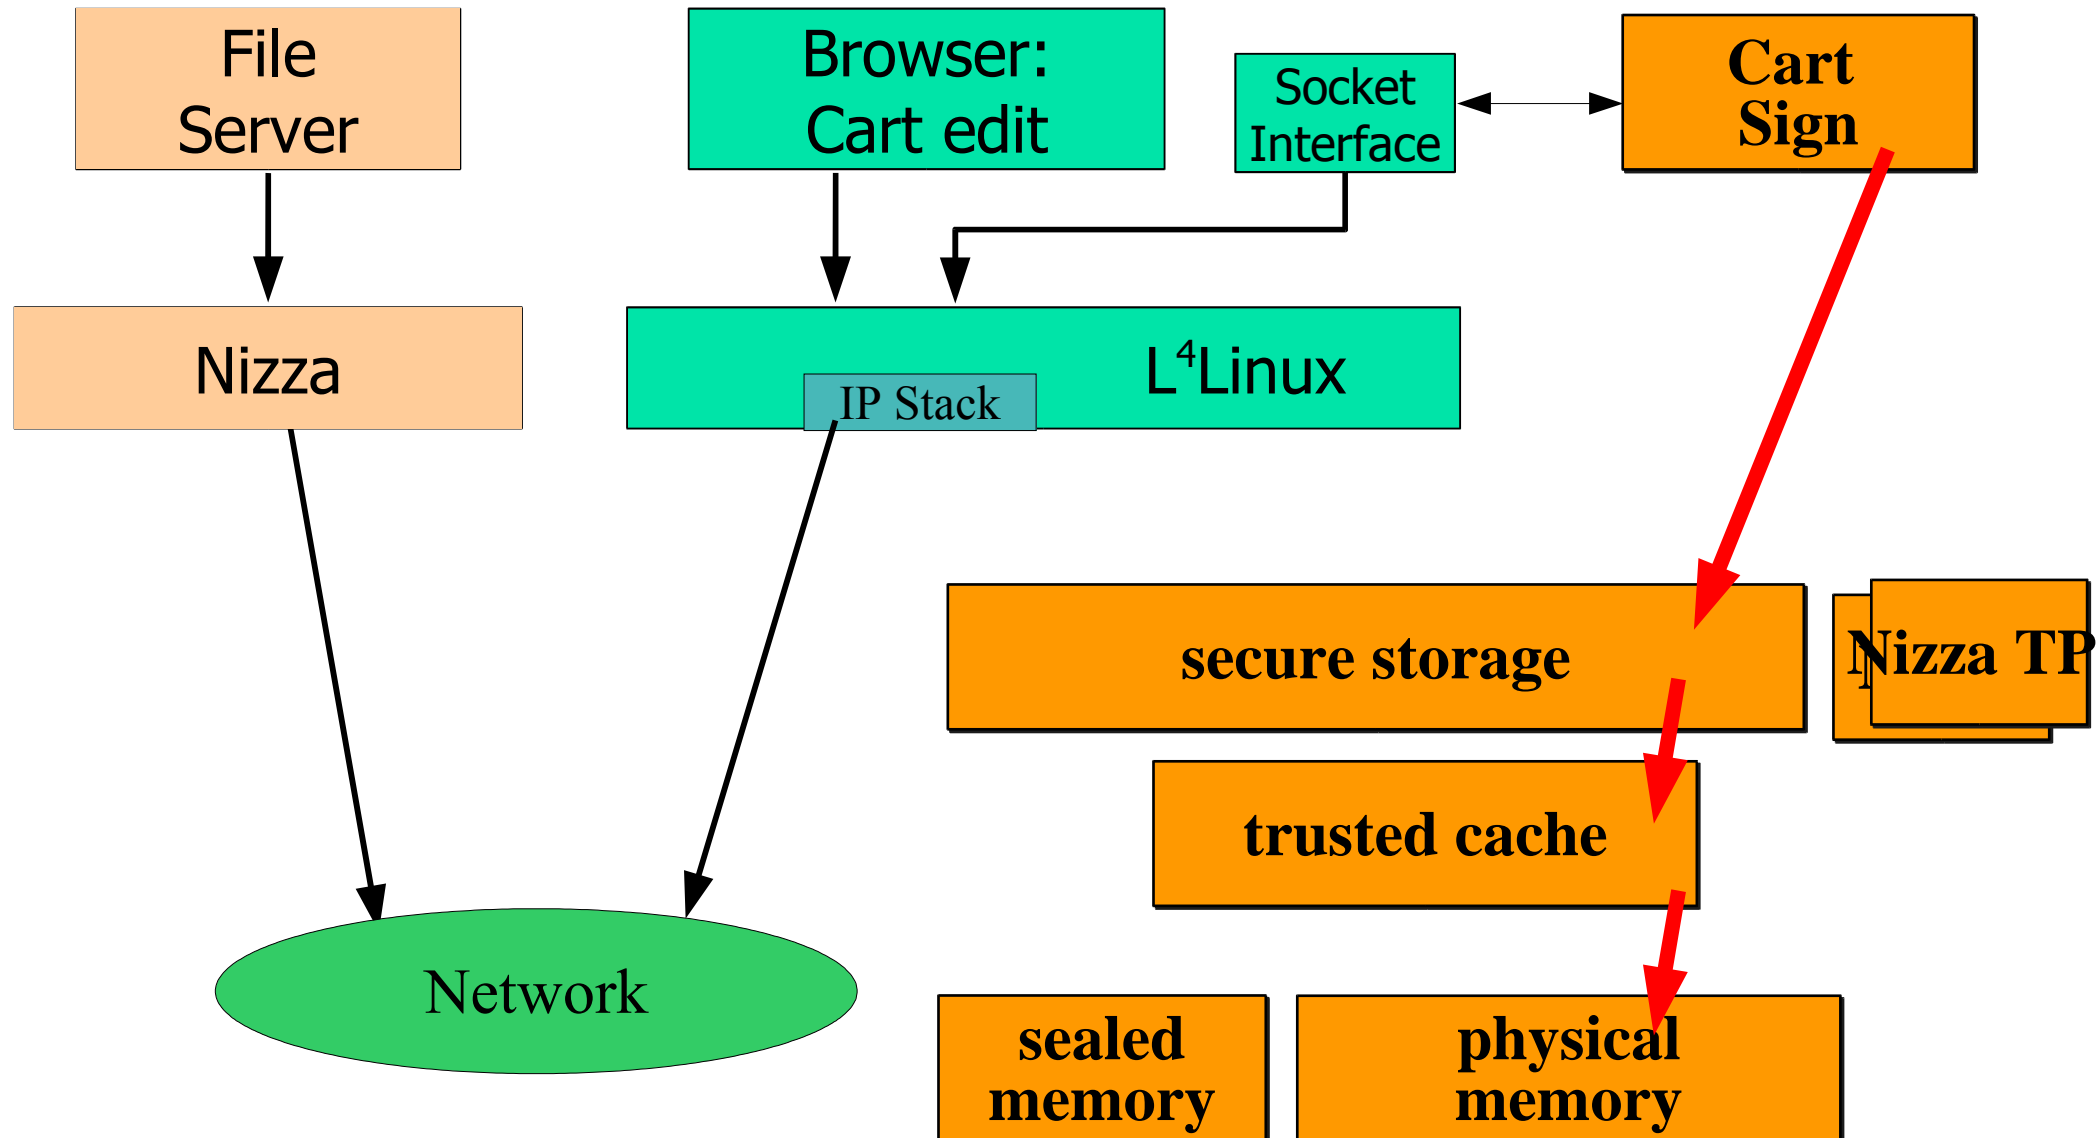
•



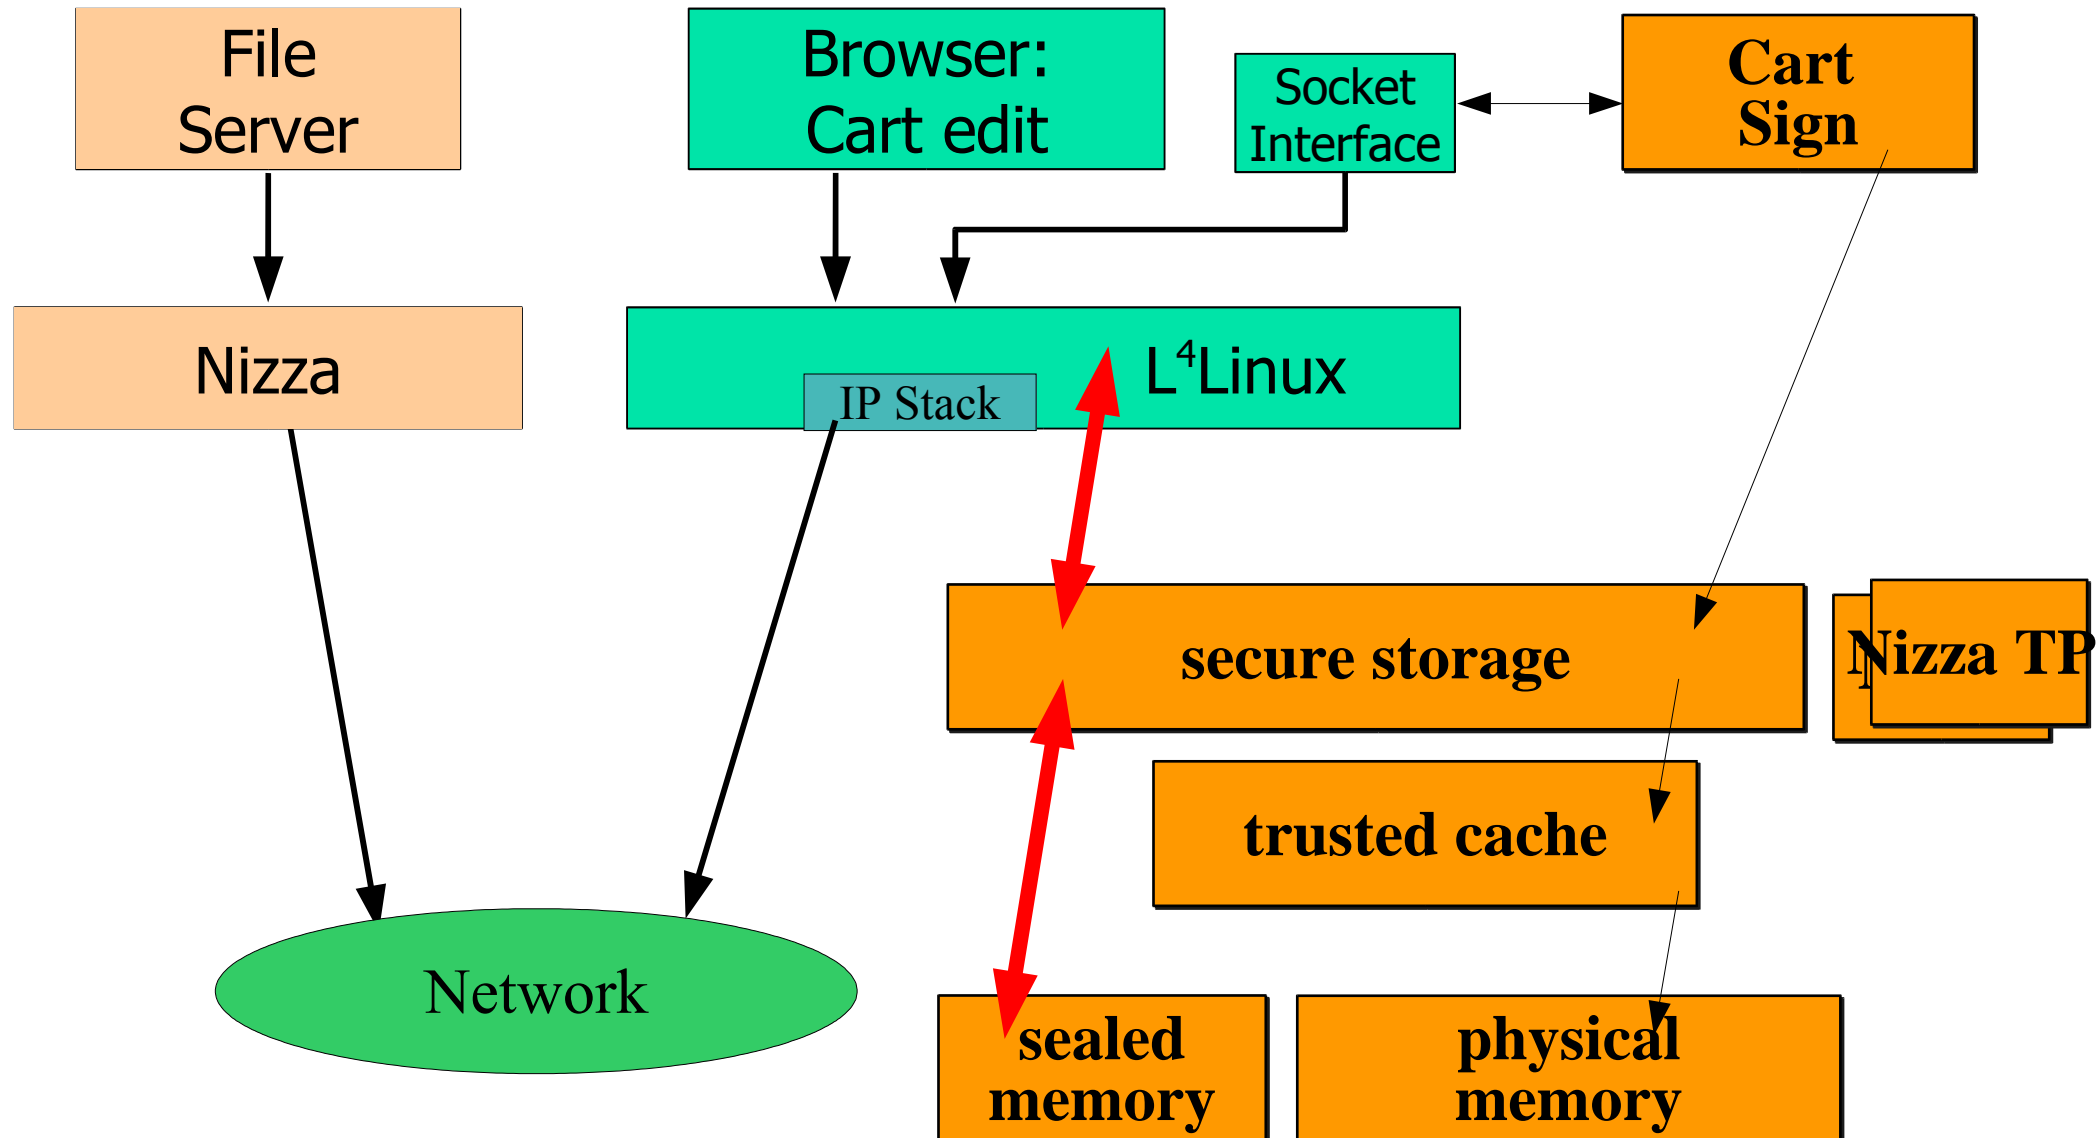
Secure Storage with small TCB (future)

- techniques:
 - use Sealed Memory as key storage
 - reuse L4Linux file system as mass storage
 - use trusted file server for recoverability
 - use resource allocation for availability

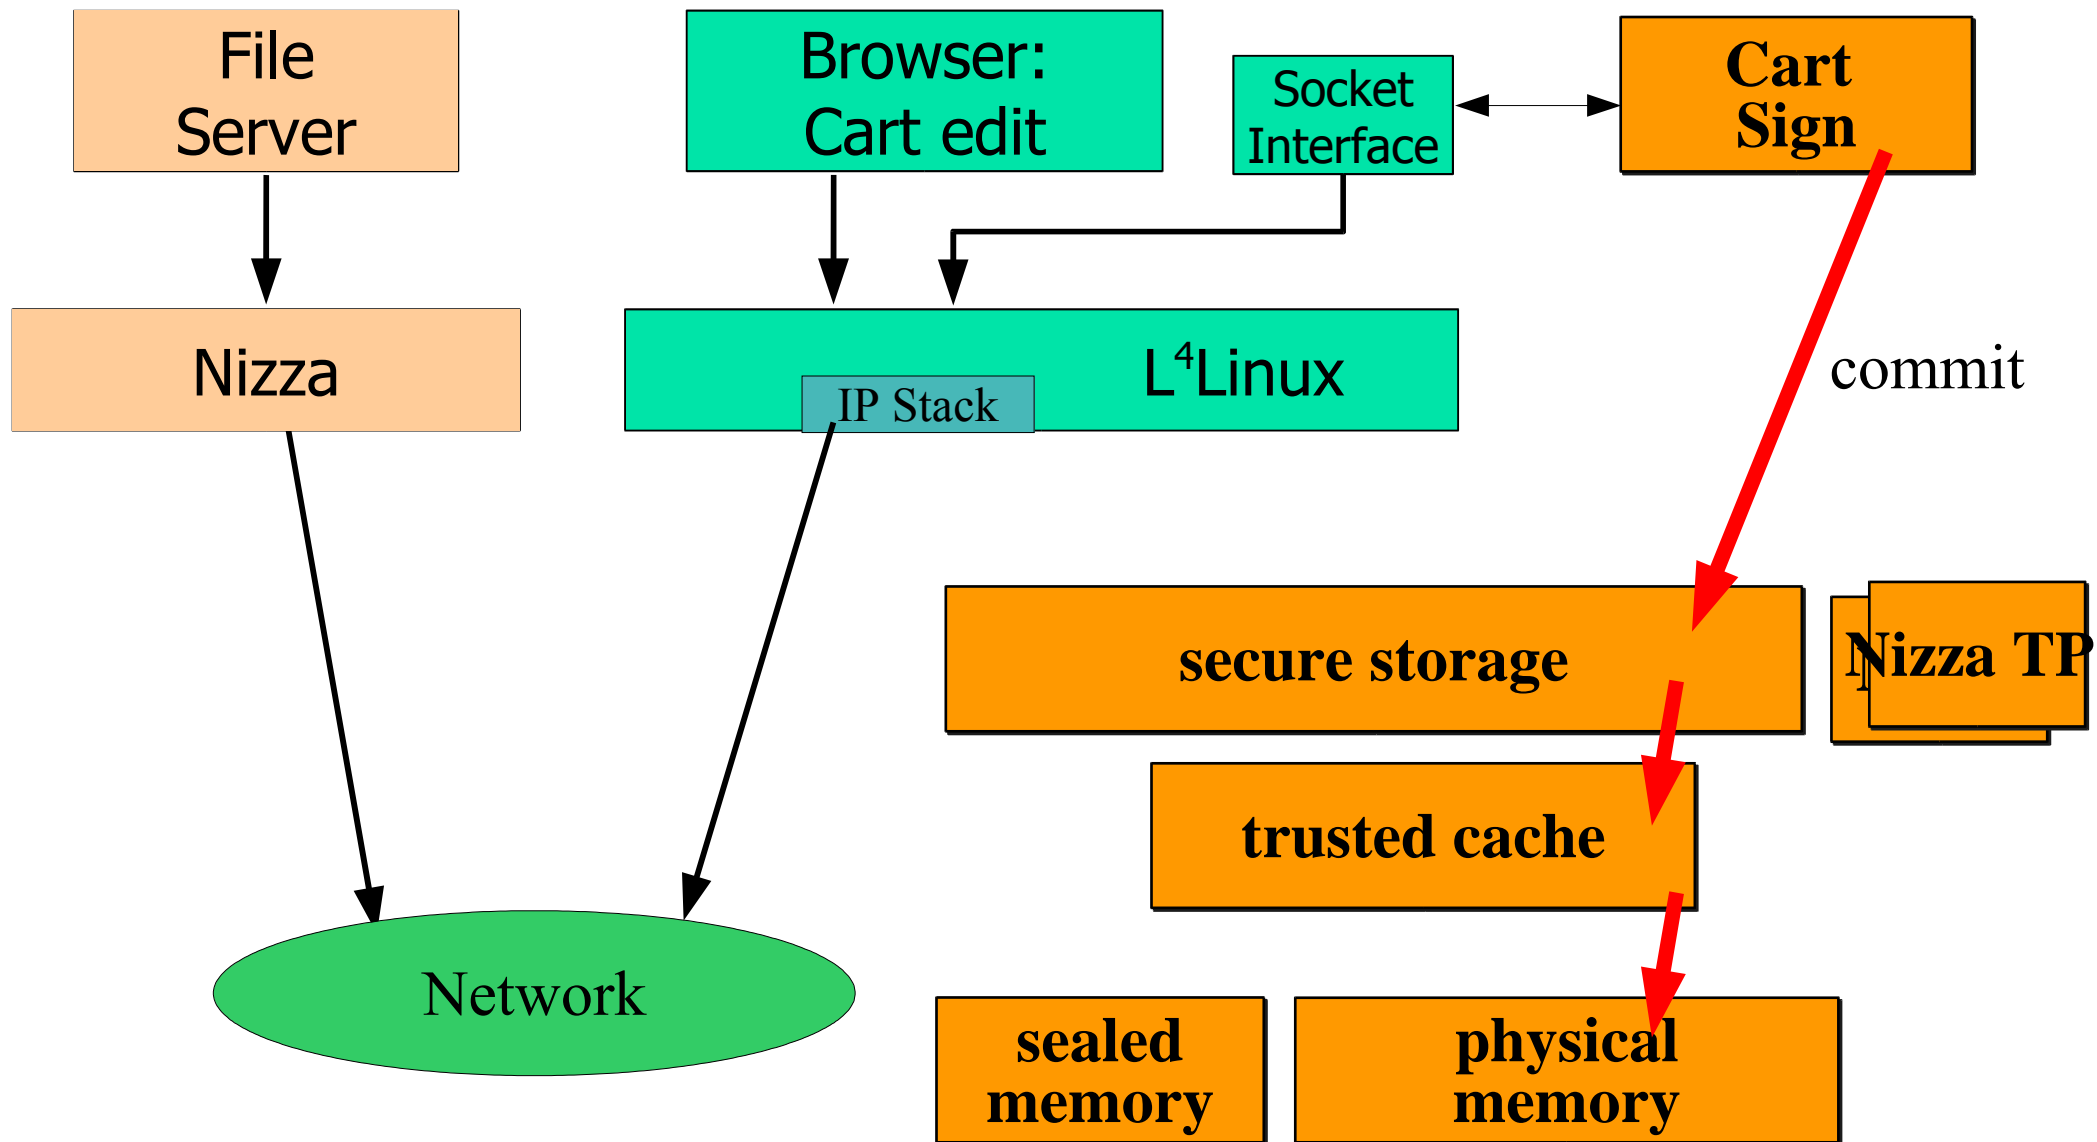
Confidentiality and Integrity



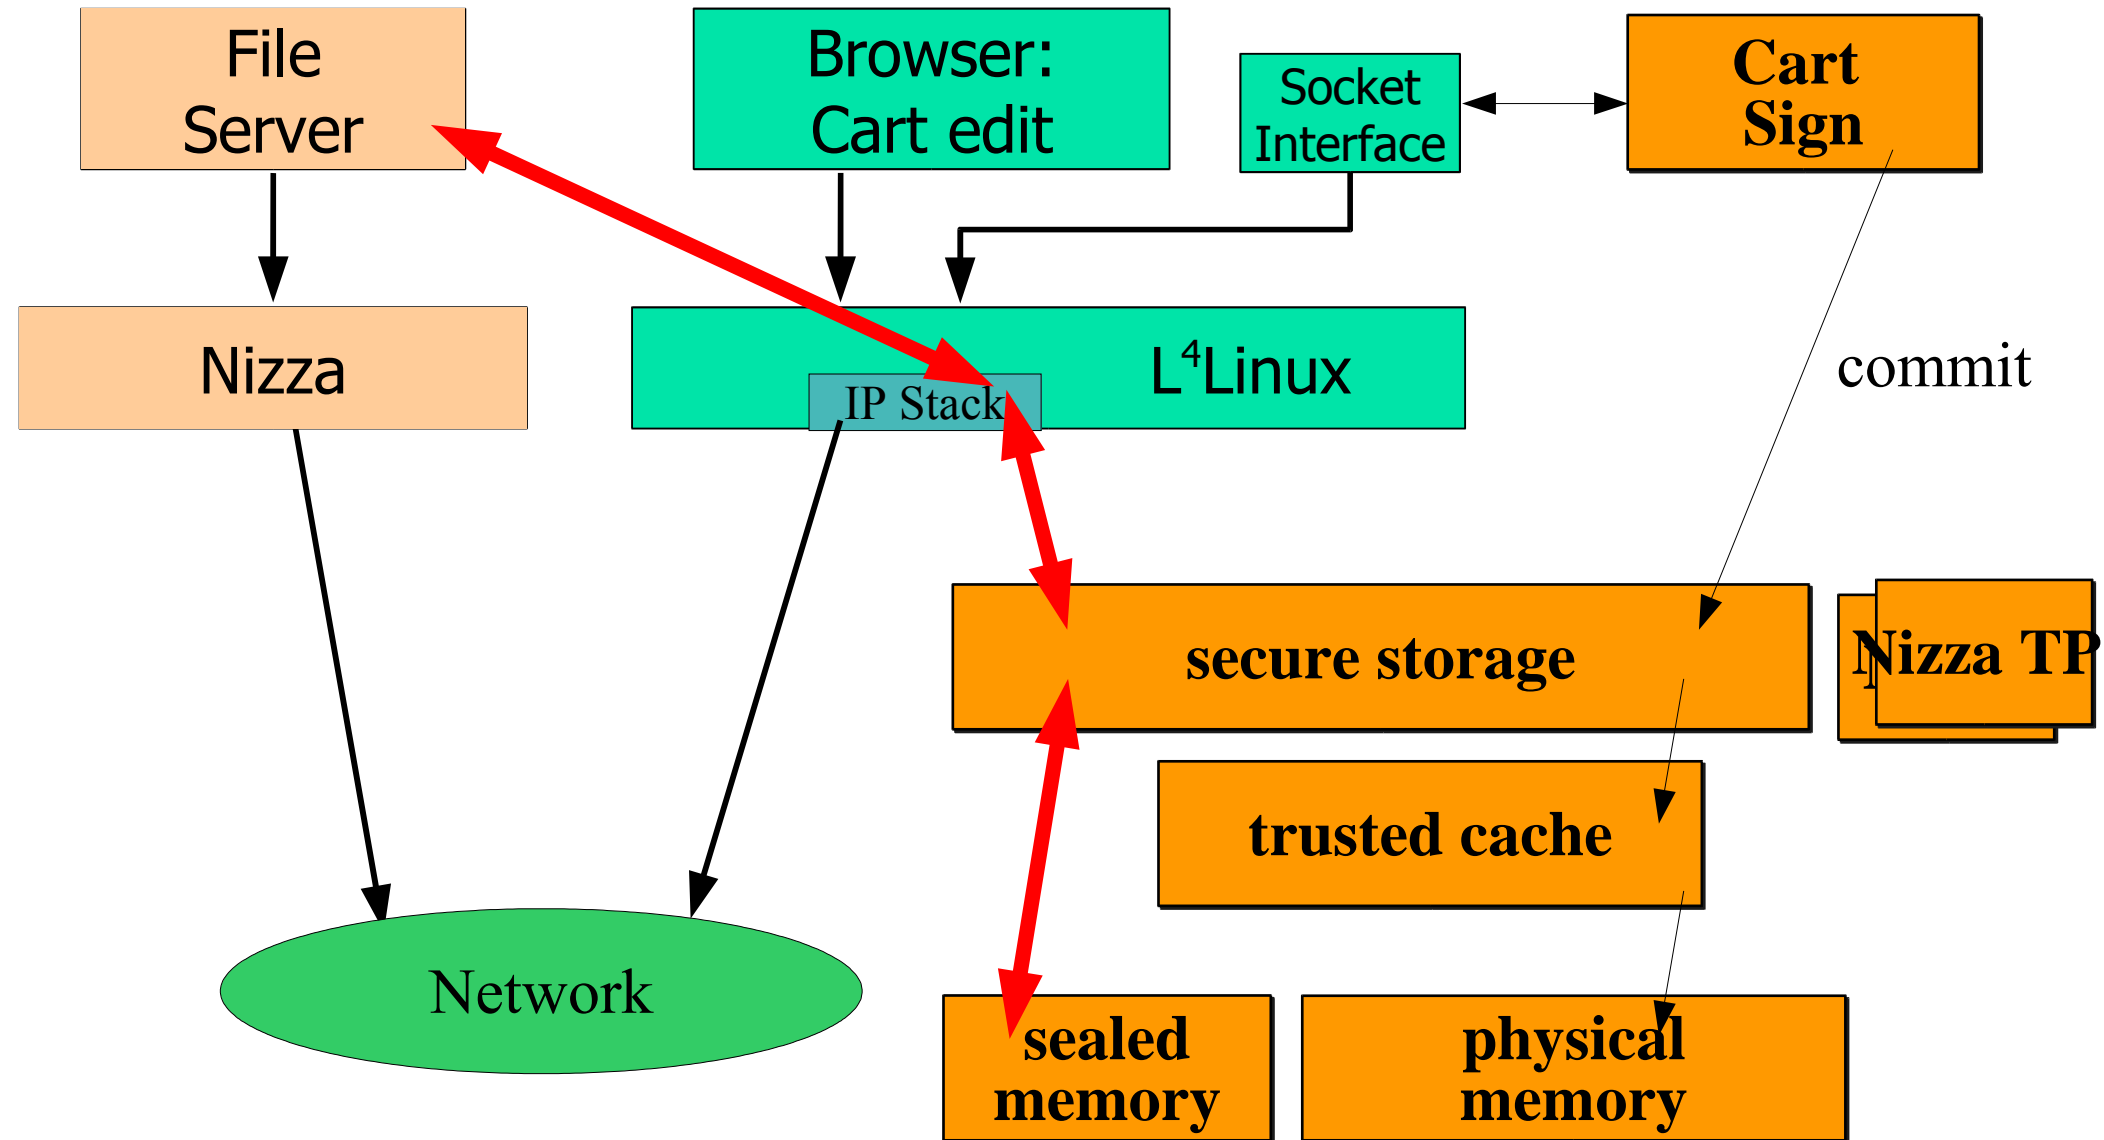
Mass Storage (File System) not part of TCB



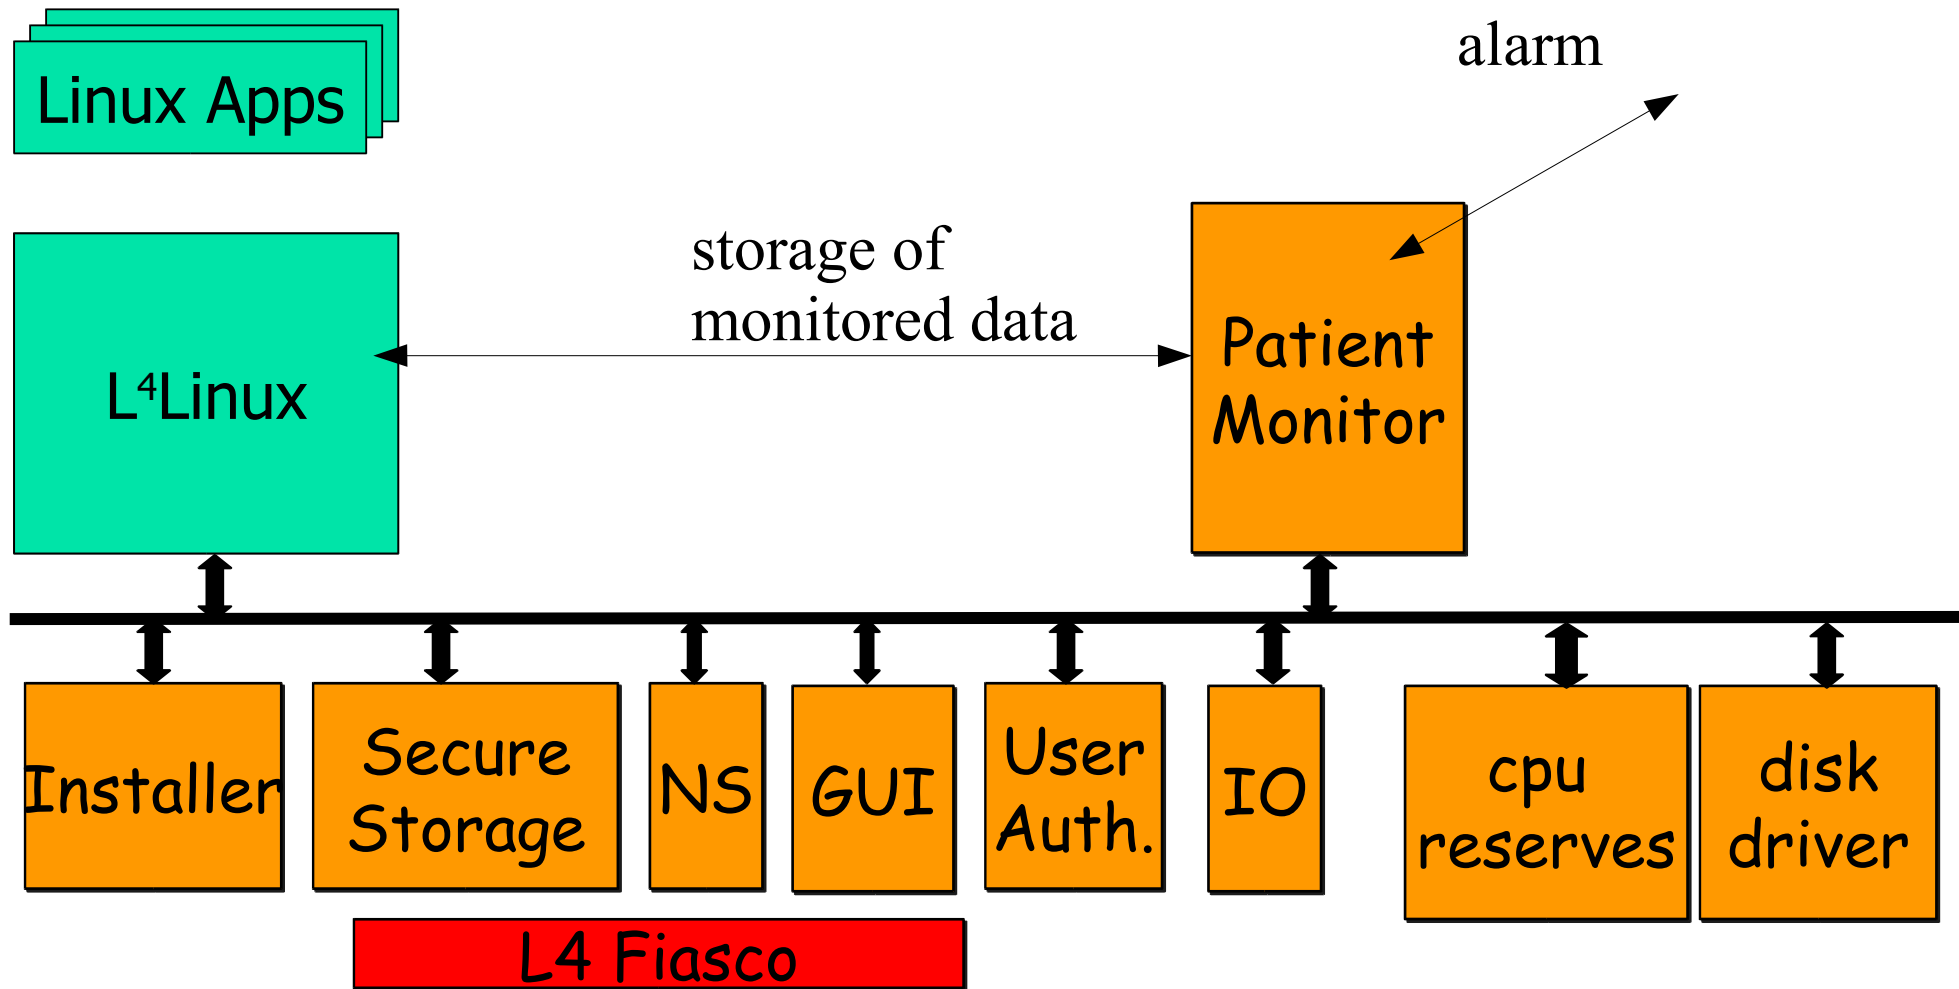
Recoverability



Recoverability



Availability for Secure Storage

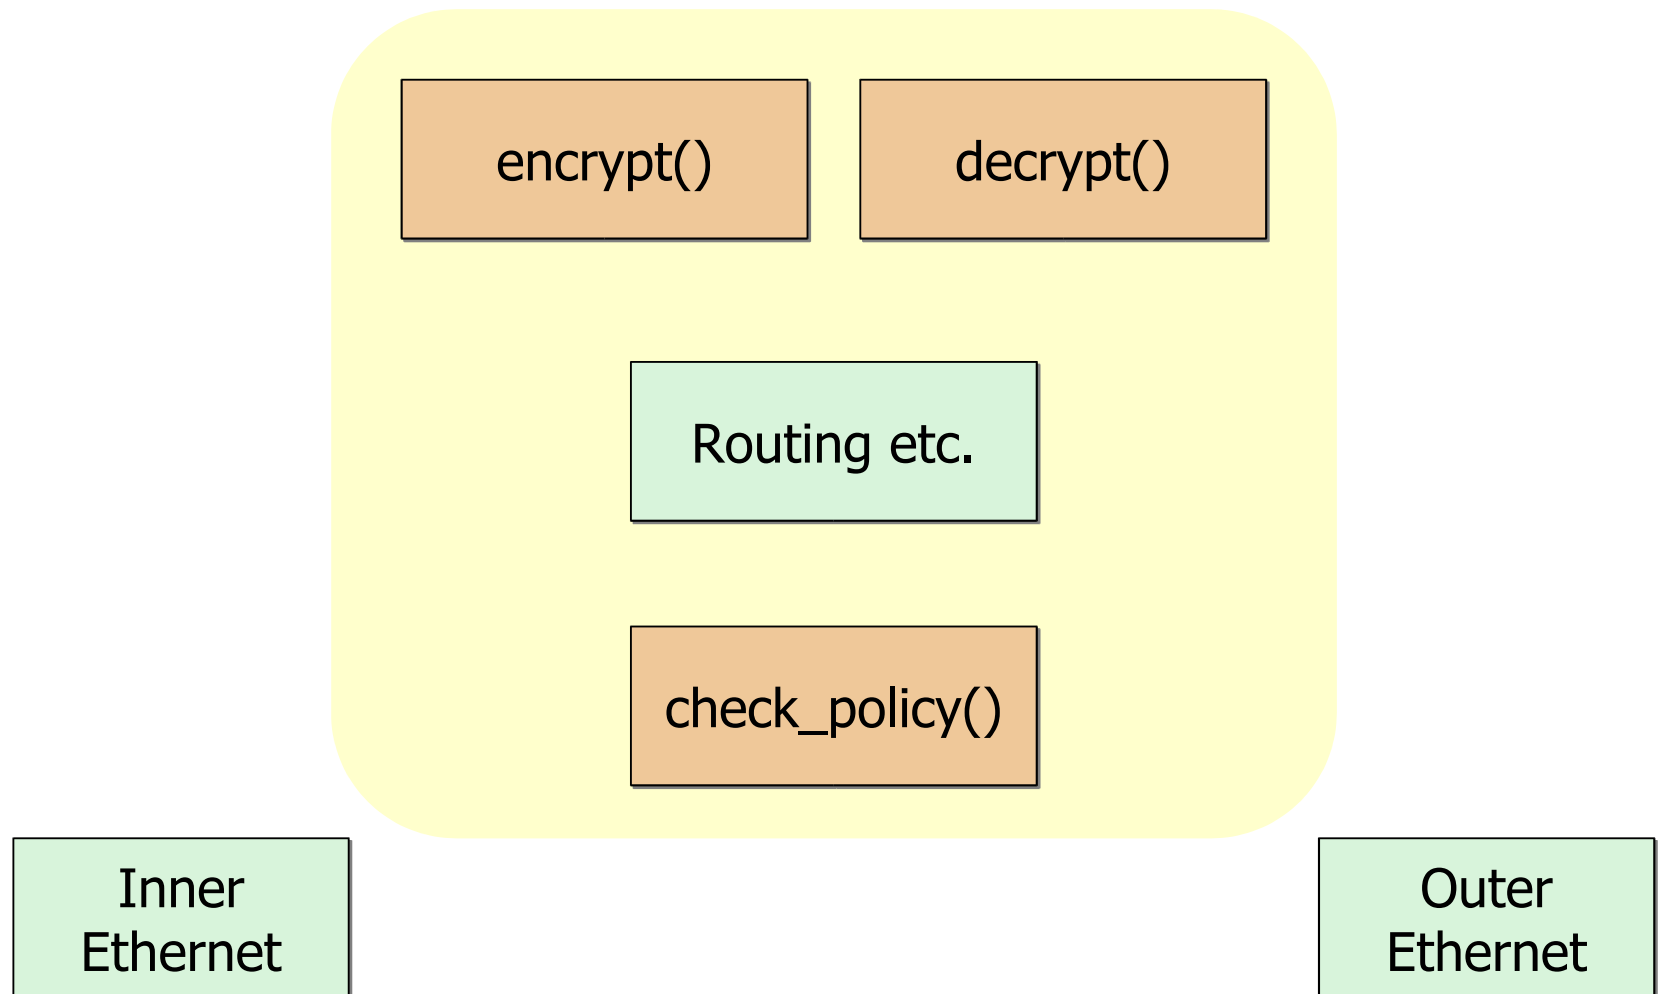


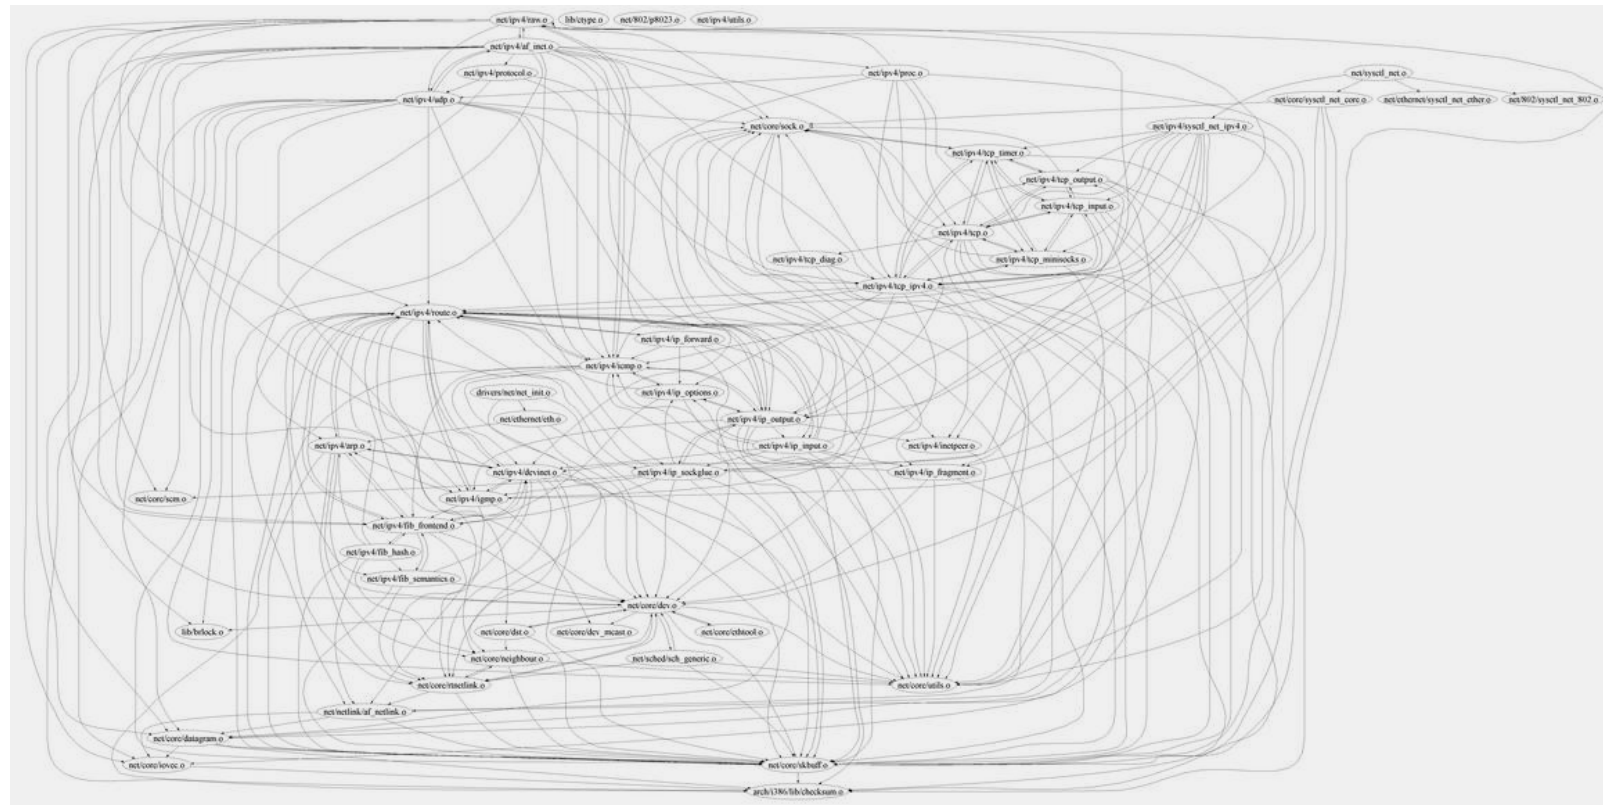
- VPN Box
- reengineering a commercial product
- first approach:
split Linux' IP-Sec



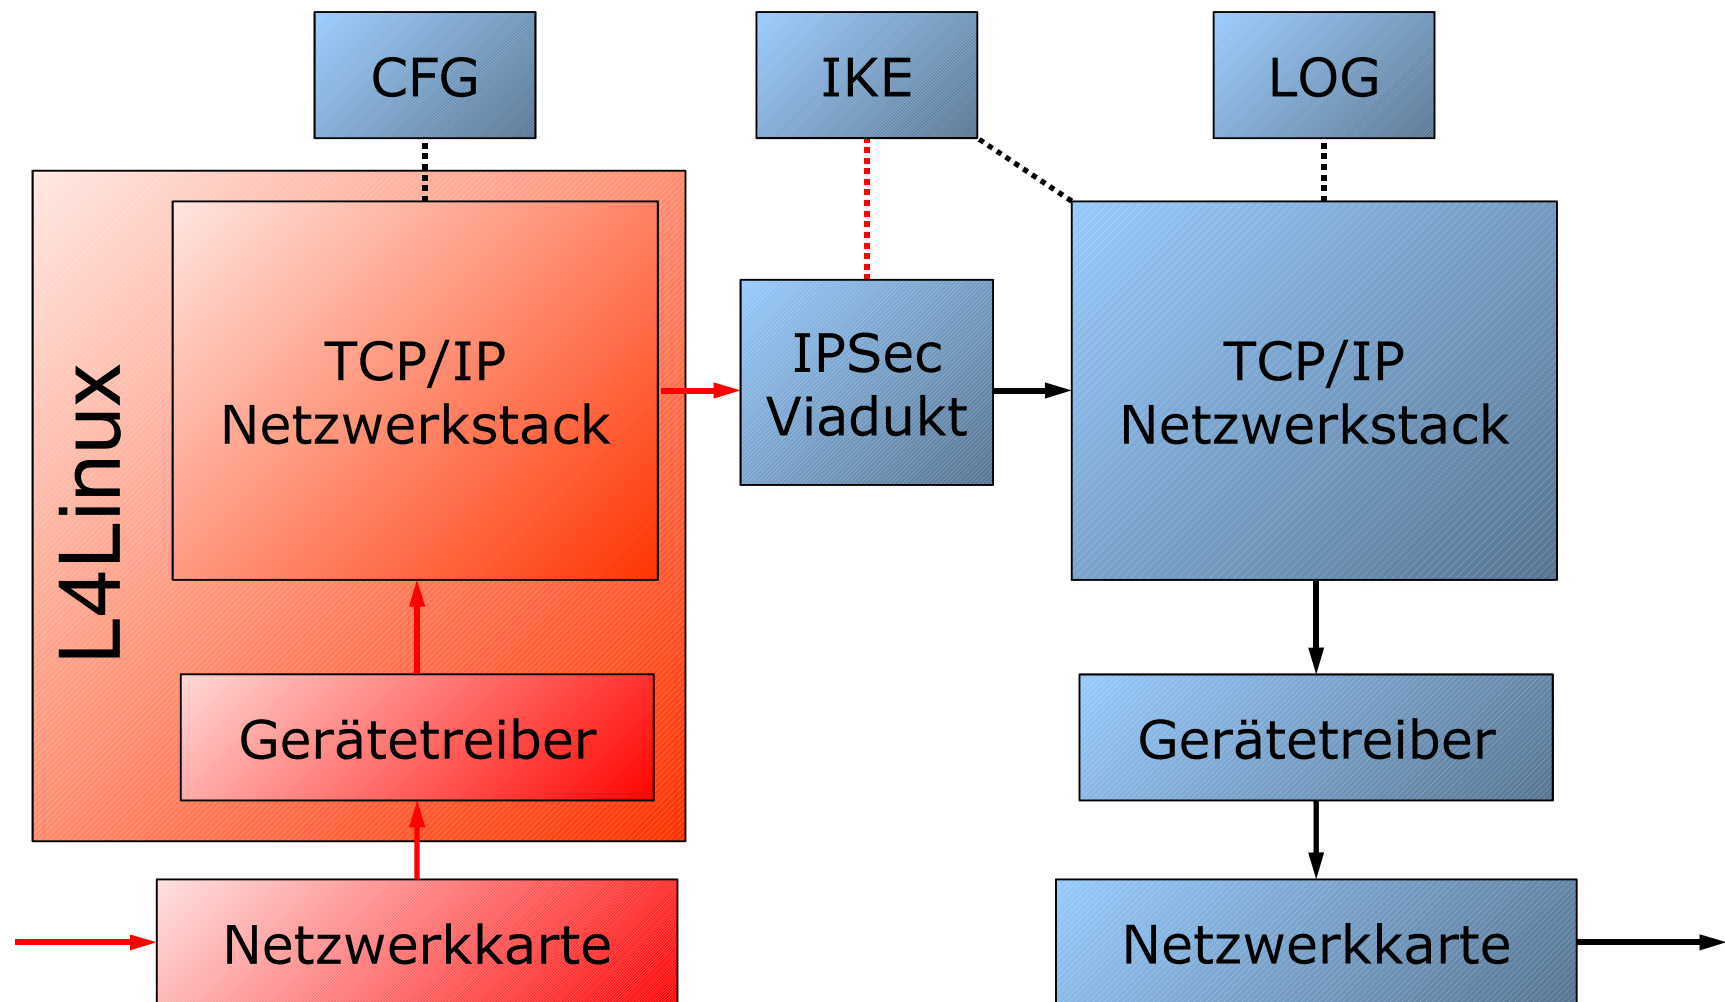
First Approach

- Decompose network stack software on the basis of IP packet flow

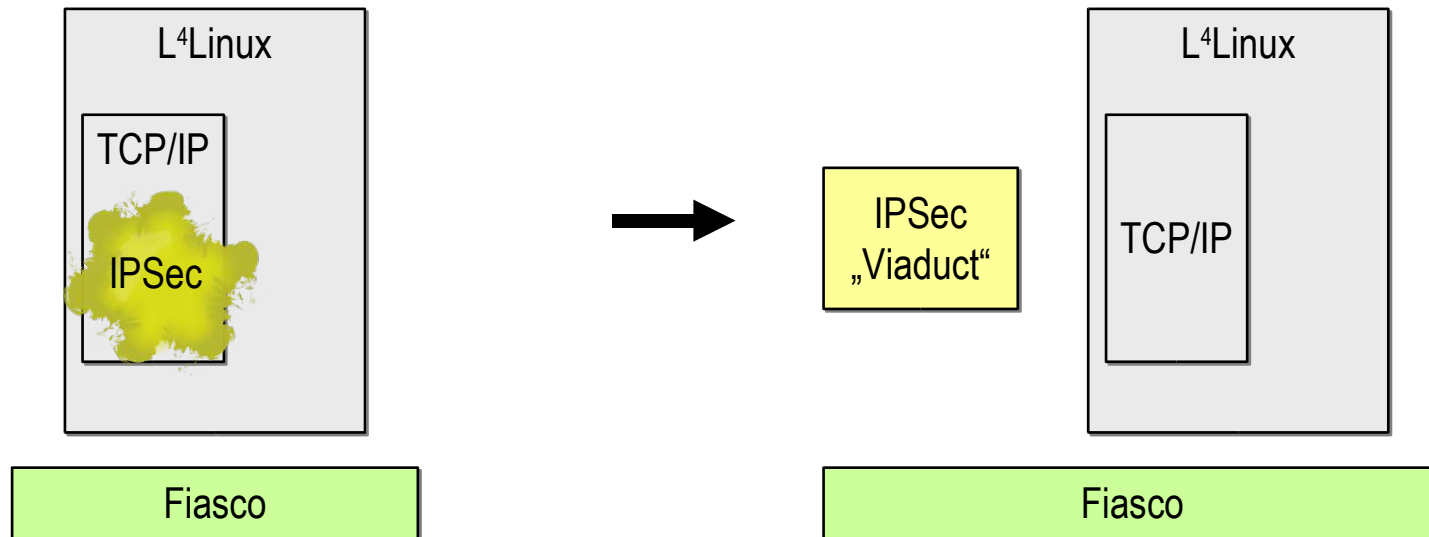




Instead: Rewrite and Trusted Wrappers

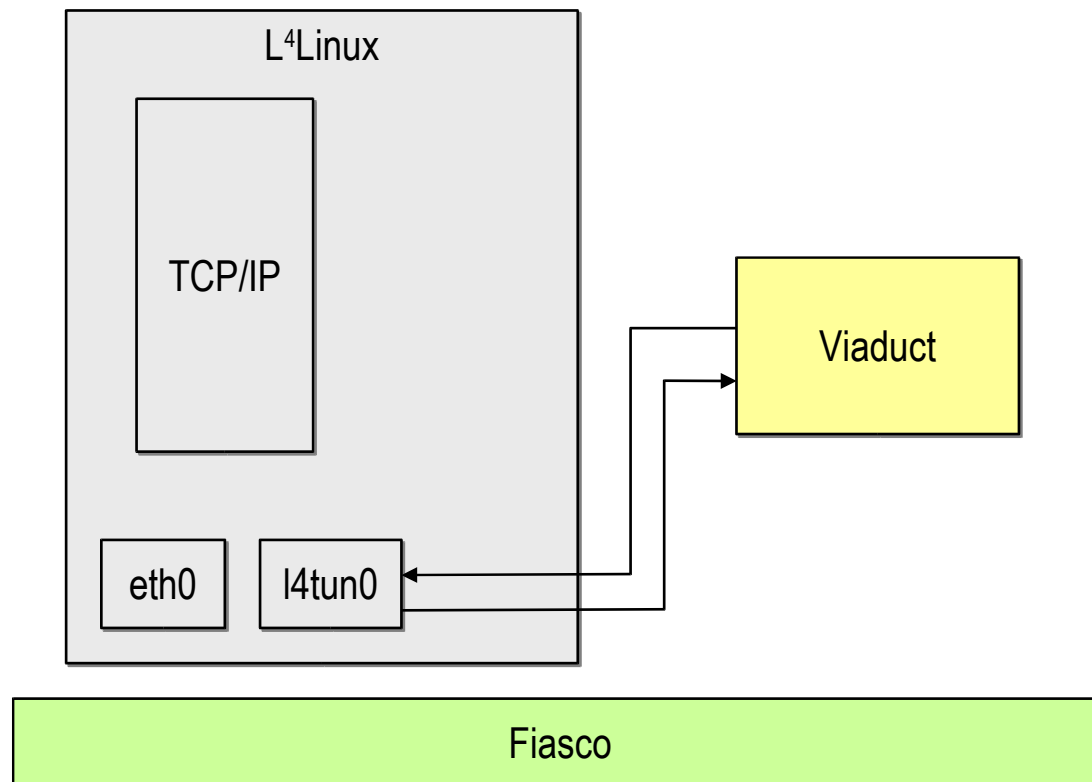


One Step Back



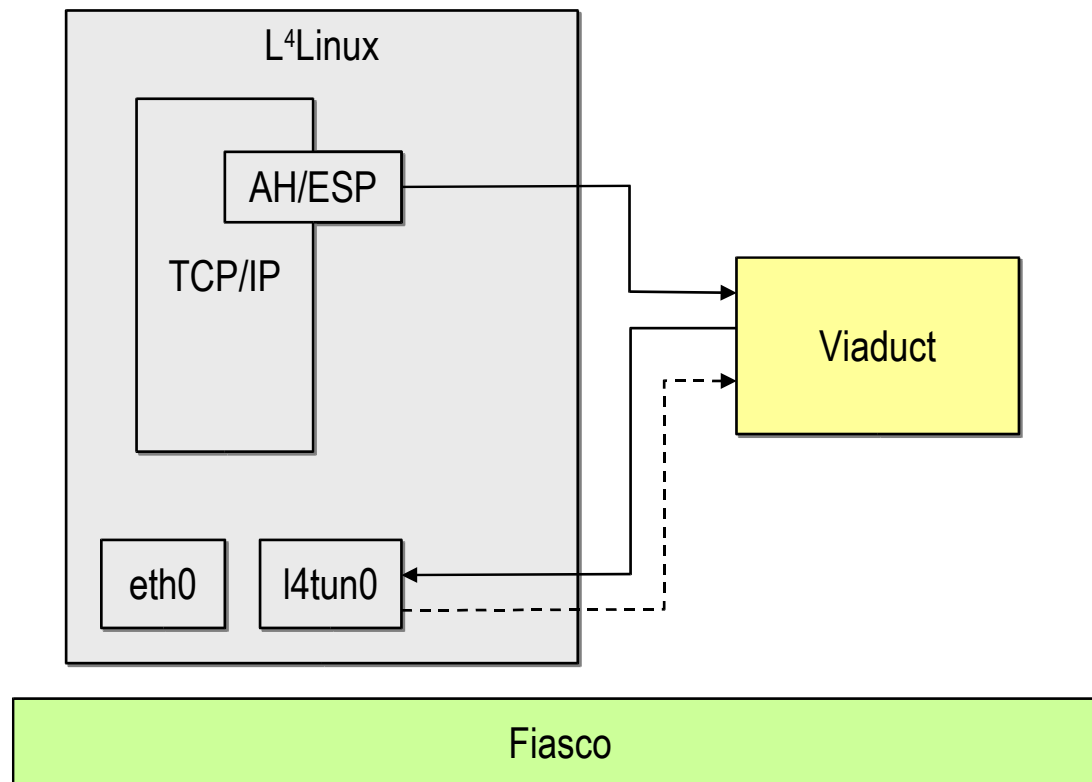
Technical Details (Viaduct)

- IP packets must be passed to the Viaduct
 - L4 IPC as Virtual network driver in L⁴Linux



Technical Details (Viaduct)

- IPSec can only handle unfragmented packets
 - Use L⁴Linux for complex reassembly





Related Work: EROS, Keykos and Nizza

- similar objectives
- moving target



Related Work: Microsoft NGSCB and Nizza

- similar objectives
- moving target



Related Work: XOM and Nizza

- XOM: take OS off trusted path
- Implementing an Untrusted Operating System on Trusted Hardware
David Lie Chandramohan A. Thekkath Mark Horowitz
SOSP 2003



Related Work: Terra and Nizza


- Terra: VMM as trusted platform
- many more projects down that line
- A Virtual Machine-Based Platform for Trusted Computing
Tal Garfinkel, Ben Pfaff, Jim Chow, Mendel Rosenblum and Dan Boneh
SOSP 2003



Nizza vs. VMM approaches

- advantages VMM:
 - support+reuse: unmodified legacy OS
 - anything else ?
- advantages Nizza
 - smaller TCP:
no network device emulation and drivers
 - fine grained sharing
 - efficiency (optimized message passing)
- VMM untrusted on Nizza ?

Challenge: Untrusted VMM





Related work “Useful” for Nizza

- Secure Storage on Untrusted Servers Secure Untrusted Data Repository (SUNDR)
Jinyuan Li, Maxwell Krohn, David Mazires, and Dennis Shasha ,
New York University
- Privtrans: Automatically Partitioning Programs for Privilege Separation
David Brumley and Dawn Song,
Carnegie Mellon University



Technical Risks

- performance
 - copying overhead
 - context switching time (hardware)
 - increased memory
(duplication of page tables)



Context switches

- Register IPC between two address spaces (1 x send, 1 x receive; kernel entry with sysenter):

Pentium-III:	600 cycles
Opteron:	700 cycles
Prescott:	2200 cycles

??



Conclusion

technologies are in place
to build much better (securer) systems

need proper integration → Nizza

